Infron

Additional experiments complete · Interpret independent batches

INFRON RESEARCH / INFERENCE PERFORMANCE

NVIDIA DGX Spark

Two-Node System and Inference Performance Evaluation

Qwen3.8 Flash Next Abliterated · Layered measurement of GB10 hardware, the two-node system and model serving

Data snapshot (Coordinated Universal Time) 2026-10-17 08:58:45 UTC

01 · Research background, abstract and findings

Large-language-model inference depends on computation, memory supply, device coordination and request scheduling. A peak arithmetic rating or single-request generation rate does not characterize serving across input lengths, concurrency and cache conditions. This study measures `qwen/qwen3.8-flash-next-abliterated` on two NVIDIA DGX Spark / GB10 nodes through Infron, connecting hardware resources to user-visible responses through layered evidence.

Synthetic tensors, fixed-length native-engine requests, public prompts and business text serve distinct experimental roles. Device timing, engine-stage timing and public responses are analyzed separately. Repeated fixed conditions retain correctness checks, sample sizes and incomplete coverage. Metric and arrival definitions draw on NVIDIA guidance and MLPerf scenario concepts; this is a custom study, not a certified benchmark submission. NVIDIA · MLPerf

Test domains, items and principal findings

The table summarizes completed measurements and supporting analyses. Domain links lead to detailed results; compute and memory-bandwidth ranges span the two node medians under fixed conditions. Prompt banks are workloads, not answer-quality tests.

Table 01-01 **Test domains, items and principal findings**

Test domain	Tests covered	Principal finding and scope
GPU compute	Scalar FP32; FP32/TF32-allowed/ FP16/BF16/FP8/NVFP4 matrices; structured sparsity.	**$8{,}192^3$** dense matrices: BF16 **94.08–94.34**, FP8 **191.88–192.64 TFLOP/s**. Rates depend on shape, data and library path.
Memory and CPU/GPU sharing	Contiguous/strided access, pointer chasing, shared/managed memory and CPU contention.	**8 GiB** contiguous copy: **223–227 GB/s**, counting reads plus writes. CPU contention can reduce GPU data supply; this is not physical DRAM traffic.
Operators and execution	Normalization, softmax, reduction, attention; launches, CUDA Graphs and streams.	Per-size latency baselines are established. Serial/parallel time ratios span **0.95–1.25×**; benefits depend on the tested condition.

Test domain	Tests covered	Principal finding and scope
Two-node communication	RDMA; AllReduce/AllGather/ ReduceScatter; serial versus overlapped compute and communication.	1 GiB AllReduce bus bandwidth: 21.81 GB/s; 840 collective and 120 overlap trials. NCCL GPU-direct networking was disabled.
Engine stages	Prefill (input processing), decode (generation), queues and two-node GPU timelines.	At concurrency one, prefill is 2,427–3,354 tokens/s and decode with 256 outputs is 49.8–60.2 tokens/s. Synthetic inputs; timings include engine scheduling.
Public serving and arrivals	60 prompts from six datasets; concurrency 1/2/4/8, reasoning settings, open arrivals and goodput.	Short prompts, concurrency 1→8: 45.1→158.9 tokens/s, first-content p95 1.65→2.39 s. Higher arrival rates cause queue rejections.
Lengths and business shapes	1K–256K length design space; eight business shapes, input/output ratios and realized lengths.	Longer inputs increase waiting. 290/299 extension calls completed; the largest attained output target was 4,096 tokens. Actual 256K generation is unverified.
Prefix cache and retention	Six prefix-sharing ratios; first/repeat, idle retention, pressure/recovery and periodic reuse.	32K inputs with a nominal 300 s wait: 91.7% cached-token fraction; six 8K periodic rounds were about 100%. These windows do not establish long-term hit rates.
Sustained work and thermals	Separate compute, bandwidth and mixed phases; clocks, temperature, power, cooling and retention.	Both nodes completed 60/60/120 minutes. Mixed-phase cooling occupied 16.4% / 6.0% (nodes 0/1). Cooling affects wall-clock throughput; phases do not establish four-hour continuous serving.
Observability and attribution	Client/engine/node telemetry, GPU profiling, hardware-cache and memory-counter diagnostics.	Timing boundaries are separated. Counters are partly readable; representative physical-memory traffic is missing, so a unique model bottleneck is not established.
H200 reference	Compute, memory supply/capacity, device links, network and power compared by function.	H200 has higher specified memory bandwidth and different link/power boundaries. This is an official-specification reference without matched H200 inference measurements.
Cost and investment return	Paid API charges; US$10,000 hardware, electricity, self-use/sales and 1/3/5-year returns.	For short prompts at concurrency eight and 75% productive utilization, self-use payback is 5.48 years, five-year ROI −8.7%; sales at twice the blended API reference give 2.65-year payback. Other operating costs are excluded.

Coverage boundary: execution and archiving have ended. The public-length list retains eight unsent requests and one incomplete response. Three of five candidate interference pairs ran, with zero fully observed pairs. Hardware, native long output, profiling, immediate cache reuse and the continuous hour retain their completed lists; these do not establish validation of every proposed condition.

Additional native-engine result: fixed output has been validated through 260,096 tokens. Public natural-stop data remain a separate cohort; native results do not replace end-to-end measurements.

02 · Experimental system and deployment

Model identity and weight provenance

Table 02–01 Model identity and weight provenance

Item	Definition and scope
Public model identifier	qwen/qwen3.8–flash–next–abliterated
Weight repository and revision	Weights published by dealignai; `be794b990578ef3031eccf9f28e675a289a09ee9`
Declared base and modification	The author declares Qwen/Qwen3.8–Flash–Next as the base and a weight–level abliteration modification. This provenance statement is not evidence of behavioral or quality changes in this study.
Internal architecture identifier	Qwen4ExpForConditionalGeneration / qwen4_exp; Internal class names differ from the public model name and are preserved verbatim.
Text–network structure	48 layers, hidden dimension 2,560; 512 routed experts with 10 selected per token; 36 linear–attention and 12 full–attention layers. Expert counts are not parameter counts; verified total/active parameter counts are unavailable.
Quantization scope	The card specifies NVFP4 W4A4 for routed experts (four–bit weights and activations); attention, shared experts and excluded modules are not all four–bit. Quantizer metadata identifies ModelOpt 0.46.0 / NVFP4.
Modalities and context	The source includes text and vision components; only text input and streamed text output are tested. The configured context ceiling is 262,144 tokens, accommodating input and generation; it is not a demonstrated serving limit.
Provenance verification boundary	Model card, configuration, tokenizer and chat template are archived; the baseline node path points to this repository revision. All weight shards were not independently rehashed. The card records qwen–community–license–1.0.

Pinned model configuration · Local provenance manifest and hashes

Endpoint and data scope at a glance

Table 02–02 Endpoint and data scope at a glance

Hardware	Specification
Nodes and chip	2 × NVIDIA DGX Spark / GB10 Grace Blackwell
CPU per node	20 Arm cores: 10 × Cortex–X925 + 10 × Cortex–A725
GPU per node	Blackwell; 6,144 CUDA cores; fifth–generation Tensor Cores
Unified memory	128 GB LPDDR5x per node; 256 GB physical capacity across two nodes
Memory bandwidth per node	273 GB/s; 256–bit
Rated compute peak per node	Up to 1 PFLOP/s (FP4 with sparsity)
Network interfaces	ConnectX–7: rated 200 Gb/s; Ethernet: 10 GbE
Power ratings per node	GB10 thermal design power 140 W; power supply 240 W

NVIDIA hardware specifications · NVIDIA networking specifications

Table 02–03 Endpoint and data scope at a glance

Core inference settings	Value
Platform endpoint	POST `https://api.example.invalid/v1/chat/completions`
Model and route	`qwen/qwen3.8-flash-next-abliterated; infron/cn;` cross–provider fallback disabled
Engine and parallelism	SGLang 0.5.21; two nodes, tensor parallelism = 2
Quantization and loading dtype	NVFP4 / `modelopt_fp4; dtype=auto; kv_cache_dtype=auto`
Compute backends	Attention: FlashInfer; mixture of experts: FlashInfer CUTLASS
Context and prefill chunk	262,144–token context ceiling (input + output); 8,192–token prefill chunks
Running slots and memory allocation	Maximum running requests 48; static memory allocation fraction 0.75
Multi–token prediction (MTP)	Enabled; EAGLE path; 3 steps, 1 branch, 4 draft tokens
Prefix cache and scheduling	RadixCache enabled; overlapping scheduling enabled
Test request settings	Temperature 0; streaming; reasoning effort none / low; client concurrency 1 / 2 / 4 / 8

Recorded node configuration

The full launch command and parsed parameters are in Appendix C.

Experimental materials and scope

Table 02–04 Experimental materials and scope

Material / scope	Definition in this study
Natural prompt bank	Six public datasets, ten prompts each: 60 unique inputs used as performance workloads.
Controlled serving loads	1K—256K input/output design space, eight business shapes, prefix ratios and retention diagnostics; realized coverage is in section 05.
Engine and hardware	231 formal native–engine requests; two–node compute, memory, operators, communication and segmented sustained loads.
Reference and economics	H200 official specifications provide an architectural reference; purchase price, electricity and utilization define cost scenarios.

03 · Experimental design, materials and workloads

System performance reflects computation, data supply, cross–device coordination and request scheduling. Four measurement layers are used: hardware microbenchmarks, standalone operators, the two–node system and end–to–end model serving. Each fixes its workload and timing boundary before matched conditions support interpretation. Hardware tests use synthetic tensors; public prompts supply serving loads. Operator numerical checks establish implementation correctness, separately from model–answer scoring.

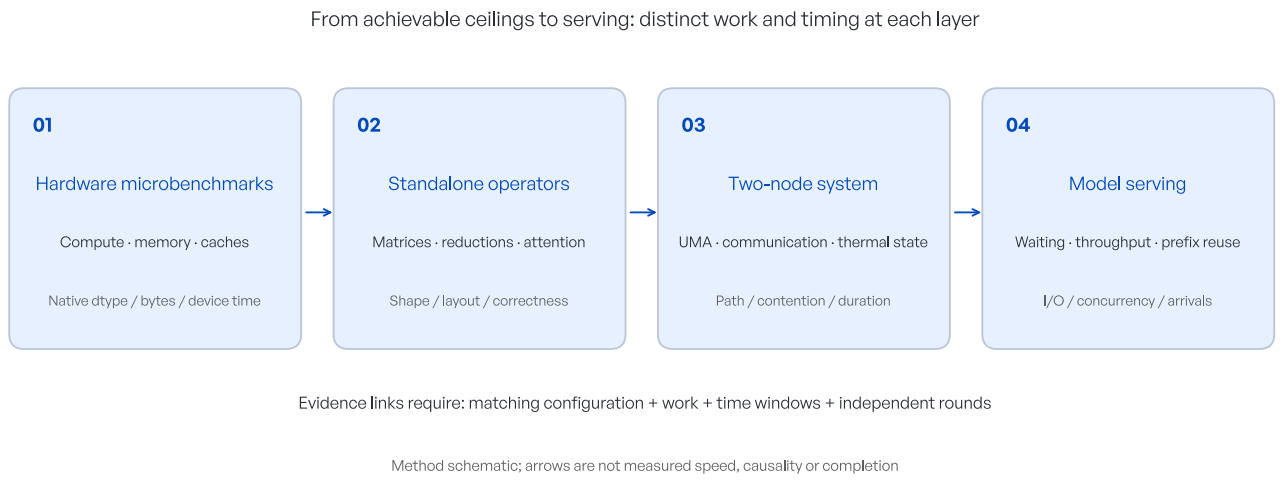


Figure 03–01 Study layers and evidence links

Schematic. Arrows indicate interpretive dependencies, not completed tests or directly subtractable timings.

Table 03–01 Experimental design, materials and workloads

Layer	Measurement boundary	Supported scope
GPU hardware	Device arithmetic and memory access	Achievable ceiling at a specified precision, size and state
Standalone operators	Separate preparation, kernel execution and synchronization	Implementation efficiency, shape sensitivity and data–movement cost
Two–node system	CPU/GPU cooperation, collectives and sustained load	Contention, communication cost and retained performance
Model serving	Client through Infron to the two nodes and back	Waiting, throughput, cache reuse and cost on the actual path

GB10 uses shared LPDDR5X unified memory. Its 128 GB is not an exclusively allocatable GPU pool; 256 GB across two nodes is not one shared address space. The 273GB/s memory rating is per node; 1PFLOP/s is a sparse–FP4 rating. Efficiency comparisons must match precision, sparsity, clocks and traffic definitions. GPU L1/L2 caches count memory accesses; serving prefix caches count reused input tokens. They are reported separately. NVIDIA

Concepts needed to read hardware results

Table 03–02 Concepts needed to read hardware results

Concept	Definition and measurement relevance
CUDA cores and Tensor Cores	CUDA cores execute general parallel arithmetic; Tensor Cores accelerate specific matrix operations. FP32/FP16/FP8/FP4 identify numeric formats; fewer bits do not guarantee faster execution for every task.
Matrix multiplication and M, N, K	GEMM means general matrix multiplication: A is M×K, B is K×N and output is M×N. Shape affects parallelism and reuse and must accompany precision.
GPU kernels and CUDA Graphs	A kernel is a program executed on the GPU; a stream orders work. CUDA Graphs replay captured dependencies and can reduce repeated host submission overhead.
L1/L2, shared memory and unified memory	L1/L2 are hardware caches; CUDA shared memory is on–chip block–local storage; GB10 unified memory is CPU/GPU–shared system memory. Their capacity, latency and access rules differ.

Concept	Definition and measurement relevance
RDMA and GPU–direct paths	RDMA lets NICs remotely access registered memory; a GPU–direct path additionally requires evidence of NIC access to GPU buffers. An ordinary RDMA test does not establish that path.
NCCL collectives	NCCL is a GPU communication library. AllReduce combines inputs and gives each rank the result; AllGather collects every rank's shard; ReduceScatter combines inputs and distributes result shards.

Benchmarks: sources, task forms and use in this study

Six public benchmarks form the fixed prompt bank; length, business and cache workloads add controlled system conditions. The original benchmarks evaluate model answers. This study uses their prompts as repeatable inputs and measures waiting time, actual generation, throughput and node resources. Benchmark names identify input provenance, not completion of their full capability evaluations.

Table 03–03 **Benchmarks: sources, task forms and use in this study**

Benchmark and original purpose	Input construction and performance use	Example of a selected task
MMLU–Pro Multiple–choice knowledge and reasoning across 14 subjects; normally evaluated by answer accuracy.	Question and original choices, without few–shot worked examples. Supplies short inputs with specialist terms and option lists.	Choose the winding required by DC machines from the supplied options. `mmlu_11888`
MATH–500 A 500–problem subset of MATH covering algebra, geometry, number theory and other topics; the original task evaluates mathematical correctness.	Keeps problem text and mathematical notation, requesting a final answer in \boxed{}. Probes output length and waiting for mathematical generation.	Count integer solutions of an absolute–value inequality. `math_386`
GSM8K Grade–school word problems involving multi–step arithmetic; normally evaluated by the final numeric answer. Uses main/test.	Keeps the word problem and asks for a solution ending with a ####–marked answer. Combines narrative text with arithmetic steps.	Calculate a person's age from multiplicative relationships among three ages. `gsm_1104`
IFEval Instruction–Following Evaluation uses verifiable language, format, length and keyword constraints; the original task checks compliance.	Uses the original instructions, retaining language and format constraints. Observes response shape and actual generation volume.	Write a poem about Yoda entering another universe, using Persian only. `ifeval_1477`
HumanEval+ Python function generation; "+" denotes additional EvalPlus correctness–test inputs. The original task checks generated programs by execution.	Keeps signatures, documentation and in–prompt examples; requests a complete Python code block. Measures streamed function generation.	Implement largest_divisor(n), returning the largest divisor smaller than n. `HumanEval_24`
MBPP+ Basic Python programming tasks, using the 378–task EvalPlus version with expanded tests. The original task checks functional correctness.	Keeps task descriptions and in–prompt assertion examples; requests a complete Python code block. Supplies shorter code–generation inputs.	Implement a function returning the smaller of two numbers, matching the supplied assertion. `Mbpp_404`

Examples paraphrase selected tasks; requests use the exact frozen prompts. Reference solutions and hidden tests are excluded, while examples and assertions already in the prompt remain. HumanEval+ and MBPP+ extended tests are not executed, and "+" adds no generation requests here. Accuracy, instruction–compliance rates and code–test pass rates are not calculated.

Data versions, sample sizes and length distributions

Source size refers to the archived version or test split; the candidate frame contains prompts eligible for selection; selected counts are unique prompts. Tokens are tokenizer units, not words or characters. Lengths below use this model's tokenizer and count prompt content only; actual API input also includes the chat template and role markers.

Table 03–04 **Data versions, sample sizes and length distributions**

Pinned data source and version	Source → candidates → selected	Content tokens: min / median / max
MMLU–Pro Revision: `b189ec765aa7`	12,032 → 1,400 → 10	70 / 166 / 386
MATH–500 Revision: `6e4ed1a2a79a`	500 → 500 → 10	36 / 64 / 188
GSM8K Revision: `740312add88f`	1,319 → 1,319 → 10	46 / 76.5 / 112
IFEval Revision: `0a1b0fe7298a`	541 → 541 → 10	24 / 43.5 / 83
HumanEval+ EvalPlus 0.3.1 SHA–256: `42526ec0e7d5`	164 → 164 → 10	77 / 153 / 315
MBPP+ EvalPlus 0.3.1 SHA–256: `b54e76275524`	378 → 378 → 10	55 / 72.5 / 144

Version cells show the first 12 characters of a commit or data–file SHA–256. Links identify the source; complete revisions and local file hashes are in the source manifest. EvalPlus 0.3.1 identifies the retrieval package; each dataset file is separately pinned by its hash.

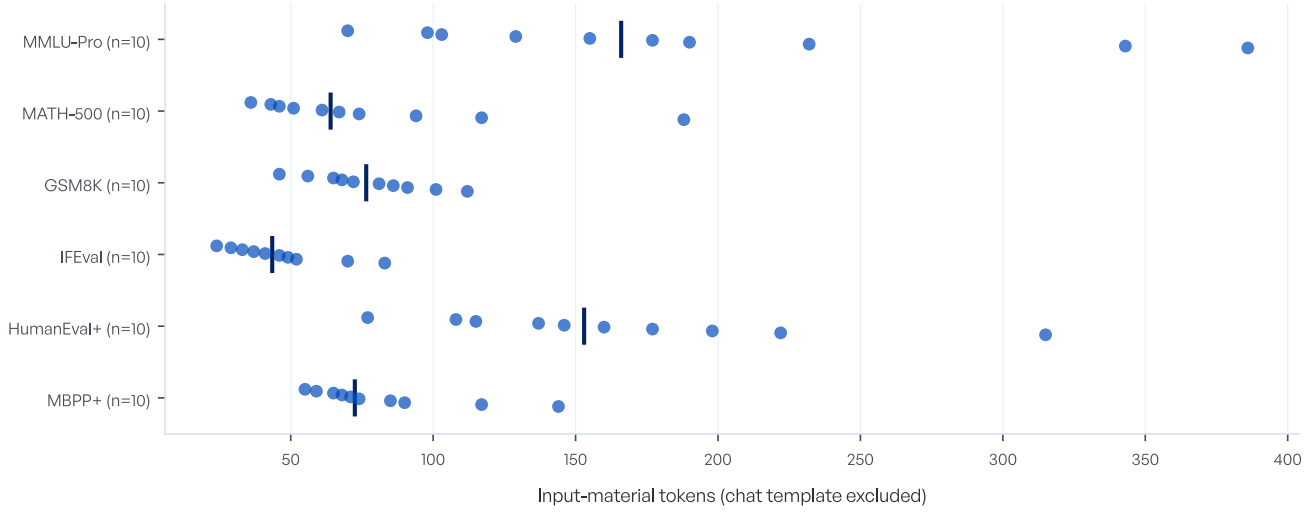


Figure 03–02 **Prompt–bank lengths**

Each point is one fixed prompt, ten per suite; vertical marks show medians. The 60 prompts span 24—386 content tokens and primarily characterize short inputs. Separate length workloads measure long–input behavior.

Sampling method and comparison scope

The candidate pool contains 4,302 items. MMLU–Pro contributes 100 items from each of 14 subjects, selected from its 12,032–item test split, yielding 1,400 candidates. The other five suites use the full versions or test splits listed above. Within each suite, candidates are sorted by content–token length and divided into ten nearly equal–count strata. A SHA–256 ranking of the fixed seed `20261014-performance-v2`, suite and item ID selects one per stratum. The result is ten per suite, 60 total, without using model outputs.

The selected MMLU–Pro items cover engineering, math, physics, chemistry, health, economics, computer science and law: two each for physics and economics, one each for the others. MATH–500 selections cover five of its seven source topics. Length stratification preserves variation in prompt size, but ten items per suite do not support subject rankings or generalization to the full suites. Equal suite counts also do not represent production traffic shares.

The same 60 prompts run for three rounds at each concurrency level 1/2/4/8, giving 720 planned measurements. Low–reasoning–budget arms add three rounds at concurrency 1/4, or 360 planned measurements. Repetitions compare performance without increasing unique prompts. The natural–bank output cap is 1,024 tokens; generation may stop earlier. Suite comparisons therefore include actual output lengths and cap–hit fractions. When content, output volume and cache state vary together, latency differences cannot be attributed solely to question difficulty.

Fixed 60 prompts · Selection rules · Source–size and sample audit

Complementary workloads: length, business and caching

The short–prompt bank does not cover long contexts, long generation or cache retention by itself. The controlled workloads below add these conditions. Counts are planned calls, including the relevant warmups, calibration or controls; Section 05 and the results report actual completion. Here 1K = 1,024 tokens and output length is a generation cap.

Table 03–05 **Complementary workloads: length, business and caching**

Workload and material source	Construction and planned scale	Added performance scope
Length and long output: Project Gutenberg's War and Peace prose and locally generated business records	Nine input/output tiers: 1/2/4/8/16/32/64/128/256K. Of 81 combinations, 63 retain safety margin, one is exactly on the boundary and 17 exceed it. Selected axes, diagonals, near–limit and related conditions total 299 extension calls.	Separates input–processing wait from sustained generation, recording output attainment. The 256K context is shared by input and output; both cannot simultaneously reach 256K.
Eight business shapes: generated tickets, orders, operations records, code and multi–turn messages	Triage, rewriting, drafting, summarization, analysis, extraction, code patching and dialogue. Each has 12 measurements plus one warmup at concurrency 1/4: 192 measurements + 16 warmups = 208 calls.	Compares response and capacity across input/output ratios. Synthetic scenario proportions do not represent production traffic.
Prefix ratios: synthetic records with shared beginnings and varied suffixes	8K/32K inputs × 0/25/50/75/90/100% shared prefixes × six independent pairs; one warmup and one probe per pair = 144 calls. Together with business shapes: 352 calls.	Target overlap controls the shared input prefix; service counters establish actual cache reuse. These are reported separately.
Cache retention and stability: independently generated business prefixes and competing requests	Business cross–cases: 72; idle retention: 36; competition/recovery: 24; periodic reuse: 28; 160 calls, accounted separately from the 144 prefix–ratio calls.	Observes reuse after idle gaps, other requests and repeated access using counter windows and latency. Short diagnostics do not establish long–term stability.

Long–text source: Project Gutenberg, with archived corpus hash. Section 03 gives selected length pairs, business input/output examples, cache intervals and concurrency controls. Appendix C lists all code and data paths. GPU compute, memory and communication use the separate system benchmarks in Sections 06—07 and are counted separately from these model–request workloads.

Model–serving experiments and materials

Each condition fixes input, output cap, reasoning effort and concurrency. The natural bank has three rounds and the length matrix four blocks per cell; seeded orders are saved before execution. Baseline and expanded loads retain natural stopping; realized lengths are verified against upstream counts and rates use actual work.

Fixed–concurrency tests replenish completed requests to compare capacity. Independent–arrival tests schedule requests at predetermined random times to observe queues and rejections. Both keep at most eight requests in flight. Counts, order and stop conditions are in the test plan.

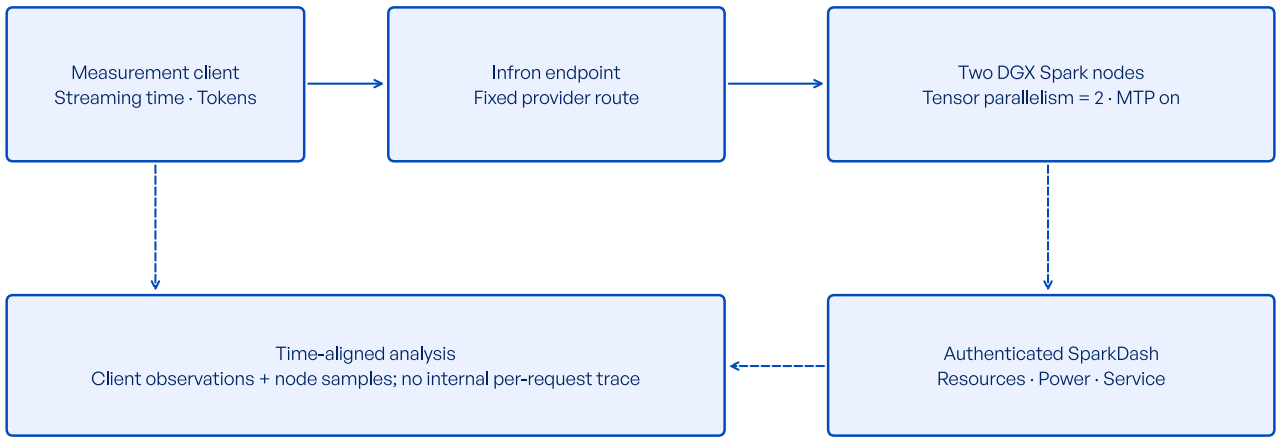


Figure 03–03 Measurement path

The client reaches the two–node service through Infron; telemetry is time–aligned. Client latency includes network and service waiting.

Alignment with SGLang and vLLM benchmark methods

The serving engine is SGLang 0.5.21; vLLM 0.31.0 is a methodology reference. The table separates upstream tool design from this experiment's rules; it is not an engine performance comparison.

Table 03–06 Alignment with SGLang and vLLM benchmark methods

Reference	Implementation and interpretation
SGLang: benchmark layers	Online serving includes HTTP, scheduling and generation. This study additionally includes Infron and the public route; offline–engine or single–batch kernel measurements do not replace service latency.
SGLang: sample count	The guide recommends at least five times maximum concurrency. New repeat blocks use 8/20 calls at concurrency 1/4, with three blocks each. Block variation is still checked; count alone does not establish steady state.
SGLang 0.5.21: EOS · vLLM 0.31.0: EOS	The tools have different EOS defaults, so settings are explicit. Compare natural stopping with ignore_eos=true and verify the forwarded request and upstream output count. Only verified fixed outputs enter length–controlled comparisons.
vLLM: arrival / goodput	Record arrivals, client concurrency and goodput separately. Preserve scheduled arrival, dispatch, first fragment and completion; include queuing and rejection in user waiting and offered load. Current arrivals are Poisson; bursty arrivals are future work.
SGLang: shared prefixes	Shared prefixes are controlled inputs; observed hits require valid fields or counter deltas over isolated windows. Natural tasks and synthetic long output remain separate populations.

Reference documents and source snapshots retain versions and SHA-256 hashes. Framework commands are references; one collector preserves route, budget and evidence consistency. Source archive

Public length-extension requests retain natural stopping: calibration did not forward ignore_eos upstream. Reaching 98%—100% of a cap establishes realized length for that request only. The native-engine experiment separately verifies fixed output in section 08.

Experimental material and length construction

The earlier tables in this chapter document the six public benchmarks, pinned versions, selection rules and length distribution of the 60 prompts. The following describes control of input length, output caps and prefix structure beyond that short-prompt bank.

Length-control material is ordinary prose from Project Gutenberg's War and Peace, sampled at fixed offsets with distinct leading identifiers to reduce cross-request prefix reuse. Source and SHA-256 are retained. This isolates length effects without claiming a real business demand distribution. Business amendments use locally generated orders, tickets, operations records and code, with no actual customer content. Project Gutenberg · Natural selection manifest · Long-text provenance and hash

Length coverage: feasibility from 1K to 256K

1K = 1,024 tokens. Input and output each use nine tiers: 1/2/4/8/16/32/64/128/256K. Input includes the complete chat template; output denotes a requested cap, with actual generation recorded separately. The 262,144-token context is shared by input and output, so neither 256K input plus 1K output nor 1K input plus 256K output fits.

Of 81 design cells, 63 fit with a 1K margin; 128K+128K has no margin and is a boundary; 17 exceed context. Execution selects 21 axis/diagonal cells plus 254K→1K and 1K→254K near-boundary probes. "Length eligible" does not establish gateway output limits, available memory or runtime feasibility.

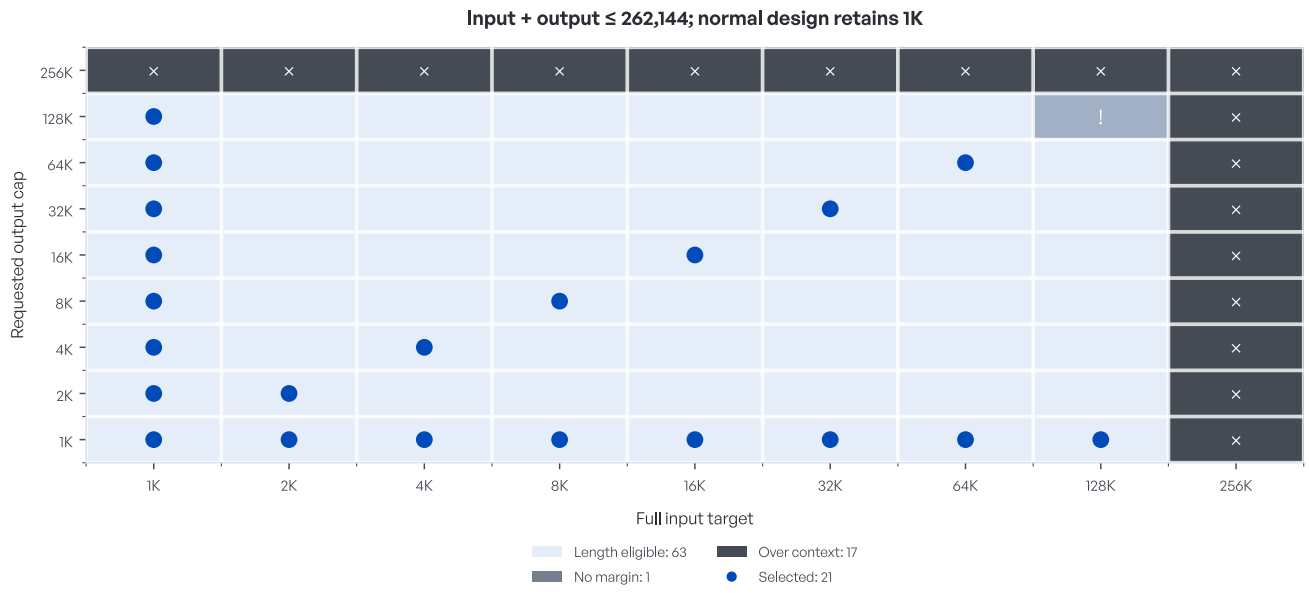


Figure 03-04 Length feasibility

Dots mark selected conditions; colors express context arithmetic, not completed measurements. The 254K near-boundary probes are separate.

Table 03–07 Length coverage: feasibility from 1K to 256K

Group	Input → output cap	Calls	Purpose
Input axis	1–128K → 1K	24	Hold output cap to examine input processing and initial waiting.
Output axis	1K → 2–128K	13	Hold input; three calls per cell through 8K output, one from 16K.
Balanced lengths	2 / 4 / 8 / 16 / 32 / 64K → same	12	Observe joint input/generation growth; high tiers are single probes.
Near boundary	254K → 1K; 1K → 254K	2	One call per cell, conditionally admitted; no tail percentiles.
Long prefixes	64 / 128K → 1K	36	0/50/100% target overlap; three seed/probe pairs per cell.
Business anchors	below	32	One warmup plus three measurements per shape; natural EOS.
Repeated blocks	1K → 4K; 32K → 1K	178	Concurrency 1/4, three blocks of 8/20 calls per cell, plus ten warmups.
Parameter calibration	1K → 1K	2	The same short-answer instruction with natural EOS and ignored EOS.

The extension fixes 299 calls across axes, diagonals and near-boundary combinations; sections 05 and 08 report realized coverage and attainment.

Representative points for the eight business shapes

Table 03–08 Representative points for the eight business shapes

Shape / length	Context	Typical input → illustrative output
Batch ticket triage 4K → 1K	Route several tickets, producing a classification list.	Input: 20 tickets; output: [{id, category, reason}, …]
Document rewriting 4K → 4K	Rewrite similar information volume while preserving structure.	Input: operations procedure; output: customer-guide sections.
Long-form drafting 1K → 8K	Expand a brief into a handbook, emphasizing generation.	Input: launch goals and constraints; output: preparation, execution and recovery.
Long-document summary 32K → 1K	Compress retrieved material or logs into concise findings.	Input: one month of records; output: key events and actions.
Evidence analysis 32K → 8K	Long evidence plus detailed reasoning stresses input and output.	Input: ticket batches; output: trends, evidence and alternatives.
Batch extraction 8K → 1K	Convert documents to structured lists, without constrained decoding.	Input: orders; output: [{order_id, amount, status}, …]
Code editing 16K → 4K	Several functions supply context for a patch and explanation.	Input: billing module; output: idempotency patch and explanation.
Multi-turn assistant 16K → 2K	History plus follow-up retains a multi-message chat structure.	Input: ticket history + next-step question; output: actions.

Outputs above illustrate form. Synthetic long output uses varied identifiers, dates and records; classification is not required to generate a 256K answer. Existing short business workloads remain baselines; new anchors are separate cohorts, not one averaged score.

Admission and interpretation

Fixed-length mode verifies forwarded ignore_eos=true and max_tokens, requiring upstream output to reach 98—100% of cap. Tiers escalate in a fixed protective order, so order and time effects may be confounded. After fixed-output calibration fails, unsent conditions use natural stopping; an unqualified preceding output tier blocks escalation. Input-axis cells have three calls and output probes ⩾16K have only one: neither supports steady-throughput or p95/p99 claims. Representative blocks enter capacity comparisons only after completeness, length and count checks.

Before dispatch, verify full input + cap ⩽262,144 and a conservative concurrency × request budget ⩽262,144. The latter is an experiment guard, not measured KV capacity. Admission estimates 120 seconds + cap×0.08 seconds + input/300 seconds per request, plus 300 seconds drain reserve per block; 0.08 seconds is a policy assumption. Monitoring and cost guards remain; early stopping and skips are retained. Long-prefix seeds/probes each have isolated counter windows; three pairs per cell do not establish long-term retention.

Short-workload baselines for the eight business shapes

Input/output ratio is defined prospectively as target input tokens divided by the output-token cap. It controls experimental shape, but generation may end early; realized input/output ratios are reported separately. Each shape has 12 measurements and one warmup at each of concurrency one and four: 208 calls total. Each condition has three blocks of four measurements, in a fixed randomized order. Concurrency arms share source material but use distinct early identifiers to reduce prefix carryover. This sample supports diagnostics, not stable tail estimates.

Table 03-09 **Short-workload baselines for the eight business shapes**

Shape	Input target → output cap	Nominal ratio	Context and interpretation
Ticket triage	512 → 64	8:1	Short request and answer for routing; emphasizes waiting over generation.
Balanced rewriting	2,048 → 2,048	1:1	Similar input/output budgets for rewriting; examines decode occupancy.
Short brief, long draft	512 → 2,048	0.25:1	A brief expands into long-form drafting; examines output cost.
Long-document summary	32,768 → 128	256:1	Long material compressed into a short summary; emphasizes prefill and first-fragment waiting.
Long-document analysis	32,768 → 2,048	16:1	Long input plus extensive analysis; combines prefill and decode pressure.
Structured extraction	2,048 → 128	16:1	Text to short JSON for extraction; no JSON mode or native tool API is enabled.
Code editing	8,192 → 512	16:1	Code context with a focused patch; medium/long input followed by short output.
Multi-turn assistant	8,192 → 512	16:1	Follow-up assistance with history; uses multiple actual chat messages.

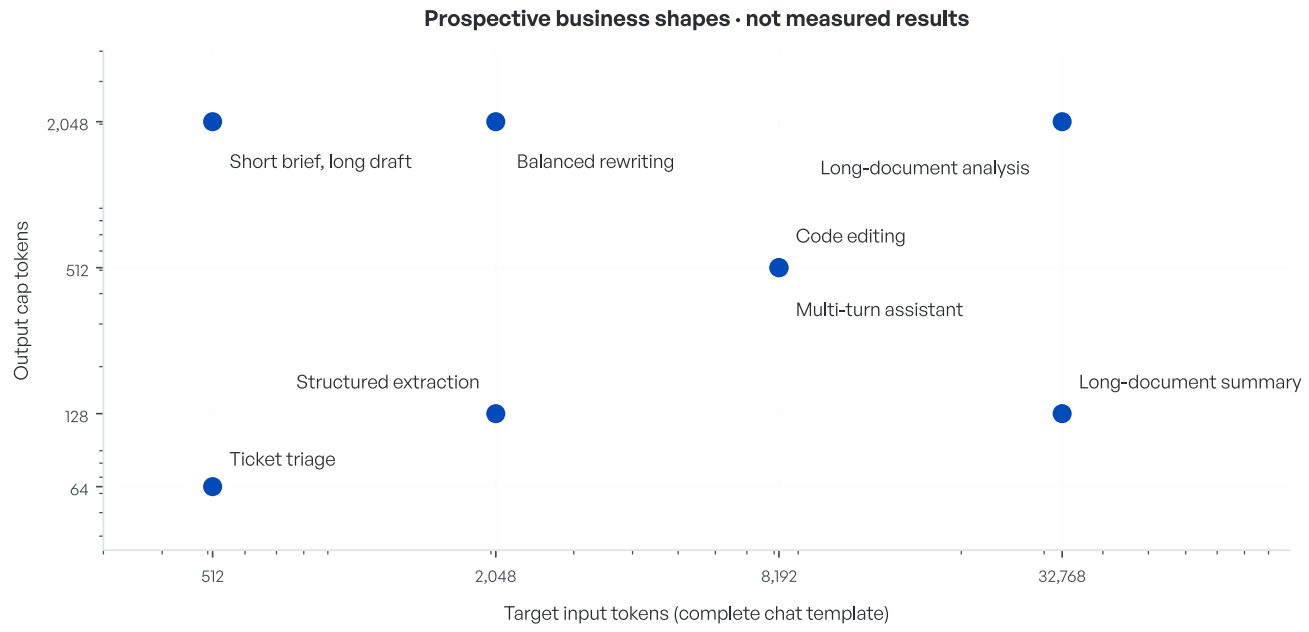


Figure 03–05 **Business–shape design**

Target input and output caps on logarithmic axes. These are prospective conditions; realized output is reported in the results after execution.

Typical input/output examples

The following are shortened teaching examples. Outputs illustrate format and are **not measured model answers**. Actual prompts contain length–controlled generated records; complete prompts and responses remain in per–request evidence.

Table 03–10 **Typical input/output examples**

Shape	Typical input excerpt	Illustrative output excerpt
Ticket triage	The customer paid but access reset failed. Classify the ticket.	{"category":"access","reason":"Access failed after payment"}
Balanced rewriting	Rewrite the operations note as a customer guide, preserving steps and constraints.	Step 1: verify the order identifier. Step 2: check payment status…
Short brief, long draft	Expand the launch brief into a handbook covering preparation, execution and recovery.	Preparation: confirm owners, rollback version and alert rules…
Long–document summary	Summarize the long operational record in five bullets with dates and unresolved issues.	1. Access restoration remains open. 2. Order payment was verified…
Long–document analysis	Analyze operational records for trends, evidence, alternative explanations and next steps.	Access incidents cluster after payment; distinguish permission synchronization from client state…
Structured extraction	Order INV–204: USD 120, paid. Return the four–field JSON object.	{"order_id":"INV–204","amount":120,"currency":"USD","status":"paid"}
Code editing	Patch the billing module to avoid processing an invoice twice after a timeout.	if invoice_id in processed_ids: return existing_result

Shape	Typical input excerpt	Illustrative output excerpt
Multi-turn assistant	History: payment succeeded, access reset failed. Follow-up: what should operations do next?	Verify order and payment status, check permission synchronization, and update the customer.

Cache ratios: controlled input versus observed reuse

The amendment crosses input lengths 8,192 and 32,768 with target shared-prefix fractions 0/25/50/75/90/100%, using six independent pairs per cell: 72 pairs and 144 calls. Each pair executes a warmup and then a measurement serially, both with a 256-token cap and the same generation parameters. Exact longest common prefixes are computed over the full tokenized chat template. A nominal zero may still share template tokens, so realized overlap is recorded.

Table 03-11 **Cache ratios: controlled input versus observed reuse**

Overlap target	Business context	Input variation and illustrative output
0%	Unrelated new tickets	Warm: ticket A. Measure: distinct record B. Output: summary of B.
25%	Small common rules, mostly dynamic data	Same initial rules + different order details. Output: new order actions.
50%	Common knowledge and user context of similar size	Same first half + a new second half. Output: new-context summary.
75%	Long shared knowledge, several customer inquiries	Same knowledge body; final ticket differs. Output: current-ticket advice.
90%	Large common material with short personalization	Retain most prefix tokens, replace trailing records. Output: updated summary.
100%	Repeat access to identical input	Repeat identical full input. Output: summary of the same material; generation still executes.

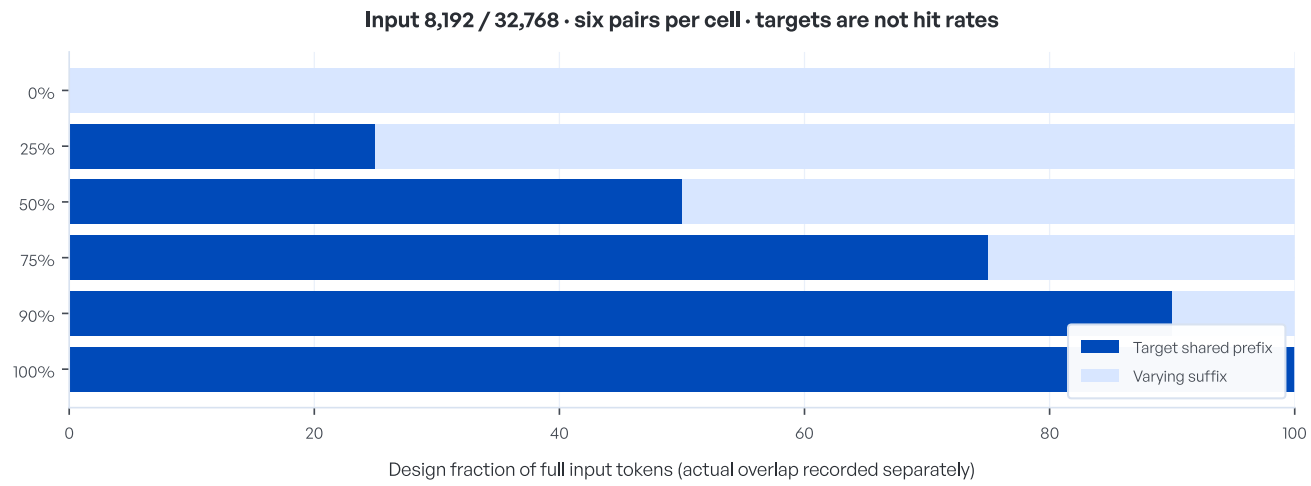


Figure 03-06 **Prefix-control design**

Dark blue is the target shared portion and light blue the varying portion; this describes input construction, not an achieved cache-hit rate.

Prefix overlap describes input similarity; actual cache hits require service fields and are not imputed when missing.

GPU and system controls

An exclusive window covers compute, memory, caches, scheduling and standalone operators; the second node repeats key conditions and both run collectives. Freeze versions, dtypes, layouts, memory state and power policy. The base matrix uses ten rounds per condition; extensions use five generally and ten for predeclared key conditions. Pair ordinary/graph launches, serial/overlapped work and CPU memory/compute interference. Retain correctness and profiling separately. Section 03 of the test plan gives full conditions, steady-state rules and scheduling.

04 · Metric definitions and statistics

Table 04-01 **Metric definitions and statistics**

Metric	Meaning and scope
Visible-content TTFT	Seconds from sending to the first visible content fragment; the user's initial wait.
First-generated latency	Send to the first reasoning, content or tool fragment; may precede visible content.
End-to-end latency	Seconds from sending to the end of the complete response stream.
Estimated time per output token	(Completion time − first-generated latency) / (generated tokens − 1); includes stream-finalization overhead, not exact inter-token timing.
Aggregate generation throughput	Generated tokens in eligible blocks / total block duration, in tokens/s.
Goodput	Complete requests meeting wait and generation-speed thresholds / measurement duration. Waiting includes client queuing from planned arrival; rejected arrivals remain in the denominator.
Length and cache	Record input, actual output, output cap, token-prefix overlap and reported cached tokens separately; a cap is not realized output.
Node resources	GPU activity, memory, temperature, power and service state aligned to the same load windows; estimated system power is not outlet-metered power.

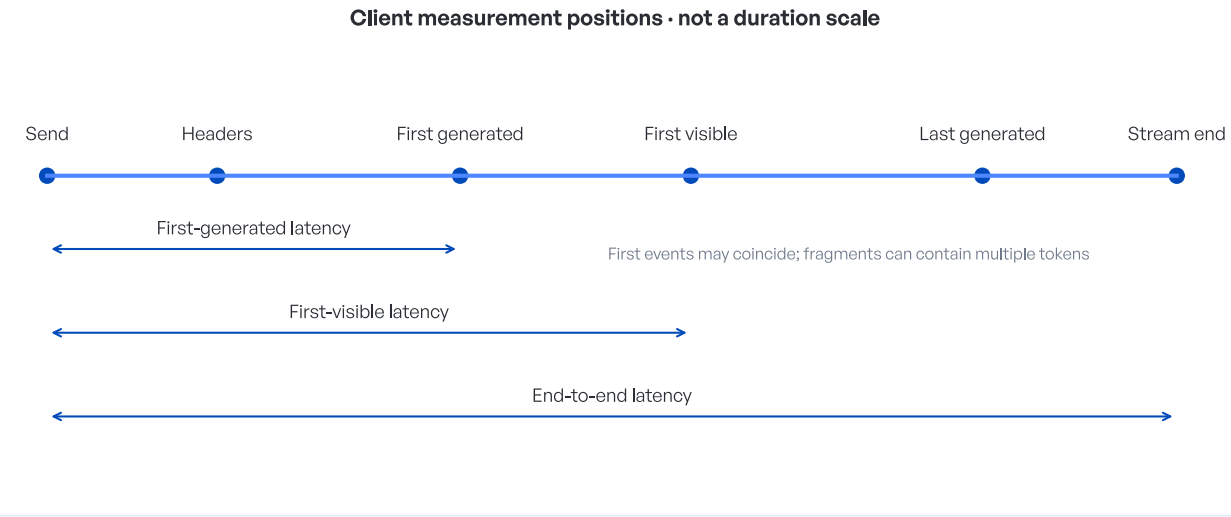


Figure 04-01 **Timing boundaries**

First generated and first visible content can differ; distances are not measured durations.

The median describes a typical response; roughly 95% of samples fall at or below p95. Throughput and latency comparisons use the same uninterrupted complete blocks, with warmups separate; failures remain in reliability statistics. Request and block counts are both reported. Repeated prompts do not increase the unique-prompt count. Business cells with 12 calls and cache cells with six pairs are diagnostic, not robust p99 estimates.

Token rates use verified upstream generation counts; costs use matched actual bills. Missing counts, cache fields and telemetry remain missing. Appendix A provides the complete metric dictionary and statistical details.

Streaming uses server-sent events (SSE). An event can contain several tokens, particularly with multi-token prediction: event spacing describes client delivery, not exact engine inter-token latency. The length extension additionally reports actual-output/cap attainment, full-input error, qualified counts and actual work per block.

Cache metrics: input structure, observed reuse and stability

Table 04-02 **Cache metrics: input structure, observed reuse and stability**

Cache metric	Definition and interpretation
Input-prefix overlap	Longest common token prefix over the full chat template / probe input length. A controlled input property.
Per-request token-cache fraction	Sum of cached input tokens / sum of input tokens for requests with valid fields, accompanied by field coverage / complete responses.
Fraction of requests with a hit	Requests with cached tokens >0 / requests with valid cache fields. A small template hit counts; it does not imply full-input reuse.
Window token-cache fraction	Endpoint cached-counter delta / prompt-counter delta; check freshness, resets, identity and receipt agreement. Count the head endpoint once.
Paired waiting change	Seed visible TTFT minus repeat TTFT; positive means the repeat was faster. Relative changes are separate. Network, order and load may contribute; this is not isolated causal acceleration.
Retention, competition and recovery	Compare cache fractions by actual idle gap; compare post-pressure with baseline for cache loss and extra waiting, then examine recovery. Findings apply only to tested pressure and intervals.
Temporal stability and coverage	Report valid / all windows, observed span, token-weighted mean, range and SD for at least three valid windows. Missing or reset intervals are neither zero-filled nor connected.

GPU and system measurement boundaries

Compute uses operations at the specified precision per device time; bandwidth separates useful read/write bytes from physical traffic. Hardware caches count accesses, while prefix caches count input tokens. Kernel, host-submit, collective and public-serving times are separate and cannot be subtracted across experiments. Sustained performance uses a same-workload warm baseline; efficiency uses integrated energy at a matching boundary. Appendix A defines formulas, units and numerical checks.

Use independent rounds or paired blocks as statistical units; report all samples, valid denominators, medians, interquartile ranges and 95% intervals. Key paired comparisons use round-cluster resampling with raw round points. Five-round screens are descriptive; ten rounds still do not guarantee precise intervals. Calculate ratios within paired rounds before aggregation. Randomize conditions within fixed stages and separate cold/warm-cache and first-touch tests. Microbenchmarks do not estimate serving tails; few rounds cannot support robust p99. A >5% decline lasting two ten-minute windows is an investigation flag, assessed with clocks, temperatures and work, not a hardware-failure verdict. Preserve timing and intrusive profiling runs separately.

05 · Realized coverage and analysis samples

Execution has ended, with explicit evidence gaps. Valid completion of fixed experiments, unexecuted conditions and incomplete observations are counted separately; no single overall completion percentage is calculated. A complete response does not establish answer correctness, target–length attainment or a production service objective.

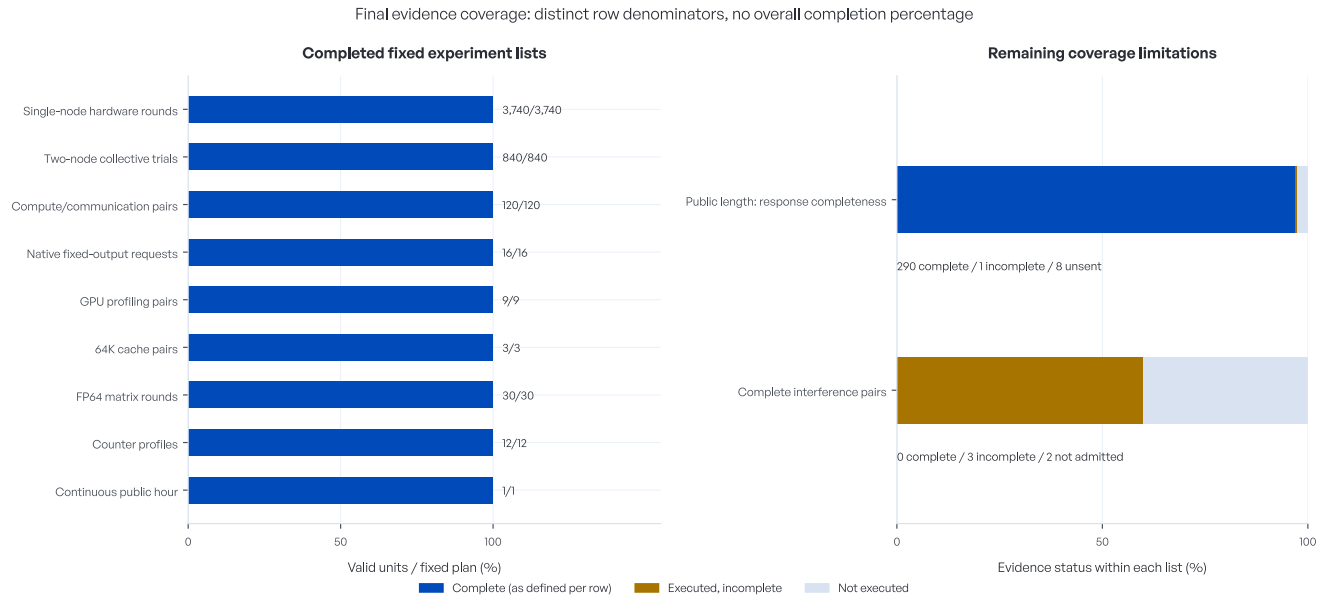


Figure 05–01 Final evidence coverage

Left: completed fixed lists. Right: public–length gaps and incomplete interference pairs. Each row retains its own sample unit and denominator; a complete public response does not establish target input/output length attainment.

Initial public batches: outcomes and count definitions

Table 05–01 Initial public batches: outcomes and count definitions

Initial public batch	Design count and meaning	Logical requests sent	Complete / expected boundary outcome	Incomplete outcomes
Interface and measurement calibration	12 (Fixed plan)	12	12	0
Natural workloads and concurrency scaling	720 (Fixed plan)	720	718	2
Input/output length matrix	1,440 (Fixed plan)	1440	1435	5
Paired reasoning budgets	360 (Fixed plan)	360	360	0
Long context and boundaries	37 (Fixed plan)	37	36	1
Prefix caching	90 (Fixed plan)	90	90	0
Business shapes and cache–ratio amendment	512 (Fixed plan)	512	507	5
Open arrivals and mixed workloads	2,000 (Conditional cap)	783	783	0

Initial public batch	Design count and meaning	Logical requests sent	Complete / expected boundary outcome	Incomplete outcomes
Length expansion and sustained operation	3,000 (Conditional cap)	929	925	4

Logical requests are deduplicated by planned sample; retry attempts are retained separately. Original batch outcomes are preserved rather than backfilled with later successes. The 2,000–call open–arrival and 3,000–call length/soak counts are conditional ceilings, not mandatory sample totals. The separate public–length list fixes 299 calls, with unsent/incomplete records below. The independent additional one–hour run has ended.

Hardware work is excluded from API–request denominators. Section 06 lists the batch scope, phase status and constrained conditions; execution manifests and status records remain separate.

Independent engine experiment: 78 archived analyzable blocks and 231 formal requests, plus two pilot/profile requests and four GPU traces across both ranks. These direct–origin observations do not enter public–API request or cost denominators; definitions and results are in section 08.

Coverage of the independent completion experiments

The additional experiments started on 17 October 2026. Status: Execution ended; delivered with actual coverage. Each batch retains its own requests, dates and denominators. New measurements do not rewrite earlier failures or success rates. Financial expenditure has no preset ceiling; actual charges are still reconciled per request.

Table 05–02 Coverage of the independent completion experiments

Additional scope	Evidence acquired	Definition
Fixed long output	16 / 16	Native engine; includes one calibration, then output–axis and balanced–length requests
Profiling pairs	9 / 9	Three input/concurrency conditions, three profiled/unprofiled pairs each
64K cache	6 / 6	First and immediate reuse for three independent prefixes
Complete public responses	1028	Independent reference blocks, affected conditions, mixed–length interference and a one–hour run
FP64 GEMM	30 / 30	Two nodes, three sizes, five rounds per cell
Representative counters	12	Compute/memory profiles returning valid counters

Unexecuted items and limits on conclusions

Table 05–03 Unexecuted items and limits on conclusions

Scope	Final evidence status	Interpretation limit
Public length extension	299 planned; 291 sent, 290 complete; 8 unsent and 1 sent without a complete response.	The 16 native fixed–output requests use a separate path and do not fill missing public combinations; natural completion is not fixed–length attainment.
Mixed–length interference	Of 5 candidates, 3 ran and 2 failed arrival–window admission; 0 fully observed pairs.	Executed groups include local short–request rejection. Retain rejection/waiting observations; no fully paired interference multiplier is supported.
Continuous public service	3600.1 continuous seconds; 600 arrivals, 599 complete and 1 locally rejected.	The one–hour window is complete; arrival success is below 100% and does not establish all–day availability.

Scope	Final evidence status	Interpretation limit
Matched H200 measurements	Not run: no H200 device was available.	Official specifications only; no measured inference–speedup ratio.
Wall–outlet energy	Not measured: no wall–meter data.	Sensor power and cost scenarios do not establish measured wall–electricity cost.
Physical DRAM traffic and per–token timing	Available counters/interfaces do not provide the required complete observations.	Report logical–byte bandwidth and event spacing without substituting them for physical traffic or exact per–token latency.

06 · GPU compute, memory and operators

Resource conditions and comparison. The resident condition retains model weights and serving processes. The unloaded condition stops both inference processes in a dedicated window, retaining the driver, clock policy and hardware. Each is warmed independently and runs the same base matrix with separate valid–round counts. Time and memory state change together; differences describe the overall resource condition, not an isolated causal cache, clock or chip effect.

Table 06–01 GPU compute, memory and operators

Node	Resource condition	Planned rounds	Archived valid timings	Memory–guard skips
0	Model resident	1070	684	386
0	Model unloaded	1070	1070	0
1	Model resident	1070	990	80
1	Model unloaded	1070	1070	0

Service restoration: verified; both nodes retain their original arguments, environment, working directory and core engine configuration.

This segment adds independent time–round evidence: 107 base conditions per node with ten rounds each, plus scalar FP32 FMA, random pointer chasing, shared and managed memory, CPU contention, multiple streams and supported low–precision/2:4 sparse GEMM. Results are checked before and after timing; guard skips and unsupported conditions are excluded from performance denominators. The original 49 valid records remain historical observations outside these repeats.

Table 06–02 GPU compute, memory and operators

Node	Suite	Planned rounds	Archived valid	Guard skips	Unsupported	Failed	Pending
0	Base matrix	1070	1070	0	0	0	0
0	Extensions	800	800	0	0	0	0
1	Base matrix	1070	1070	0	0	0	0
1	Extensions	800	800	0	0	0	0

8 records with invalid API calls or validation boundaries are audited separately and excluded from performance samples and valid coverage. Raw records and exclusion evidence remain in the analysis summary.

Definitions. TFLOP/s is trillions of floating–point operations per second; GEMM counts 2MNK. Sparse 2:4 data retain nonzero work separately; plots use dense–equivalent work. Logical GB/s counts requested reads/writes, not physical DRAM traffic. CUDA–event intervals include host–submission gaps, also affecting pointer–access nanoseconds; these are not intrinsic single–load latency. Summaries use condition–round medians and observed ranges; inner iterations and two devices are not a large independent device sample.

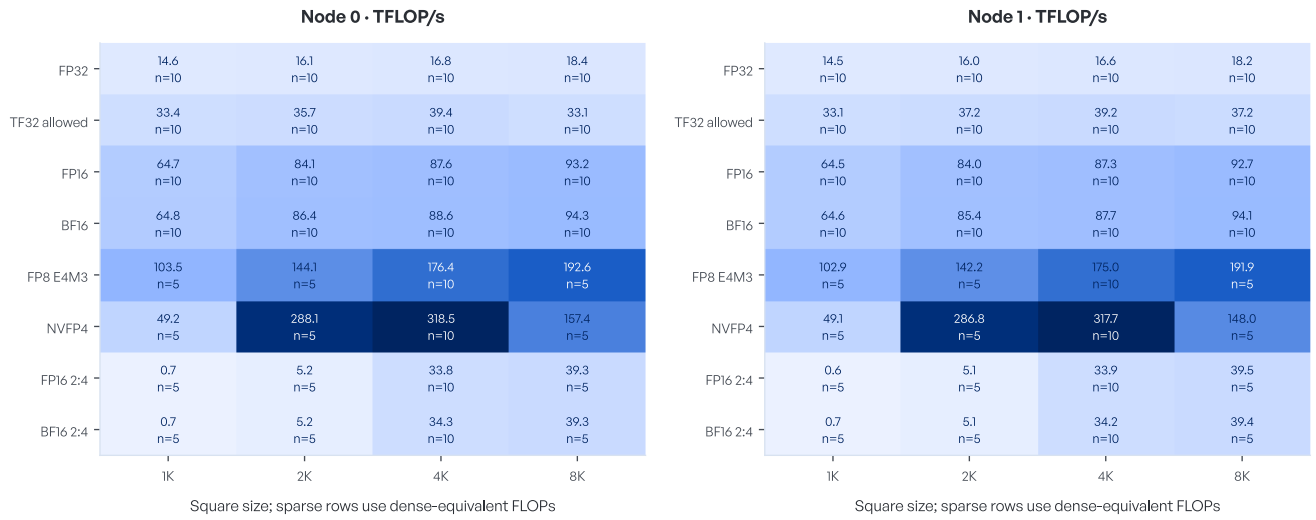


Figure 06–01 **Precision and size**

Cells show measured median TFLOP/s; blanks indicate no valid archived data. FP8, NVFP4 and sparse rows have explicit operand distributions and library paths; cross–row differences are not precision–only causal effects.

Precision–by–node statistics for key shapes are in Appendix B; the figure retains the size trend.

Independent operators and call boundaries

Reduction, Softmax and RMSNorm use the installed PyTorch library; RMSNorm includes library output allocation. Attention forces Flash SDPA with BF16,16heads,128dimensions per head and no causal mask, comparing full–length and single queries. These fixed synthetic tensors are not complete model prefill or decode. The 1M–element add compares ordinary submission with graph replay at equal work.

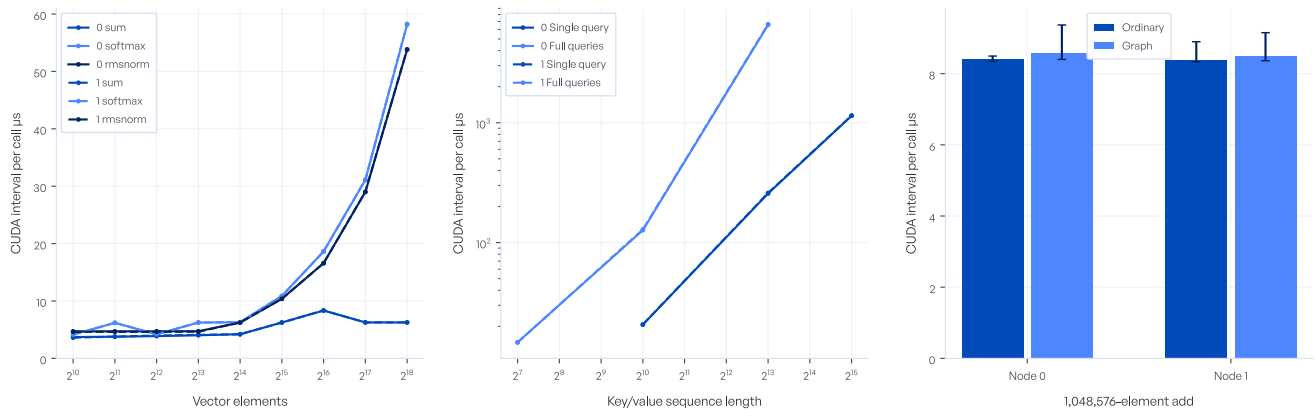


Figure 06–02 **Operator scale and launch mode**

Condition–median CUDA intervals; add–operation error bars show observed ranges. Attention uses a logarithmic latency axis.

Scalar compute, data supply and resource contention

Scalar FP32 uses explicit FMA instructions to compare one dependent chain with eight independent chains, holding per–thread FMA work fixed while varying blocks and parallelism; assembly is archived. Pointer chasing creates a complete random cycle with 128–byte spacing and ca/cg cache hints, where each load depends on the previous result. Shared memory varies address stride within a warp. CPU controls use 0/1/4/10/20 pinned worker processes with memory–streaming and compute controls. Each serial/parallel stream pair preserves the same matrix count.

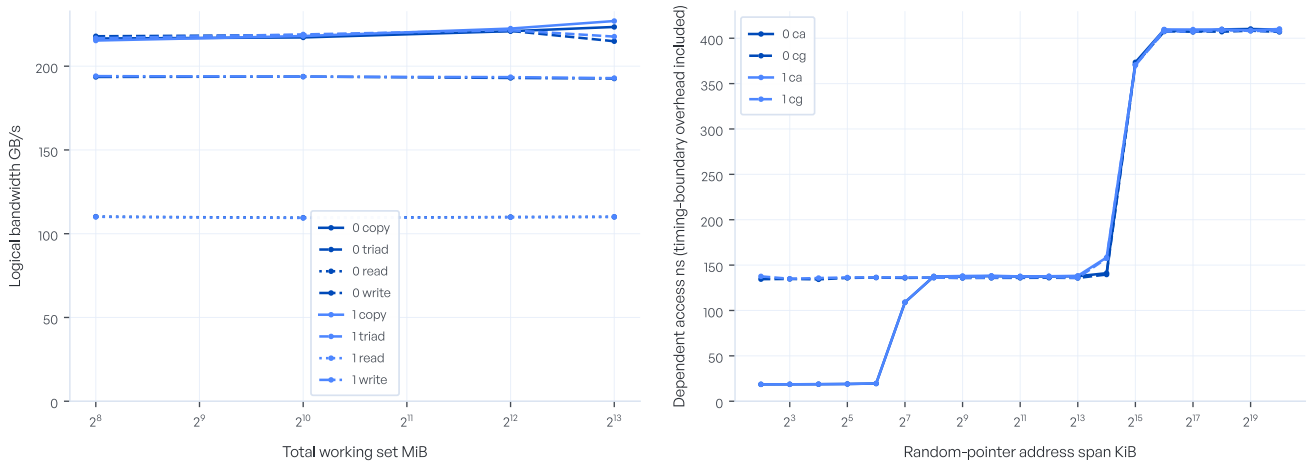


Figure 06–03 Data–supply curves

Logical read/write bandwidth on the left and dependent random–access time on the right. Address span is not physical DRAM traffic; a curve transition alone does not establish cache capacity or hit rate.

Table 06–03 Scalar compute, data supply and resource contention

Experiment	Node–condition groups	Valid rounds	Interpretation
Scalar FP32 FMA	36	180	Per–condition CSV gives workload–specific units and resource boundaries
Shared–memory strides	12	60	Per–condition CSV gives workload–specific units and resource boundaries
Empty–kernel launch	4	40	Per–condition CSV gives workload–specific units and resource boundaries
Multiple streams	16	80	Per–condition CSV gives workload–specific units and resource boundaries
CPU contention	18	90	Per–condition CSV gives workload–specific units and resource boundaries
Managed first/repeat access	2	10	Per–condition CSV gives workload–specific units and resource boundaries

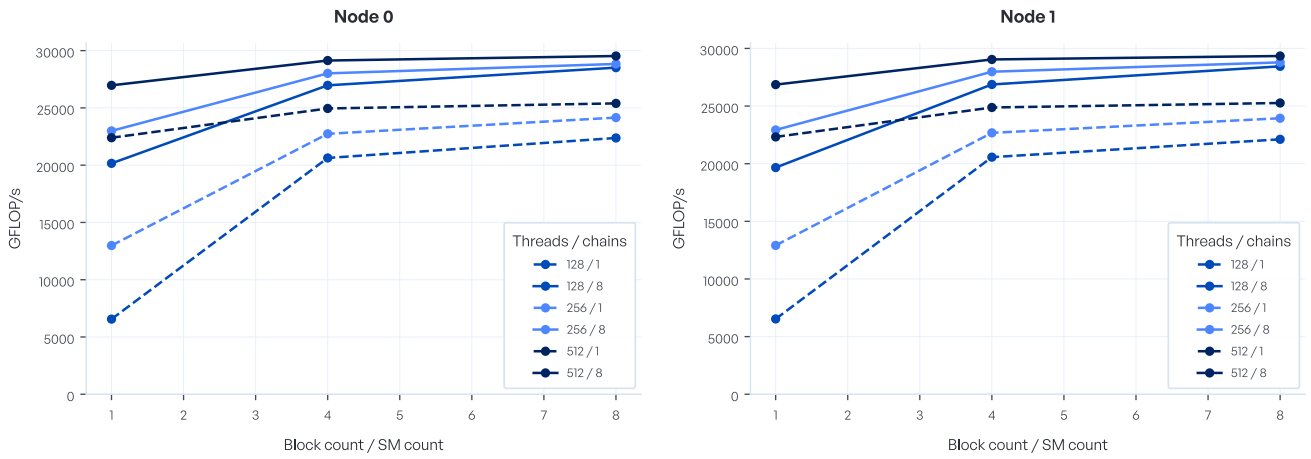


Figure 06–04 Scalar FP32 parallelism

Fixed per–thread work with varied block size, block count and dependency chains; median rates are not Tensor Core peaks.

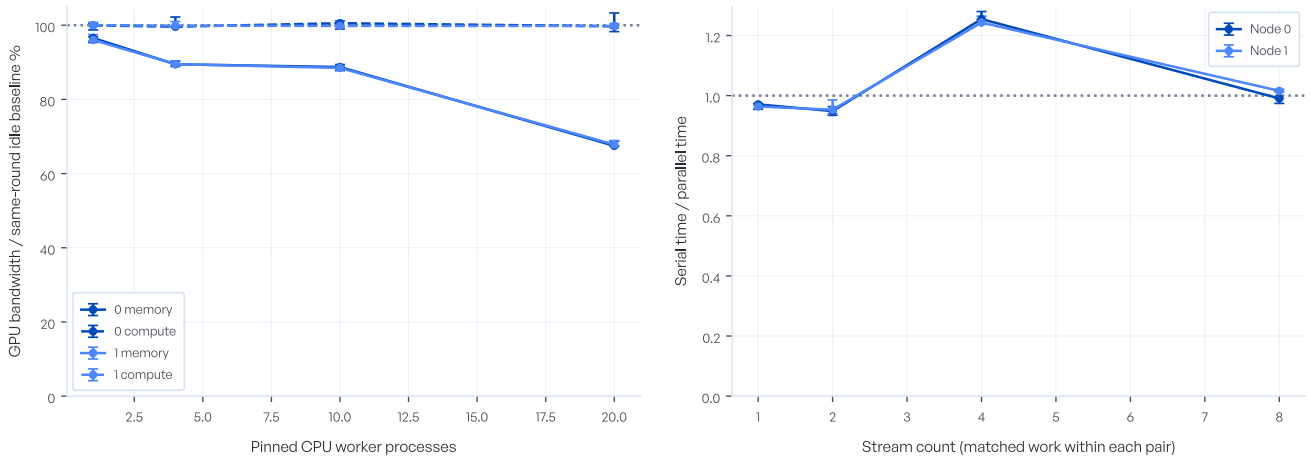


Figure 06–05 Contention and streams

Left: same–node, same–round ratios to the zero–contention baseline. Right: matched–work serial/parallel time. Points are paired–ratio medians; bars are observed ranges.

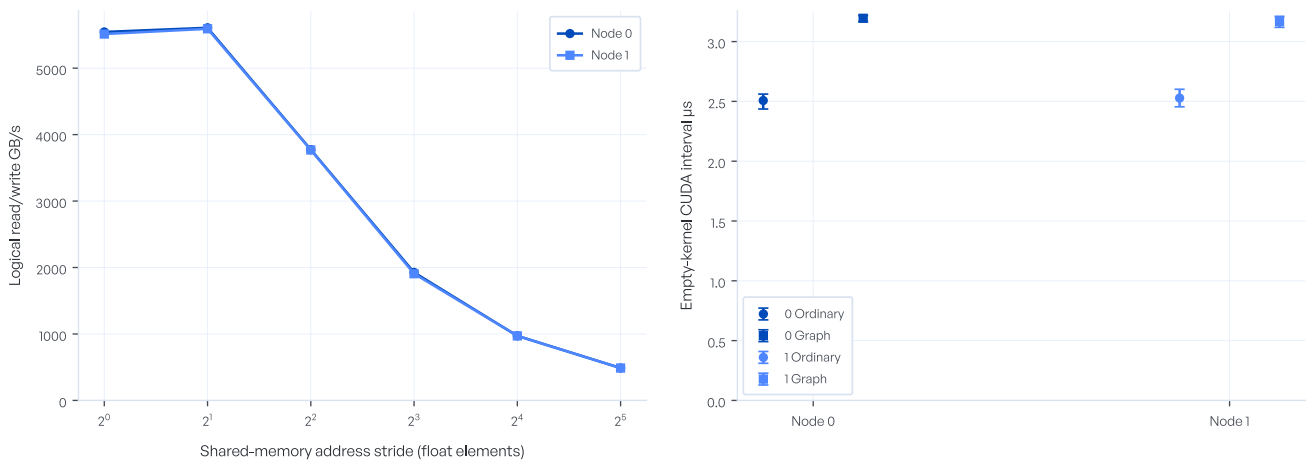


Figure 06–06 Shared memory and launch

Strides change address mapping; bank–conflict counts require hardware counters. Launch intervals include repeated host–submission gaps, not isolated instruction latency.

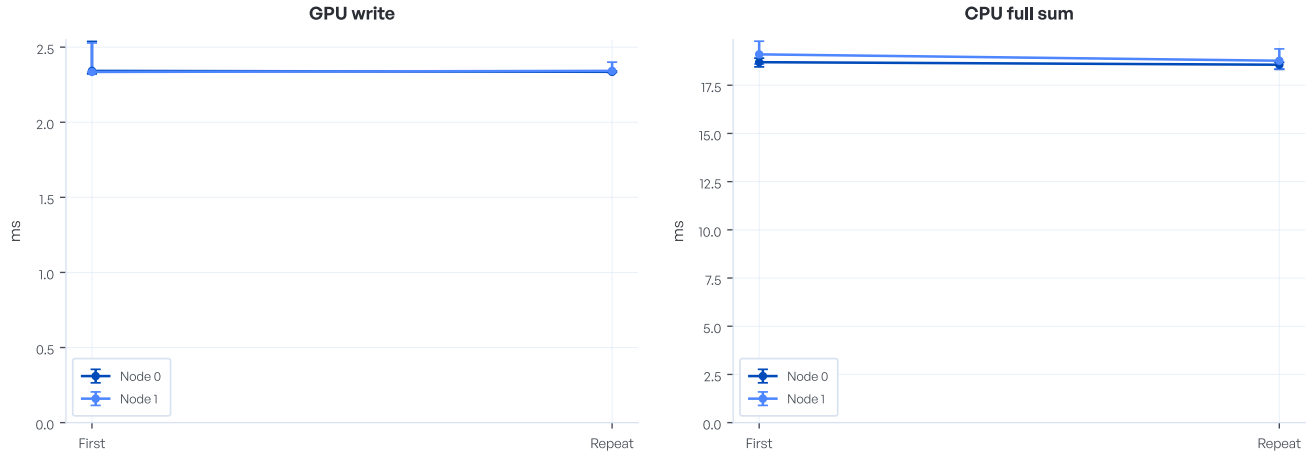


Figure 06–07 Managed–memory access order

CPU initialization, two GPU writes, then two full CPU sums within the same allocation per round. CPU sums include reduction work and are not pure migration latency.

Table 06–04 Scalar compute, data supply and resource contention

Node	Observation	Result	Conditions
0	Highest scalar–FP32 condition median	29529.7 GFLOP/s	512 threads, 8×SM blocks, 8 chains; 5 rounds
0	CPU–contention retention range	67.49–100.55%	Condition–median ratios across 8 conditions; see figure for ranges
0	Matched–work multi–stream ratio range	0.95–1.25×	Condition–median ratios across 4 conditions; see figure for ranges
1	Highest scalar–FP32 condition median	29341.2 GFLOP/s	512 threads, 8×SM blocks, 8 chains; 5 rounds
1	CPU–contention retention range	67.82–100.00%	Condition–median ratios across 8 conditions; see figure for ranges

Node	Observation	Result	Conditions
1	Matched-work multi-stream ratio range	0.95—1.24×	Condition-median ratios across 4 conditions; see figure for ranges

NVFP4 uses exactly representable E2M1 operands with explicit block maxima and unity scales, validating packed values before GEMM. This isolates arithmetic error and does not represent production weight quantization or MoE kernel performance. Quantization/packing and sparse compression times are saved separately, as are diagnostic traces and unprofiled timing. Counter access is diagnostic without system permission changes; unavailable cache-hit and active-warp counters remain missing.

Table 06-05 Scalar compute, data supply and resource contention

Node	Hardware counter status	Observability scope
0	Profiler allocation limited	The standalone probe succeeds; its first CUDA allocation fails under profiling
1	Probe fields partially readable	One fixed FMA thread block only; representative matrix/model counters remain unmeasured

Table 06-06 Scalar compute, data supply and resource contention

Diagnostic counter	Node 0	Node 1
Active warps relative to sustained peak	Unavailable	16.56 %
L2 sector hit fraction	Unavailable	83.91 %
DRAM bytes read	Unavailable	Unavailable
DRAM bytes written	Unavailable	Unavailable

Available values come from one successful profile of a fixed single-thread-block FP32 FMA probe after service restoration and idle admission, solely to verify counter access. Nsight Compute runs administratively without changing driver permissions; clock/cache controls are none. DRAM-byte fields return n/a and remain missing, never zero. This probe does not represent large-matrix or model-average cache hits, physical bandwidth utilization or compute efficiency; shared units such as L2 may also reflect other GPU activity. It cannot establish a model roofline. NVIDIA permission guidance · Nsight Compute metric guidance

Per-record data, condition summaries and guard states · Valid timing CSV. Memory-guard skips with resident models do not establish unsupported hardware. Resident and unloaded results are separate conditions; rounds are not pooled and differences are not attributed to one hardware mechanism.

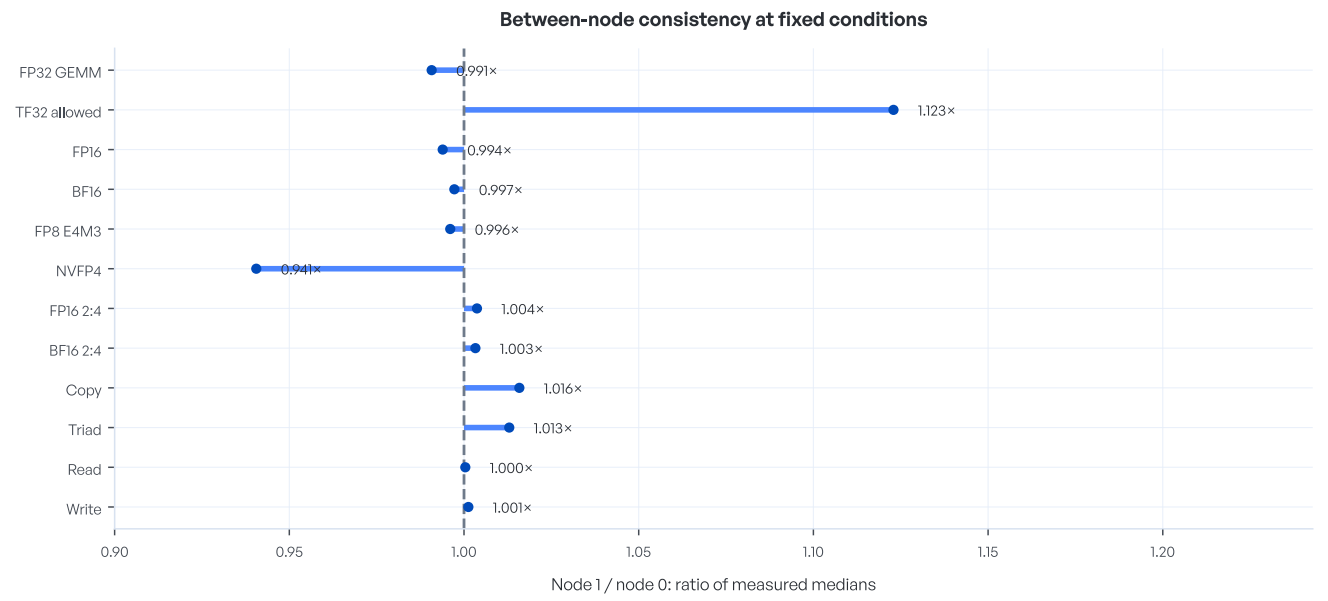


Figure 06–08 Node consistency

Node 1 median / node 0 median for fixed 8192^3 matrices and 8 GiB access. Ratios describe these two tested devices, not population variability.

Additional FP64 matrix measurements

After unloading the model, CUDA double-precision GEMM counts a multiply-add as two operations. Inputs are exactly representable signed binary fractions. Nine entries spanning rows/columns are checked against a CPU double-precision reference every round; only validated rounds enter the table.

Table 06–07 Additional FP64 matrix measurements

Node	M=N=K	n	TFLOP/s median [range]
0	1024	5	**0.4130** [0.4107–0.4157]
0	4096	5	**0.4254** [0.4246–0.4254]
0	8192	5	**0.4280** [0.4280–0.4281]
1	1024	5	**0.4086** [0.4048–0.4097]
1	4096	5	**0.4246** [0.4246–0.4250]
1	8192	5	**0.4276** [0.4276–0.4277]

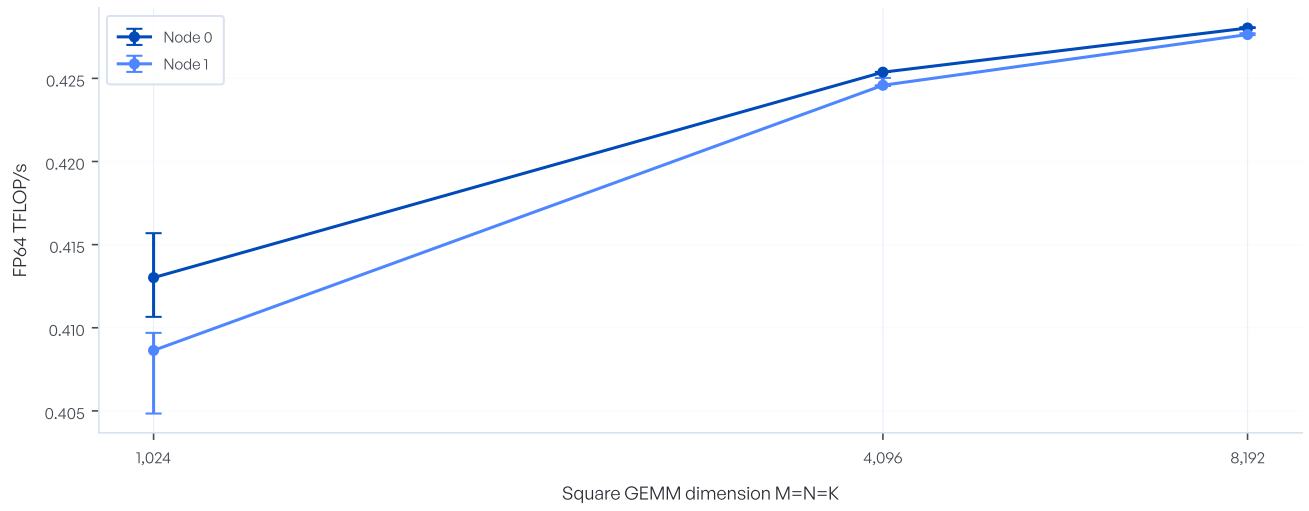


Figure 06–09 FP64 size effects

Points are five–round medians, with observed ranges; these do not predict low–precision model throughput.

Hardware counters for representative kernels

Table 06–08 Hardware counters for representative kernels

Node / kernel	Metric	n	Median [range]	Unit
0 / fma	Profiled kernel duration	3	936896.000 [936672.000–937504.000]	ns
0 / fma	L2 sector hit percentage	3	15.690 [7.020–17.090]	%
0 / fma	SM throughput percentage of peak	3	99.030 [99.020–99.040]	%
0 / fma	Active warp percentage	3	46.340 [46.340–46.350]	%
0 / triad	Profiled kernel duration	3	3831552.000 [3800544.000–3865088.000]	ns
0 / triad	L2 sector hit percentage	3	0.010 [0.010–0.020]	%
0 / triad	SM throughput percentage of peak	3	2.740 [2.690–2.740]	%
0 / triad	Active warp percentage	3	88.920 [88.470–92.660]	%
1 / fma	Profiled kernel duration	3	938752.000 [937888.000–939232.000]	ns
1 / fma	L2 sector hit percentage	3	18.230 [18.000–45.320]	%
1 / fma	SM throughput percentage of peak	3	99.030 [99.020–99.040]	%
1 / fma	Active warp percentage	3	46.340 [46.340–46.350]	%
1 / triad	Profiled kernel duration	3	3999008.000 [3891552.000–4325664.000]	ns
1 / triad	L2 sector hit percentage	3	0.020 [0.010–0.070]	%
1 / triad	SM throughput percentage of peak	3	2.590 [2.390–2.650]	%
1 / triad	Active warp percentage	3	87.300 [84.800–90.130]	%

The compute kernel uses eight independent FMA chains per thread and four blocks per SM; the memory kernel uses a **768 MiB** three–array working set. Administrative Nsight Compute captures one selected kernel without changing global driver permissions, clocks or cache control. Counters describe these kernels, not average model cache behavior. The current interface advertises no physical DRAM–byte metric, so a model Roofline requiring that count remains unsupported.

07 · Two-node communication and cooperation

Device data exchange: port, RDMA and collectives

One GPU communication process (rank) runs per node over ConnectX-7/RoCE. The active physical port specifies 200 Gb/s=25 GB/s in one direction; two logical interfaces share it. RDMA measurements use host buffers; NCCL measurements use GPU collectives, each with its own work definition. Functionally aligned H200 interconnect references are consolidated in section 10.

Table 07-01 Device data exchange: port, RDMA and collectives

Transfer / operation	Message size	Rounds	Median rate	Measurement conditions
RDMA Write / roceP2p1s0f1	8 MiB	3	13.629 GB/s (109.03 Gb/s)	2026-10-14; host-memory buffers, one logical path
RDMA Write / rocep1s0f1	8 MiB	3	13.617 GB/s (108.94 Gb/s)	2026-10-14; host-memory buffers, one logical path
all_gather	1 GiB	10	19.930 GB/s	2026-10-16; two GPU ranks, bus metric
all_reduce	1 GiB	10	21.815 GB/s	2026-10-16; two GPU ranks, bus metric
reduce_scatter	1 GiB	10	20.946 GB/s	2026-10-16; two GPU ranks, bus metric

RDMA Write uses ib_write_bw, one QP and one-way 8 MiB messages, with three repeats per logical path. Its --report_gbits output is divided by eight for GB/s. This separate historical host-buffer test neither substitutes for nor adds to GPU collectives. Current NCCL logs confirm RoCE through NET/IB and explicitly report GPU Direct RDMA Disabled; a GPU-direct RDMA path is not claimed.

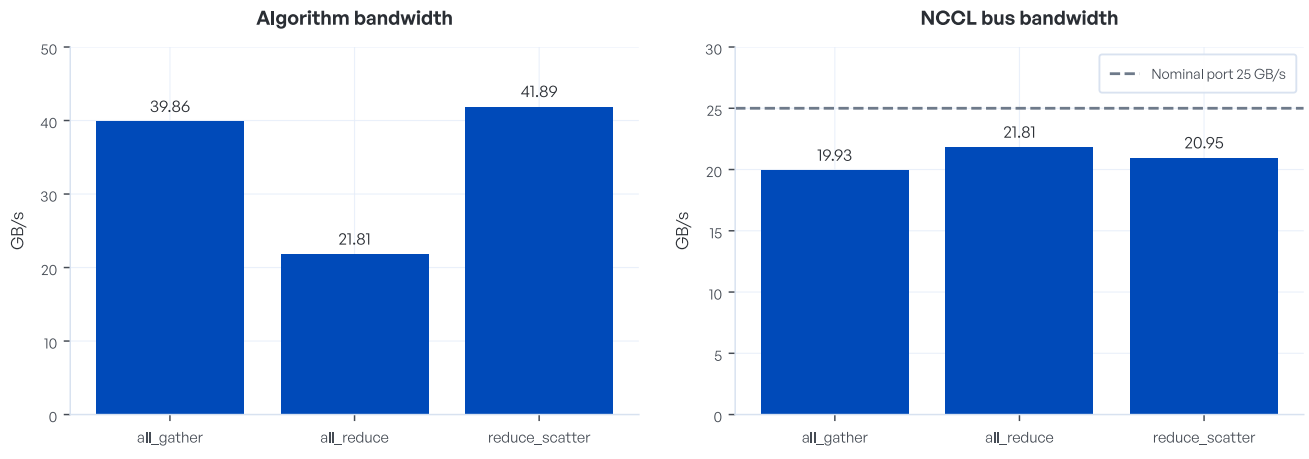


Figure 07-01 1 GiB collectives

Algorithm bandwidth counts the defined full input/output tensor. NCCL bus factors for two ranks are 1 for AllReduce and 1/2 for AllGather/ReduceScatter. Algorithm GB/s may exceed 25 without exceeding a 200 Gb/s physical link. Bus bandwidth is not a NIC hardware-counter reading; the dashed line is only a nominal line-rate reference.

First RDMA path · Second RDMA path · Fixed conditions and provenance · Official–source snapshot index.This segment has 840 archived valid two–rank paired trials. Each rank pair is one observation using the slower rank. Algorithm bandwidth is defined message bytes divided by time; bus bandwidth uses NCCL's correction, never summing ranks. Message size is full input for AllReduce/ReduceScatter and full output for AllGather. The original 434 trials are retained separately.

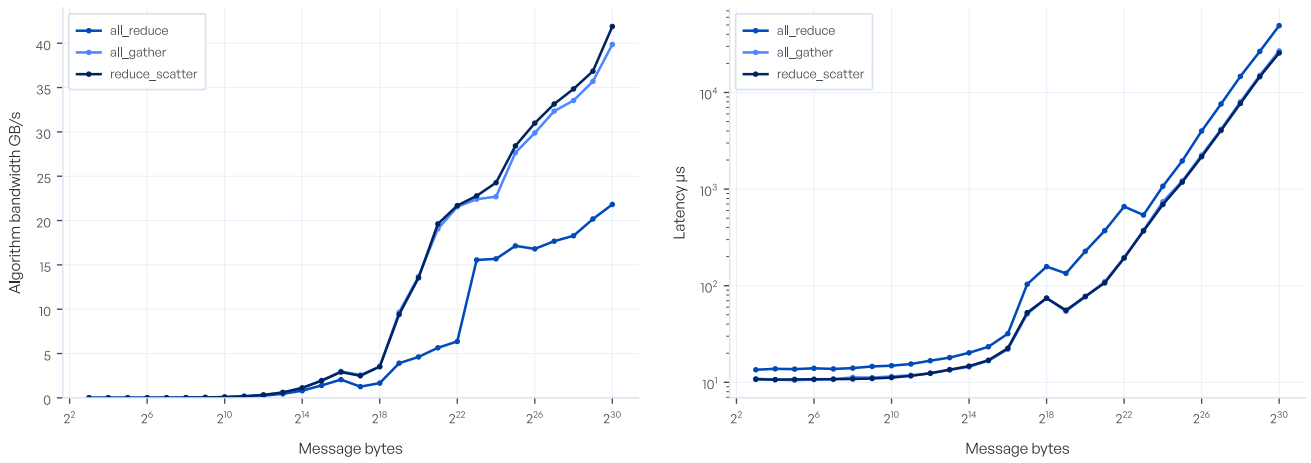


Figure 07–02 Collective size curves

Algorithm bandwidth and latency are shown separately for three operations; both nodes must pass memory admission.

Table 07–02 Device data exchange: port, RDMA and collectives

Operation	Message bytes	Paired rounds	Median latency µs	Algorithm GB/s	Bus GB/s
all_gather	8	10	10.89	0.001	0.000
all_gather	8,192	10	13.46	0.608	0.304
all_gather	8,388,608	10	374.38	22.408	11.204
all_gather	67,108,864	10	2245.77	29.882	14.941
all_gather	1,073,741,824	10	26937.75	39.860	19.930
all_reduce	8	10	13.49	0.001	0.001
all_reduce	8,192	10	18.02	0.454	0.454
all_reduce	8,388,608	10	539.32	15.555	15.555
all_reduce	67,108,864	10	3993.42	16.805	16.805
all_reduce	1,073,741,824	10	49220.55	21.815	21.815
reduce_scatter	8	10	10.74	0.001	0.000
reduce_scatter	8,192	10	13.51	0.606	0.303
reduce_scatter	8,388,608	10	368.19	22.784	11.392
reduce_scatter	67,108,864	10	2165.91	30.984	15.492
reduce_scatter	1,073,741,824	10	25631.79	41.891	20.946

Controlled compute/communication overlap

Hold one BF16 GEMM and one AllReduce fixed, measuring compute-only, communication-only, serial combination and two-stream overlap. Work is matched for each message/matrix pair. Serial/overlapped time above one means this synthetic combination finishes sooner; below one means overlap is slower. It is not a two-node model speedup.

Table 07-03 Controlled compute/communication overlap

Message bytes	Square size	Paired rounds	Serial / overlap	Observed range
8,192	2048	5	1.059	1.001–1.113
8,192	4096	5	1.020	1.011–1.029
8,388,608	2048	5	1.168	1.093–1.193
8,388,608	4096	5	1.035	1.020–1.094
67,108,864	2048	5	1.023	1.016–1.027
67,108,864	4096	5	1.096	1.071–1.122

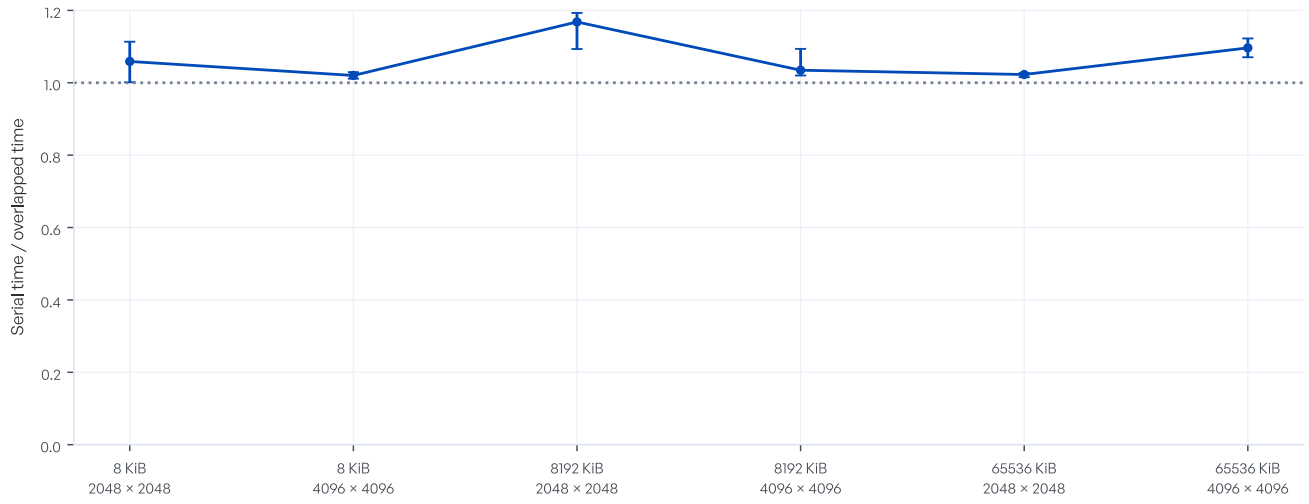


Figure 07-03 Overlap benefit

Median within-round four-arm ratios with observed ranges; above one indicates shorter time for this workload.

Current statistics and sources · NCCL bandwidth definitions. Transport is verified from this segment's NCCL logs; interface names alone do not establish a GPU Direct path.

08 · From engine stages to serving performance

This section first defines and measures engine stages, then examines public-endpoint experience and input/cache effects. Prefill processes input; decode generates output. Workloads and timing boundaries differ, so results inform interpretation without subtracting times across experiments.

Independent engine prefill and decode experiment

This experiment calls the native SGLang endpoint on the inference host, excluding the public network, Cloudflare and the Infron gateway from its timing path. Weights and two-node tensor parallelism are unchanged, with MTP enabled. Concurrency is capped at eight; direct-origin calls are outside Infron API billing. These observations are not pooled with public-endpoint requests.

Table 08-01 Independent engine prefill and decode experiment

Experiment	Work and repeats	Measurement scope
Prefill	1K/8K/32K/64K/128K/250,000 inputs, one output; C=1 with C=4/8 at shorter inputs; three rounds per condition	Scheduler prefill start-to-finish; includes chunk scheduling gaps, not pure GPU time
Decode	1K/8K/32K inputs × C=1/2/4/8, plus 64K/128K at C=1; fixed 256 outputs; three rounds	Output rate from API-process first token to completion, MTP acceptance length and verify count
Prefix reuse control	8K/32K, three initial/immediate-repeat pairs each, one output	Validate actual cached tokens; does not measure retention
GPU timelines	8K input, 128 output, C=1; both ranks capture prefill and four decode scheduler steps	GPU activity intervals, NCCL-named kernels and MTP draft/verify stages; profiled data stay separate from normal rate estimates

The manifest contains 78 blocks and 231 formal requests, plus two pilot/profile requests. Combinations exceeding concurrency × (input + output) = 262,144 are excluded. Seeded synthetic order records are tokenized with the installed model tokenizer to exact token-ID lengths, bypassing chat templates. Temperature is zero and EOS is ignored to fix output length. This controlled engine workload does not represent natural-chat length or MTP acceptance. Conditions are shuffled within each round; requests use distinct early prefixes, with hits verified from returned fields.

Timing definitions and validation

Table 08-02 Timing definitions and validation

Metric	Definition and scope
Scheduler prefill duration	prefill_finished_time − forward_entry_time within the scheduler process; includes chunk gaps and excludes the initial admission queue.
Prefill processing rate	Σ(input − cached tokens) / (last prefill finish − first prefill start) in the same block; overlapping request durations are not summed.
Per-request decode rate	(actual output − 1) / (engine E2E − engine TTFT), cross-checked against native decode_throughput; includes scheduling, result delivery and MTP execution.
Concurrent decode-window throughput	Σ(actual output − 1) / (last finish − first token), timed in one API process; staggered prefill admission can remain within this window.
MTP acceptance length and rate	Installed acceptance length = completed output tokens / verify count; acceptance rate = correct drafts / proposed drafts. Their denominators differ; acceptance length is neither a speedup nor an actual per-captured-step output count.
GPU activity decomposition	Union kernel/copy intervals separately per rank, then split NCCL-only, other-only and overlap. Nested annotations and cross-rank durations are not added.

Scheduler and API-process absolute timestamp conversions have different offsets; intervals are computed only within one process clock domain. Three-block medians and observed min/max ranges describe repeatability, without tail-percentile or confidence-interval claims. Retracted requests require separate analysis and are excluded from this timing decomposition.

Independent rate measurements

The archive contains 78 analyzable blocks and 231 formal requests; 0 failed timing validation.

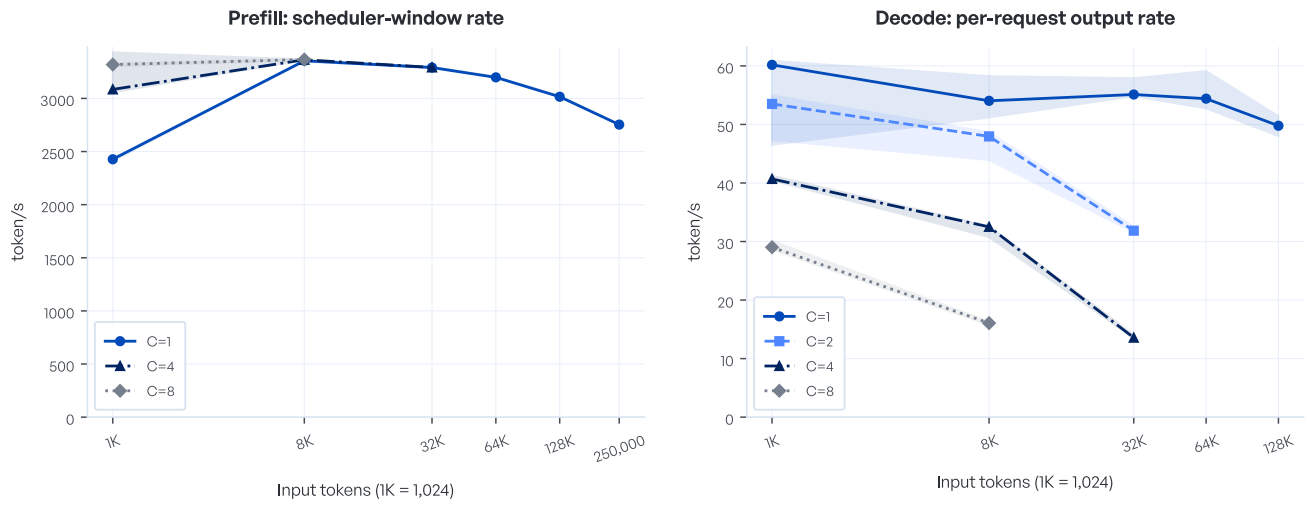


Figure 08-01 Engine-stage rates

Left: scheduler prefill-window rate. Right: per-request decode rate. Shading is the observed block range, not a confidence interval.

Table 08-03 Independent rate measurements

Input tokens	Concurrency	Blocks / requests	Median prefill s	Window rate tokens/s	Cached tokens
1,024	1	3 / 3	0.422	2,427.4	0
1,024	4	3 / 12	0.915	3,084.9	0
1,024	8	3 / 24	2.059	3,318.9	0
8,192	1	3 / 3	2.442	3,354.2	0
8,192	4	3 / 12	4.837	3,363.7	0
8,192	8	3 / 24	4.861	3,366.1	0
32,768	1	3 / 3	9.957	3,290.8	0
32,768	4	3 / 12	12.351	3,291.9	0
65,536	1	3 / 3	20.491	3,198.3	0
131,072	1	3 / 3	43.462	3,015.8	0
250,000	1	3 / 3	90.787	2,753.7	0

Table 08–04 Independent rate measurements

Input tokens	Concurrency	Blocks / requests	Per request tokens/s	Concurrent window tokens/s	Median MTP acceptance length
1,024	1	3 / 3	60.2	60.2	3.71
1,024	2	3 / 6	53.5	107.0	3.74
1,024	4	3 / 12	40.7	144.9	3.61
1,024	8	3 / 24	29.0	206.6	3.58
8,192	1	3 / 3	54.0	54.0	3.32
8,192	2	3 / 6	48.0	93.2	3.44
8,192	4	3 / 12	32.5	86.2	3.46
8,192	8	3 / 24	16.1	81.1	3.37
32,768	1	3 / 3	55.1	55.1	3.41
32,768	2	3 / 6	31.9	38.0	3.18
32,768	4	3 / 12	13.6	29.0	3.37
65,536	1	3 / 3	54.4	54.4	3.41
131,072	1	3 / 3	49.8	49.8	3.12

Across measured C=1 input lengths, median per–request decode rates range from 49.8 to **60.2 tokens**/s. This range applies to synthetic records and fixed output; it does not substitute for public–endpoint capacity or arbitrary business generation rates.

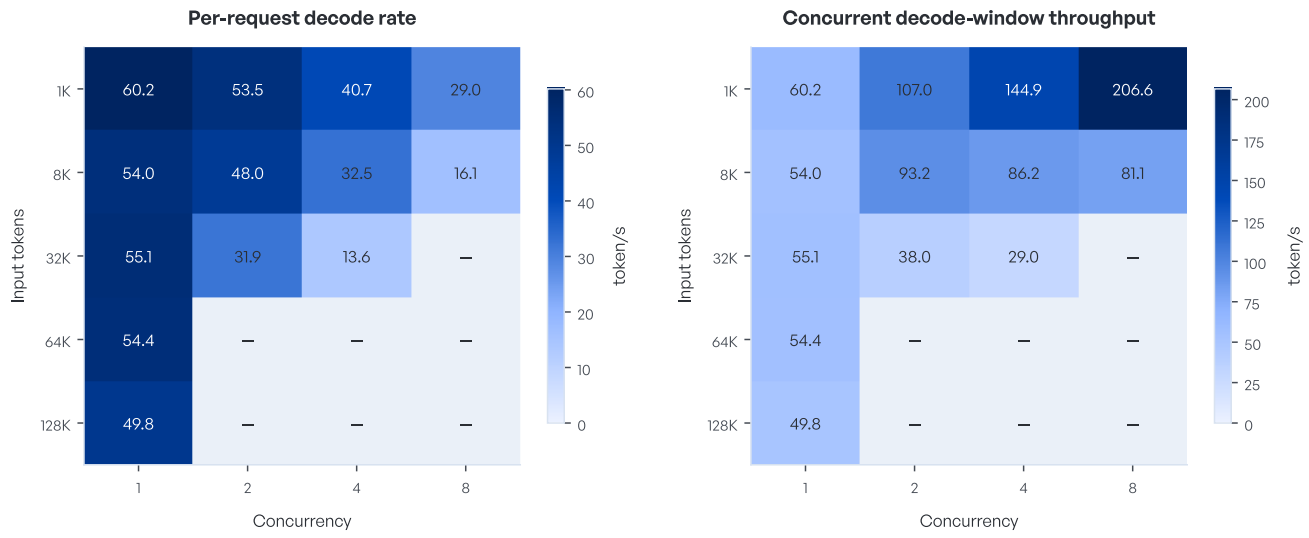


Figure 08–02 Decode length × concurrency

Each cell summarizes three blocks; gray cells are unmeasured and panels use independent scales. Output is fixed at 256 tokens; more concurrency need not increase aggregate throughput at longer inputs.

Immediate prefix–reuse prefill control

Table 08–05 Immediate prefix–reuse prefill control

Input	Pairs	Repeated–input cached–token fraction	Initial / repeated median prefill s	Median paired duration ratio
8,192	3	0.00%	2.4343 / 2.4507	0.99
32,768	3	75.00%	9.9719 / 2.5517	3.91

Cache fraction uses total actual repeated–input tokens as its denominator; duration ratios are calculated per pair before taking the median. Fixed initial–then–repeat order retains order effects. Three immediate repeats do not establish long–term cache stability, and their ratios are not model compute speedups.

Reported cached–token fractions were 0% for immediate 8K repeats and 75% for 32K repeats. Identical input does not guarantee full reuse. These observations apply to the tested lengths, one–token output and current cache state; they do not establish universal cache rules or retention.

GPU–stage and communication activity

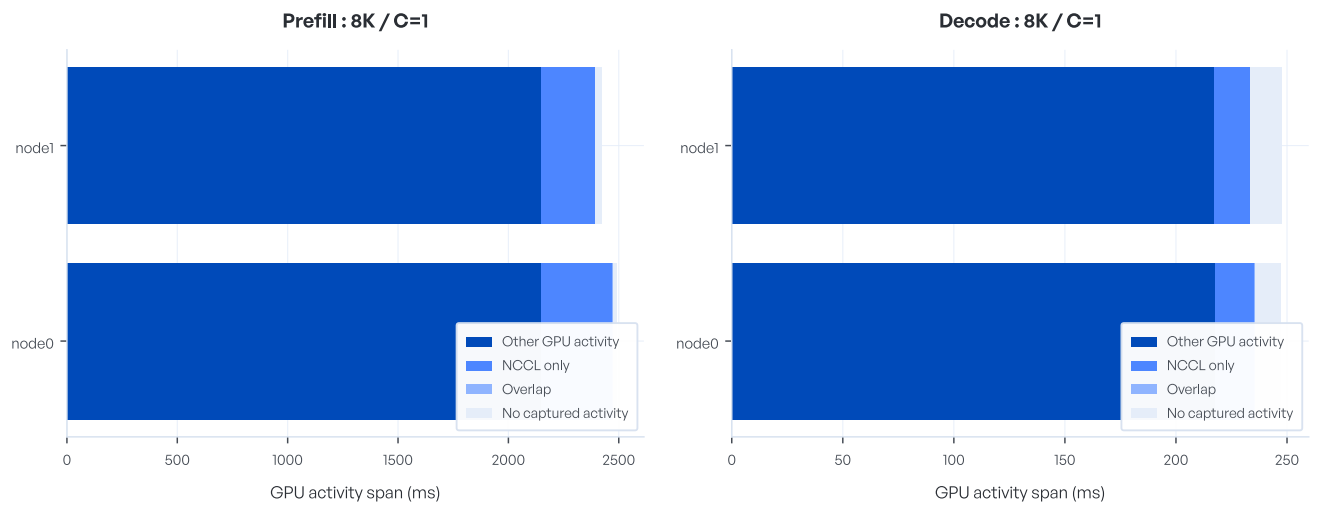


Figure 08–03 Per–rank GPU activity

Each bar covers its own captured GPU activity span. Profiling adds overhead and explains paths only; other activity may include communication without NCCL kernel names.

Table 08–06 GPU–stage and communication activity

Rank / stage	Activity span ms	Busy union ms	NCCL–only ms	Overlap ms
node0 / decode	247.07	235.38	17.68	0.05
node0 / prefill	2,489.51	2,472.72	325.91	0.02
node1 / decode	247.42	233.38	15.80	0.03
node1 / prefill	2,423.98	2,395.04	246.97	0.00

Table 08–07 GPU–stage and communication activity

Rank / capture stage	Internal phase	Matched executions	Mean GPU–correlated span ms
node0 / decode	MTP draft	4	10.575
node0 / decode	Target–model verify	4	44.100
node0 / decode	Draft–state extension	4	6.150
node0 / prefill	Target–model prefill	1	2,401.288
node0 / prefill	Draft–state extension	1	87.003
node1 / decode	MTP draft	4	10.692
node1 / decode	Target–model verify	4	44.079
node1 / decode	Draft–state extension	4	6.159
node1 / prefill	Target–model prefill	1	2,332.348
node1 / prefill	Draft–state extension	1	87.009

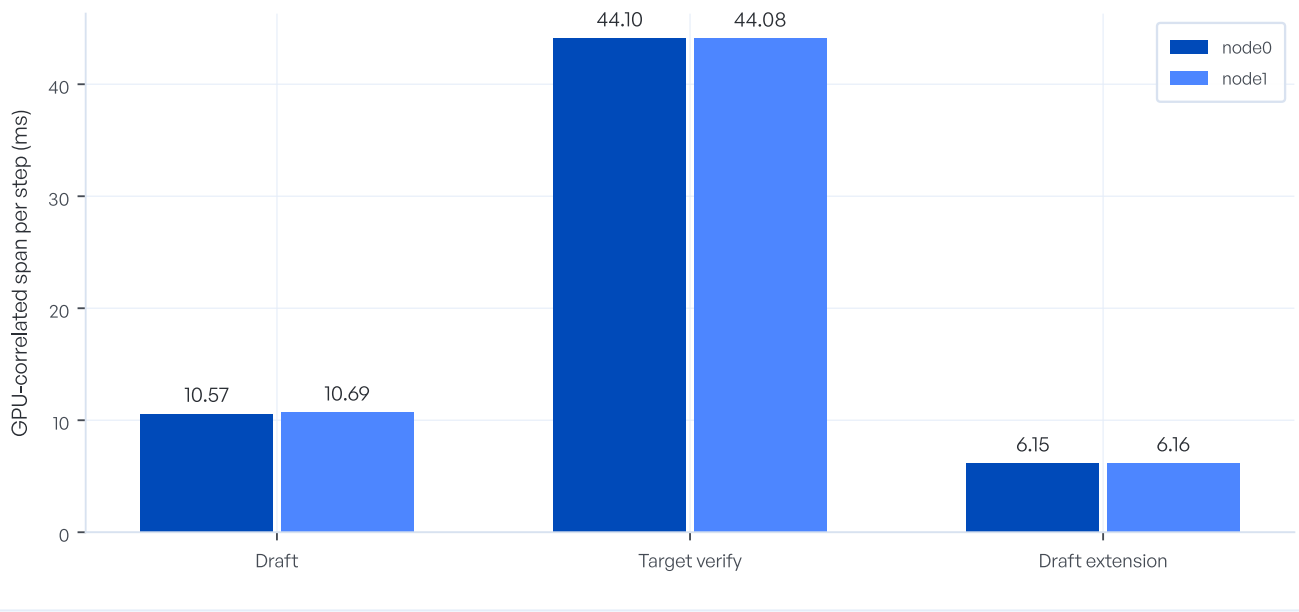


Figure 08–04 MTP phase decomposition

Four captured decode steps per rank; bars show mean GPU–correlated spans, which cannot be added as end–to–end latency.

Phase ranges match CPU and GPU annotations by External id. CUDA Graphs annotate one verify execution on multiple virtual streams, so intervals are merged per execution before calculating its span. Spans include device waits and communication and differ from busy–kernel time; draft, verify and extension spans cannot be added without overlap checks. Accepted–token counts for each captured step are unavailable, so device decode tokens/s is not fabricated by multiplying step count by a fixed MTP length.

Both ranks captured prefill/decode GPU kernels and MTP draft, verify and draft–extension annotations. GPU work and communication overlap, so summing operator or NCCL durations double–counts time. The initial trace covers 8K at concurrency 1; additional input/concurrency conditions and paired overhead are reported below. NCCL activity is not equivalent to removable communication wait.

Per–request metrics · Per–block metrics · Source and raw timelines

Implementation references: SGLang profiling · SGLang observability · Installed timing source; online documentation informs methodology, while saved installed source determines field semantics.

Public concurrency: throughput and waiting

Table 08–08 **Public concurrency: throughput and waiting**

Concurrency	Blocks / requests	Generated tokens/s	Visible TTFT median / p95 s	Completion p95 s
1	2 / 120	45.1	1.01 / 1.65	19.46
2	3 / 180	78.7	1.07 / 1.78	24.87
4	1 / 60	120.2	1.11 / 1.79	30.53
8	3 / 180	158.9	1.17 / 2.39	46.93

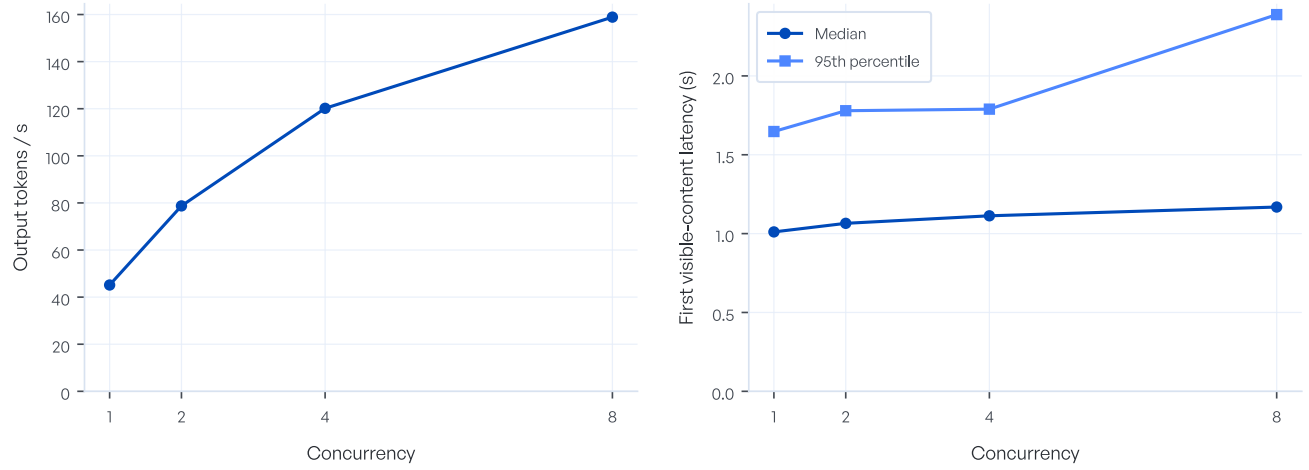


Figure 08–05 **Concurrency comparison**

Same–cohort throughput and visible–content latency; consider the waiting cost of additional concurrency.

For this short–prompt bank, concurrency eight gives **158.9 tokens/s** with mean output 400.8 tokens. Use workload–matched results for long–document capacity planning.

Length: input processing and generation

The heatmaps separate input length, output budget and concurrency. Each cell includes uninterrupted complete blocks only. Darker means a larger value; grey means no eligible block. The output budget is a ceiling, not a forced generation length. Both figures use the same population and show first–content waiting and aggregate generation respectively.

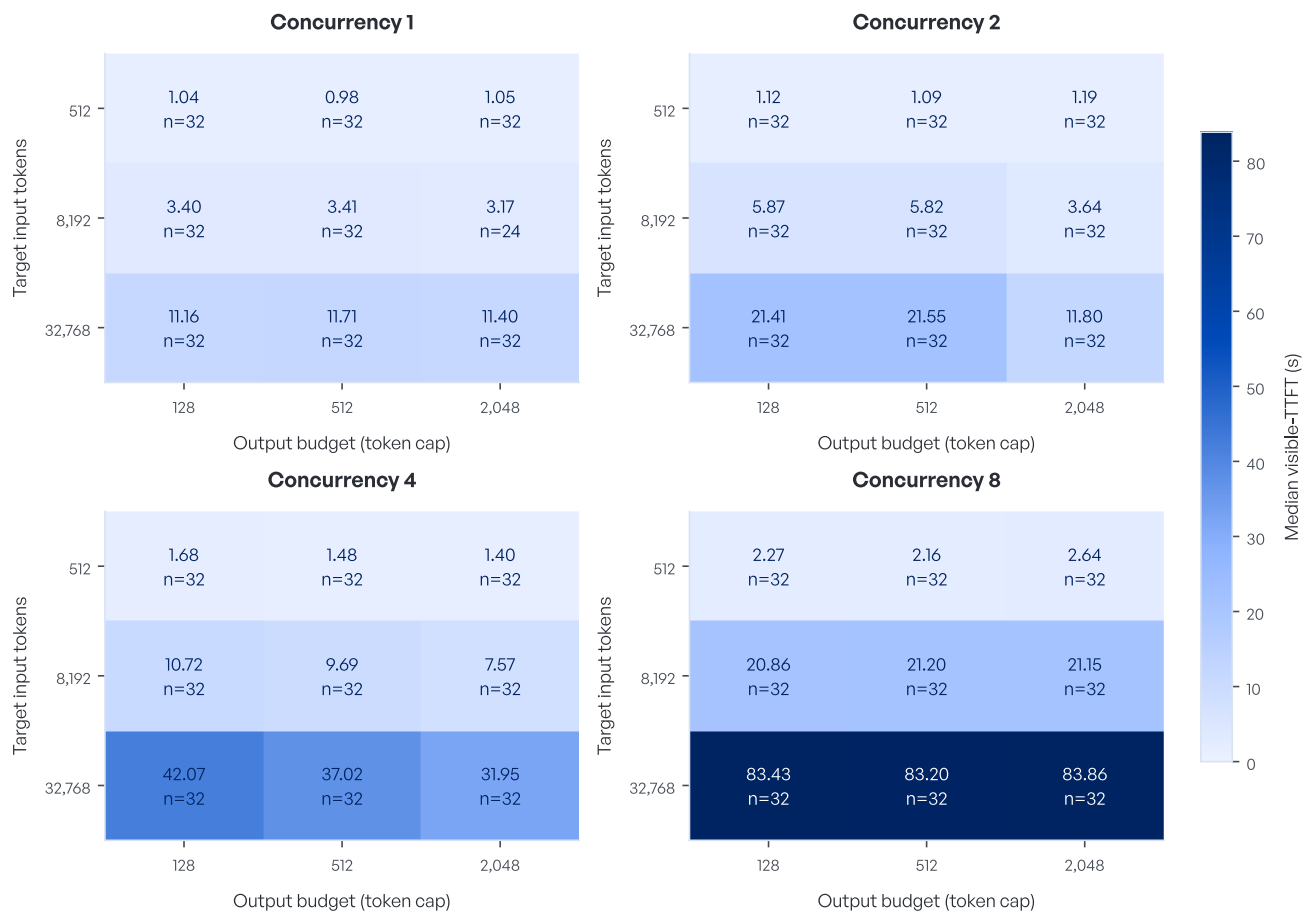


Figure 08–06 **Input/output lengths and visible waiting**

Each cell shows median visible–TTFT and response count n. All four concurrency panels share one scale. This latency is not pure GPU prefill time.

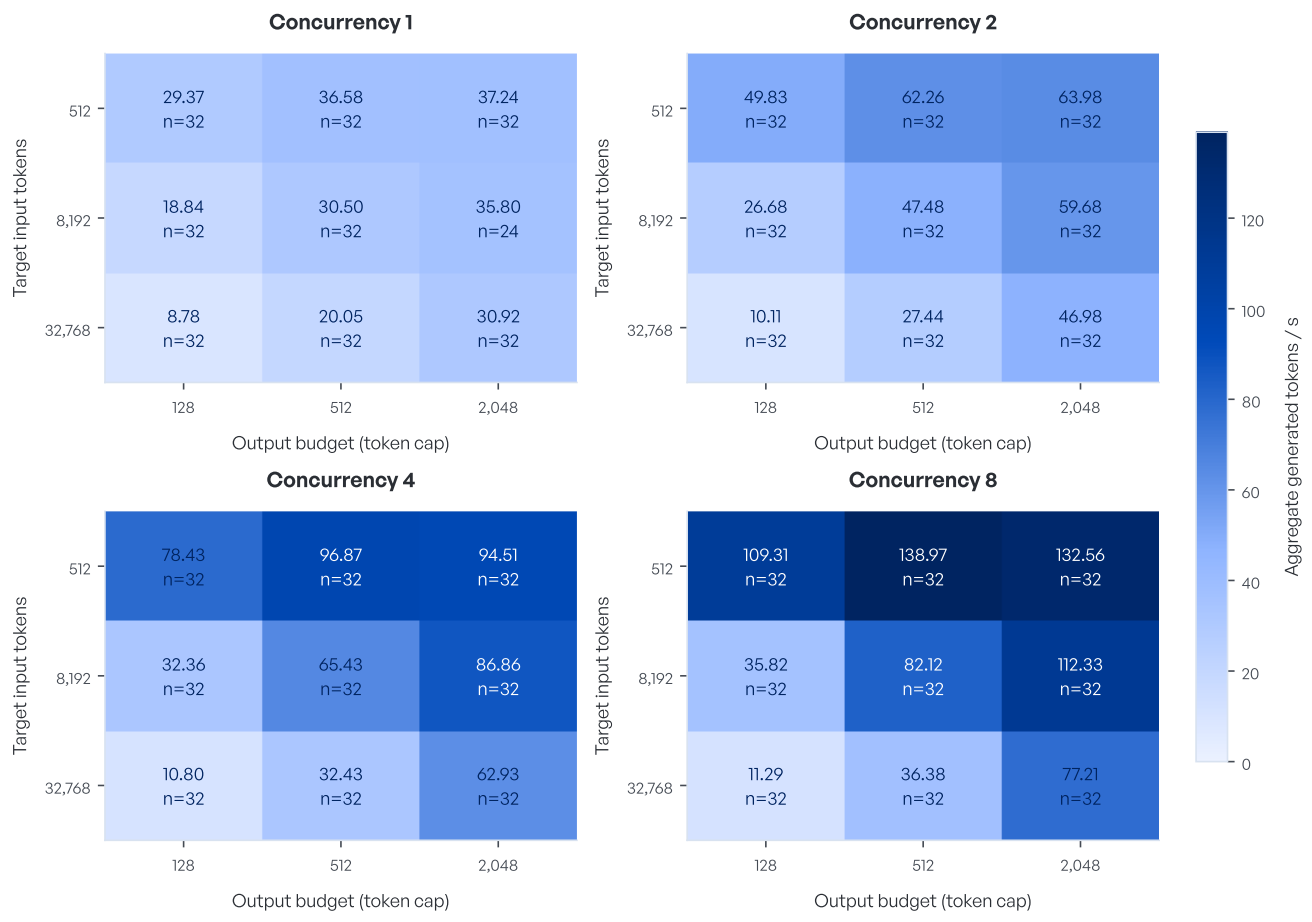


Figure 08–07 Input/output lengths and throughput

Actual generated tokens divided by total complete-block duration. Early stopping changes realized work. A higher output budget is not necessarily more useful business output.

Comparable slice at a fixed output budget

The table fixes the output budget at **512 tokens** while varying input length and concurrency. Realized input includes the chat template; mean realized output reveals differences in generation work.

Table 08–09 Comparable slice at a fixed output budget

Input target / concurrency	Blocks / responses	Mean input / output	Visible–TTFT median s	E2E median s	Aggregate tokens/s
512 / 1	4 / 32	513 / 512.0	0.98	13.97	36.6
512 / 2	4 / 32	513 / 512.0	1.09	16.12	62.3
512 / 4	4 / 32	513 / 512.0	1.48	20.53	96.9
512 / 8	4 / 32	513 / 512.0	2.16	28.51	139.0
8,192 / 1	4 / 32	8,193 / 512.0	3.41	16.78	30.5
8,192 / 2	4 / 32	8,193 / 512.0	5.82	21.24	47.5
8,192 / 4	4 / 32	8,193 / 512.0	9.69	30.69	65.4
8,192 / 8	4 / 32	8,193 / 512.0	21.20	48.44	82.1

Input target / concurrency	Blocks / responses	Mean input / output	Visible–TTFT median s	E2E median s	Aggregate tokens/s
32,768 / 1	4 / 32	32,769 / 512.0	11.71	24.89	20.0
32,768 / 2	4 / 32	32,769 / 512.0	21.55	37.09	27.4
32,768 / 4	4 / 32	32,769 / 512.0	37.02	62.01	32.4
32,768 / 8	4 / 32	32,769 / 512.0	83.20	111.55	36.4

Length effect. At concurrency one and output cap 512, increasing the input target from 512 to **32,768 tokens** changes median visible–TTFT from 0.98 to **11.71 s** and aggregate output from 36.6 to **20.0 tokens**/s. This is end–to–end length sensitivity under the observed prefix state and serving path.

Reproducible data for all length conditions

Table 08–10 **Comparable slice at a fixed output budget**

Context or prefix condition	Complete / failed	First–generated median s
Long context and boundaries · Context boundary probe	0 / 1	—
Long context and boundaries · Target input 131,072 tokens · Concurrency 1 · Generation after long input	3 / 0	48.34
Long context and boundaries · Target input 131,072 tokens · Concurrency 2 · Generation after long input	3 / 0	90.65
Long context and boundaries · Target input 16,384 tokens · Concurrency 1 · Generation after long input	3 / 0	6.77
Long context and boundaries · Target input 16,384 tokens · Concurrency 2 · Generation after long input	3 / 0	11.28
Long context and boundaries · Target input 200,000 tokens · Concurrency 1 · Generation after long input	3 / 0	74.05
Long context and boundaries · Target input 250,000 tokens · Concurrency 1 · Generation after long input	3 / 0	94.08
Long context and boundaries · Target input 32,768 tokens · Concurrency 1 · Generation after long input	3 / 0	12.55
Long context and boundaries · Target input 32,768 tokens · Concurrency 2 · Generation after long input	3 / 0	22.01
Long context and boundaries · Target input 4,096 tokens · Concurrency 1 · Generation after long input	3 / 0	2.54
Long context and boundaries · Target input 4,096 tokens · Concurrency 2 · Generation after long input	3 / 0	2.41
Long context and boundaries · Target input 65,536 tokens · Concurrency 1 · Generation after long input	3 / 0	24.49
Long context and boundaries · Target input 65,536 tokens · Concurrency 2 · Generation after long input	3 / 2	42.67
Prefix caching · History text 0 tokens · First prefix	10 / 0	1.07
Prefix caching · History text 0 tokens · Repeated prefix	10 / 0	0.96
Prefix caching · History text 0 tokens · Appended text	10 / 0	1.06

Context or prefix condition	Complete / failed	First-generated median s
Prefix caching · History text 32,768 tokens · First prefix	10 / 0	11.69
Prefix caching · History text 32,768 tokens · Repeated prefix	10 / 0	1.55
Prefix caching · History text 32,768 tokens · Appended text	10 / 0	1.61
Prefix caching · History text 8,192 tokens · First prefix	10 / 0	3.61
Prefix caching · History text 8,192 tokens · Repeated prefix	10 / 0	1.08
Prefix caching · History text 8,192 tokens · Appended text	10 / 0	1.17

Reasoning budget: visible waiting and output shape

Compare low versus none for the same prompt, round and concurrency; differences are low minus none. Performance analysis retains missing visible content and output-limit events.

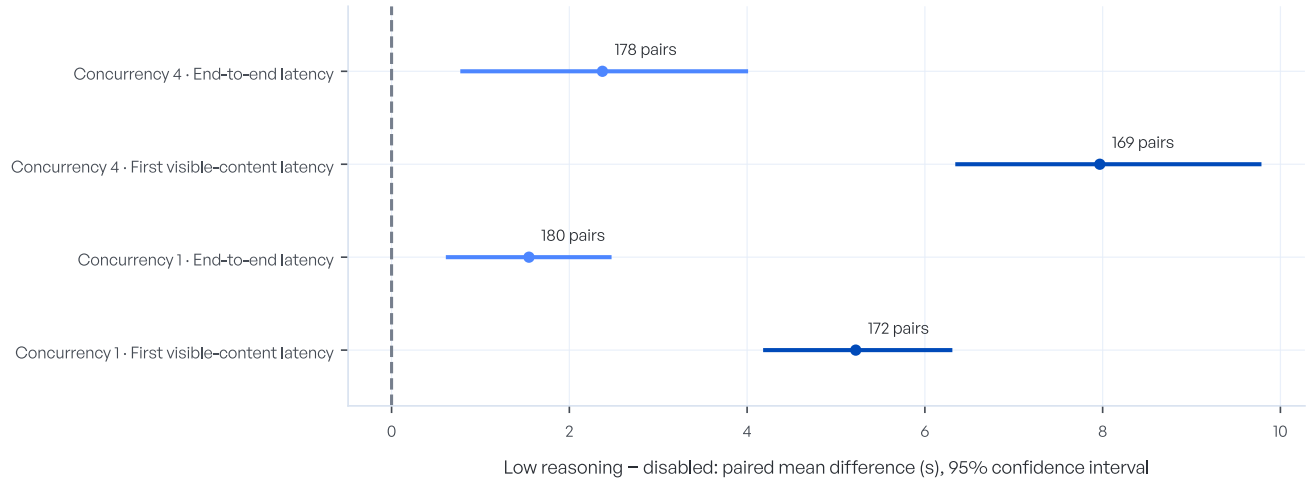


Figure 08-08 Paired reasoning delay

Points are mean differences; lines are 95% prompt-cluster intervals. Pair counts and round sensitivity are in Appendix B.

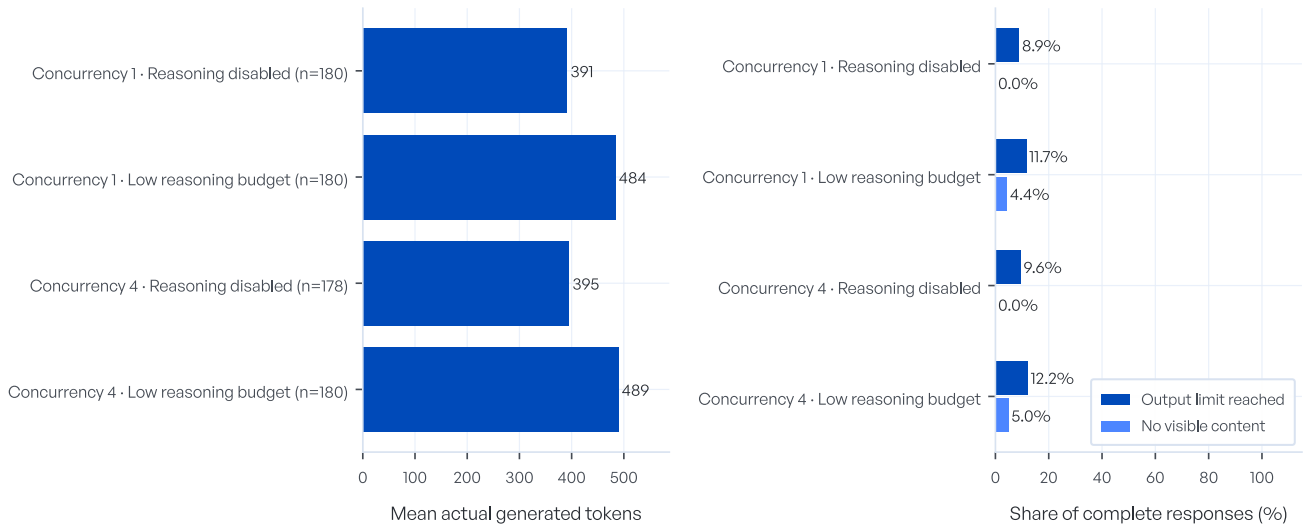


Figure 08–09 Output shape

Read generated volume with output-limit and no-visible-content fractions; the latter two can overlap.

Business shapes and cache ratios: observed coverage

At this snapshot, business-shape and cache-ratio tests have 347 / 352 planned calls completed, including 260 / 264 complete measured responses.

Table 08–11 Business shapes and cache ratios: observed coverage

Shape / concurrency	Complete	Median output	Median realized I/O	Median visible TTFT s	Median E2E s
Long-document analysis / 1	12	2,048.0	16.0	11.92	49.73
Long-document analysis / 4	12	2,048.0	16.0	42.59	101.32
Code editing / 1	12	512.0	16.0	3.64	11.54
Code editing / 4	12	512.0	16.0	11.73	22.86
Multi-turn assistant / 1	12	499.5	16.4	3.49	13.26
Multi-turn assistant / 4	12	505.0	16.2	11.63	26.16
Short brief, long draft / 1	12	1,599.5	0.3	1.10	35.31
Short brief, long draft / 4	8	1,566.5	0.3	2.04	51.05
Structured extraction / 1	12	43.0	47.7	1.49	2.15
Structured extraction / 4	12	43.0	47.7	3.80	4.61
Balanced rewriting / 1	12	510.0	4.0	1.52	13.50
Balanced rewriting / 4	12	486.0	4.2	3.83	20.04
Long-document summary / 1	12	128.0	256.0	11.59	14.60
Long-document summary / 4	12	128.0	256.0	42.35	46.74
Ticket triage / 1	12	25.0	20.6	1.10	1.73
Ticket triage / 4	12	25.0	20.5	2.03	2.80

Table 08–12 Business shapes and cache ratios: observed coverage

Shape / concurrency	Complete continuous-block token/s	Median estimated TPOT ms	Length-limited / complete
Long-document analysis / 1	36.1	21.88	7/12
Long-document analysis / 4	73.1	28.85	7/12
Code editing / 1	38.8	16.13	8/12
Code editing / 4	74.9	22.54	9/12
Multi-turn assistant / 1	31.4	23.95	6/12
Multi-turn assistant / 4	63.5	32.48	6/12
Short brief, long draft / 1	44.9	21.50	1/12
Short brief, long draft / 4	117.9	31.08	0/8
Structured extraction / 1	19.0	15.04	0/12
Structured extraction / 4	36.8	19.41	0/12
Balanced rewriting / 1	45.0	22.44	0/12
Balanced rewriting / 4	98.8	32.49	0/12
Long-document summary / 1	8.6	22.77	12/12
Long-document summary / 4	10.9	35.22	12/12
Ticket triage / 1	13.5	25.98	0/12
Ticket triage / 4	34.2	32.41	0/12

Table 08–13 Business shapes and cache ratios: observed coverage

Input / target overlap	Complete	Realized overlap %	Cache-field coverage	Reported cached %	Median visible TTFT s
32,768 / 0%	6	0.0	0/6	—	12.44
32,768 / 100%	6	100.0	0/6	—	1.91
32,768 / 25%	6	25.0	0/6	—	11.79
32,768 / 50%	6	50.0	0/6	—	9.18
32,768 / 75%	6	75.0	0/6	—	6.77
32,768 / 90%	6	90.0	0/6	—	4.56
8,192 / 0%	6	0.1	0/6	—	3.50
8,192 / 100%	6	100.0	0/6	—	1.14
8,192 / 25%	6	24.9	0/6	—	3.60
8,192 / 50%	6	49.9	0/6	—	3.47
8,192 / 75%	6	74.9	0/6	—	3.46
8,192 / 90%	6	89.9	0/6	—	3.55

Interpret in order: verify realized lengths and truncation → align concurrency and sample composition → inspect waiting and generation → check cache-field coverage → examine contemporaneous resources. Without complete blocks, valid cache fields or matching inputs, describe observations and retain alternative explanations.

Context boundaries and request reuse

One-token output probes observe input-processing waits; long and oversized contexts check the usable range; first, repeated and appended prompts check request reuse. Appendix B gives actual lengths, latency and outcomes. A configured limit is not a validated usable capacity.

Cache reuse and stability

A prefix cache retains intermediate states for processed input so later matching prefixes may reduce input computation. Reuse depends on prefix equality, residency and scheduling. Cache occupancy describes storage state; it is not a hit fraction.

Cache diagnostics: scope and samples

The existing 144-call, two-length/six-overlap design is retained. The following 160 diagnostics are added. Arrows denote target input → output cap; temperature is zero and reasoning is disabled, with identical generation settings within each reuse pair.

Table 08-14 **Cache diagnostics: scope and samples**

Calls	Test	Conditions and meaning
72	Business × reuse × concurrency	Triage 512→64, dialogue 8,192→512, summary 32,768→128; concurrency 1/4; four independent prefixes per condition, each with seed, 50% overlap and exact repeat.
36	Idle retention	8,192/32,768 input × nominal waits 0/30/300 s × three independent prefixes × seed/probe. Each prefix receives one terminal probe.
24	Competition and recovery	Three independent 8,192-token anchors: seed → baseline repeat → four distinct 32,768-token distractors → post-pressure repeat → recovery repeat.
28	Periodic-reuse stability	Four 8,192-token prefixes; four seeds, then six repeat rounds at concurrency four, spaced 120 s apart; planned span ten minutes.
0	Cache under sustained load	Reuse arrival and soak tests; approximately 60-second valid windows align cache fraction, TTFT and occupancy. No extra generation calls.

The three crossed shapes represent short input, conversation and long documents. The other five retain their original performance measurements, without extrapolated cache findings. Concurrency arms use independent prefixes; 50%/100% probe order alternates. A seed is a fresh-prefix reference, not proof of an empty engine cache. Actual idle time runs from seed completion to probe dispatch; zero denotes no intentional wait, with telemetry overhead retained. Timing deviations, retries, telemetry gaps and counter mismatches are flagged and excluded from cache attribution.



Figure 08–10 Cache evaluation sequence

Reuse, idle retention, pressure/recovery and temporal behavior. Conditions are prospective.

Observed evidence and coverage

The supplement has completed **160/160 calls** with 85 attributable counter windows. Interpret business, actual idle gaps, pressure and periodic reuse separately; incomplete or invalid cells remain unresolved. Across 30 observed first/repeat pairs, 26 repeats have lower TTFT. This supports a waiting-time observation, not isolated causal cache benefit. Across 358 valid active node-counter windows, the token-weighted cached fraction is 20.3%. This describes the input mix in those windows, not per-request hits or long-term stability.

Table 08–15 Observed evidence and coverage

Workload	Complete	Cache fields valid / missing
Interface and measurement calibration	12	0 / 12
Natural workloads and concurrency scaling	718	0 / 718
Input/output length matrix	1435	0 / 1435
Paired reasoning budgets	360	0 / 360
Long context and boundaries	36	0 / 36
Prefix caching	90	0 / 90
Business shapes and cache-ratio amendment	507	0 / 507
Open arrivals and mixed workloads	783	0 / 783
Length expansion and sustained operation	925	0 / 925

Paired waiting observations from existing first/exact-repeat sequences follow. First denotes first within the sequence, not guaranteed absence of prior cache state. Differences are calculated per pair before taking the median.

Table 08–16 Observed evidence and coverage

History tokens	Pairs	Input median	First TTFT s	Repeat TTFT s	Paired reduction s	Repeat faster
0	10	107	1.07	0.96	0.08	6/10

History tokens	Pairs	Input median	First TTFT s	Repeat TTFT s	Paired reduction s	Repeat faster
8192	10	8,298	3.61	1.08	2.49	10/10
32768	10	32,875	11.69	1.55	10.04	10/10

Coverage counts cache fields in original upstream usage only. Zeros in gateway-normalized responses remain unverified until matched to upstream fields and are excluded from hit-fraction denominators.

The archived SGLang configuration has `enable_cache_report=false` while prefix caching is enabled. Missing per-request fields therefore do not establish a miss. SparkDash exposes cumulative prompt and cached-token counters; isolated diagnostic blocks use fresh boundary deltas reconciled with the same requests. API fields, node counters and their coverage remain separate. SGLang v0.5.21 · SparkDash counter implementation

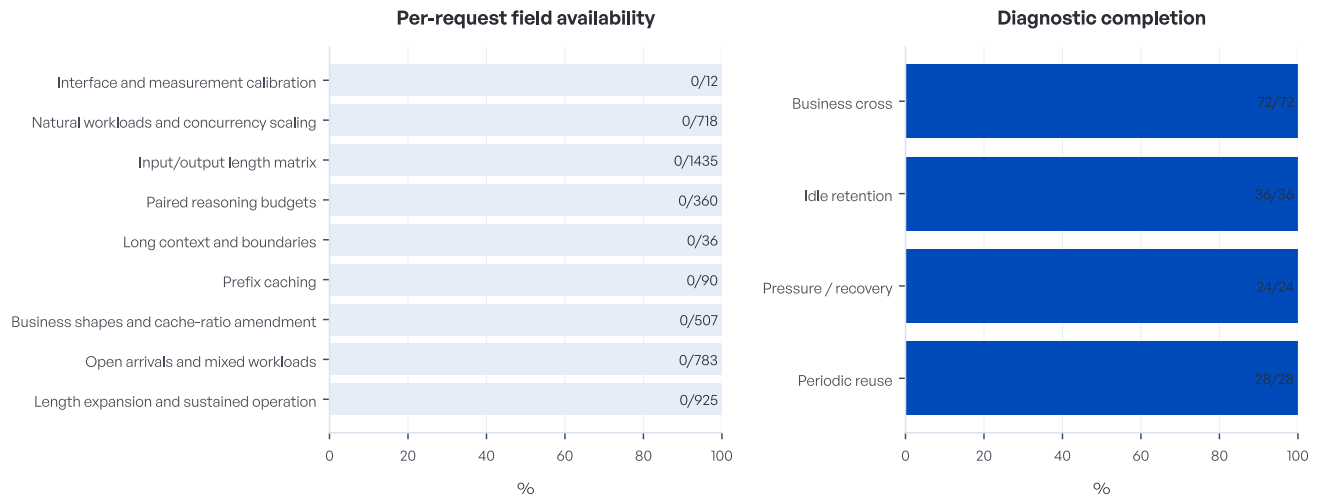


Figure 08-11 Evidence coverage

Left: per-request cache-field availability. Right: completed diagnostic calls. Neither is a cache-hit rate.

Table 08-17 Observed evidence and coverage

Condition	Conc.	n	TTFT median s	Paired reduction s	Valid windows / cache %
Business / triage / Seed	4	4	2.48	0.00	1 / 0.0%
Business / triage / 50%	4	4	2.59	−0.09	1 / 0.0%
Business / triage / 100%	4	4	1.62	0.88	1 / 99.8%
Business / dialogue / Seed	1	4	3.44	0.00	1 / 0.0%
Business / dialogue / 100%	1	4	3.76	−0.16	1 / 0.0%
Business / dialogue / 50%	1	4	3.35	−0.06	1 / 0.0%
Business / summary / Seed	1	4	13.58	0.00	1 / 0.0%
Business / summary / 50%	1	4	9.72	3.83	1 / 25.0%
Business / summary / 100%	1	4	1.78	11.62	1 / 100.0%
Business / triage / Seed	1	4	1.19	0.00	1 / 0.0%
Business / triage / 100%	1	4	0.91	0.29	1 / 99.8%
Business / triage / 50%	1	4	1.17	0.08	1 / 0.0%

Condition	Conc.	n	TTFT median s	Paired reduction s	Valid windows / cache %
Business / summary / Seed	4	4	43.15	0.00	1 / 0.0%
Business / summary / 50%	4	4	25.90	17.25	1 / 43.5%
Business / summary / 100%	4	4	2.91	40.25	1 / 99.8%
Business / dialogue / Seed	4	4	11.22	0.00	1 / 0.0%
Business / dialogue / 100%	4	4	9.89	1.34	1 / 24.8%
Business / dialogue / 50%	4	4	12.10	−0.88	1 / 0.0%
Idle / 32768 / Seed / 30 s	1	3	13.63	—	3 / 0.0%
Idle / 32768 / Probe / 30 s	1	3	2.22	10.53	3 / 100.0%
Idle / 32768 / Seed / 0 s	1	3	13.26	—	3 / 0.0%
Idle / 32768 / Probe / 0 s	1	3	2.79	10.67	3 / 100.0%
Idle / 32768 / Seed / 300 s	1	3	12.89	—	3 / 0.0%
Idle / 32768 / Probe / 300 s	1	3	3.69	9.96	3 / 91.7%
Idle / 8192 / Seed / 0 s	1	3	4.75	—	3 / 0.0%
Idle / 8192 / Probe / 0 s	1	3	2.03	3.08	3 / 100.0%
Idle / 8192 / Seed / 300 s	1	3	4.84	—	3 / 0.0%
Idle / 8192 / Probe / 300 s	1	3	2.67	2.13	3 / 100.0%
Idle / 8192 / Seed / 30 s	1	3	4.39	—	3 / 0.0%
Idle / 8192 / Probe / 30 s	1	3	1.71	2.43	3 / 100.0%
Pressure / 8192 / Seed	1	3	4.33	0.00	3 / 0.0%
Pressure / 8192 / Baseline	1	3	1.76	2.50	3 / 100.0%
Pressure / 8192 / Post−pressure	1	3	1.70	2.63	3 / 100.0%
Pressure / 8192 / Recovery	1	3	1.76	2.52	3 / 100.0%
Periodic / 8192 / Seed	4	4	11.42	—	1 / 0.0%
Periodic / 8192 / Probe / 1	4	4	2.02	9.38	1 / 100.0%
Periodic / 8192 / Probe / 2	4	4	2.16	9.25	1 / 100.0%
Periodic / 8192 / Probe / 3	4	4	2.18	9.26	1 / 100.0%
Periodic / 8192 / Probe / 4	4	4	2.26	9.15	1 / 100.0%
Periodic / 8192 / Probe / 5	4	4	2.39	9.03	1 / 100.0%
Periodic / 8192 / Probe / 6	4	4	2.33	9.08	1 / 100.0%

Cache variation over time

358 valid windows contain prompt activity, with a token−weighted cache fraction of **20.29%** and a window range of **0.00%**—100.00%. Windows may contain different workloads; they describe temporal behavior rather than proving retention or stable hits. The sample SD across window fractions is 25.77 percentage points.

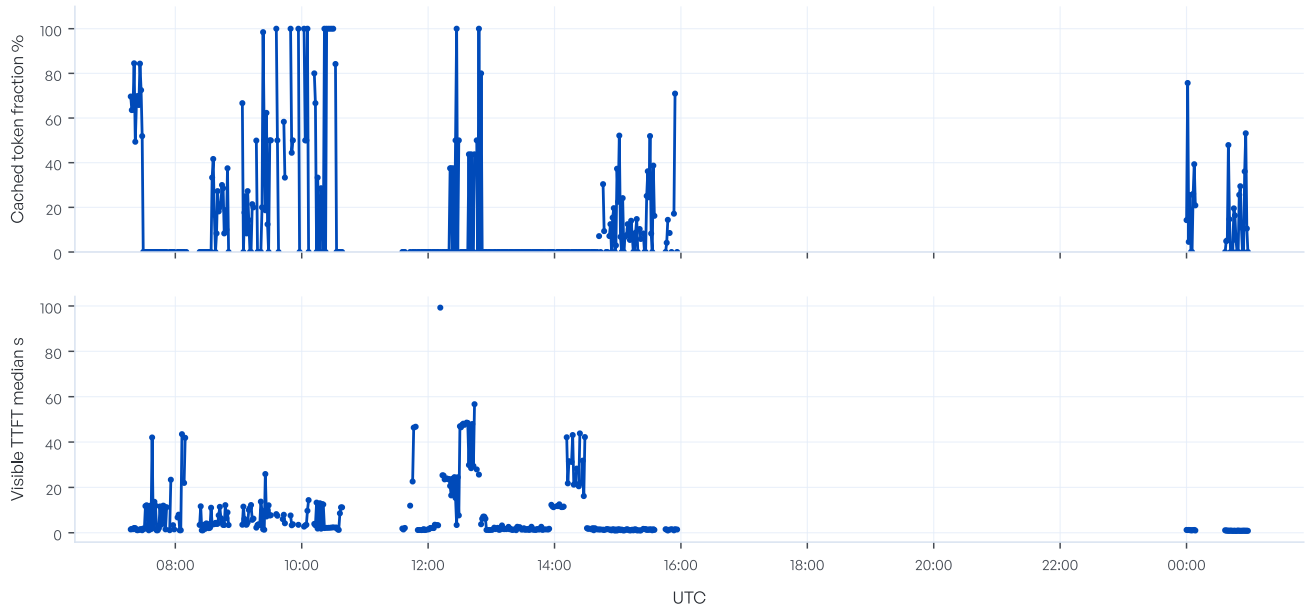


Figure 08–12 Time observation

Counter–delta fractions in approximately 60–second windows with contemporaneous TTFT; workload changes prevent simple causal attribution.

Interpretation and further validation

Only matched–parameter, complete, retry–free blocks with reconciled counters receive block token–cache fractions; concurrent blocks do not yield per–request hit fractions. Idle prefixes are probed once to avoid refreshing retention. Pressure tests insert about 131,072 distractor input tokens, without claiming cache–capacity exhaustion. Periodic reuse spans ten planned minutes; sustained load uses actual observed duration. Further work should vary business popularity, independent periods and competition strength. Capacity and cost models receive no unverified cache benefit.

Expanded lengths: coverage, realized output and repeatability

Length expansion: 290/299 complete, 205 outputs reached **98–100%** of cap. Maximum completed target–attaining output is **4,096 tokens**, limited to its input and probe conditions.

Public length expansion uses natural stopping as defined in section 03; requested caps and realized generation are reported together.

Table 08–18 Expanded lengths: coverage, realized output and repeatability

Load	Complete / planned	Upstream counts	Actual output range	Median attainment	Output length attained
Output calibration	2/2	2	2–2	0.2%	0
Input axis	24/24	24	1,024–1,024	100.0%	24
Output axis	9/13	9	1,826–4,955	100.0%	5
Balanced lengths	9/12	9	2,048–5,943	100.0%	6
Near boundary	1/2	1	1,024–1,024	100.0%	1
Long prefixes	36/36	36	1,024–1,024	100.0%	36
Business anchors	31/32	31	25–4,096	16.9%	Natural EOS

Load	Complete / planned	Upstream counts	Actual output range	Median attainment	Output length attained
Representative warmups	10/10	10	1,024–4,096	100.0%	7
Repeated blocks	168/168	168	1,024–4,096	100.0%	124

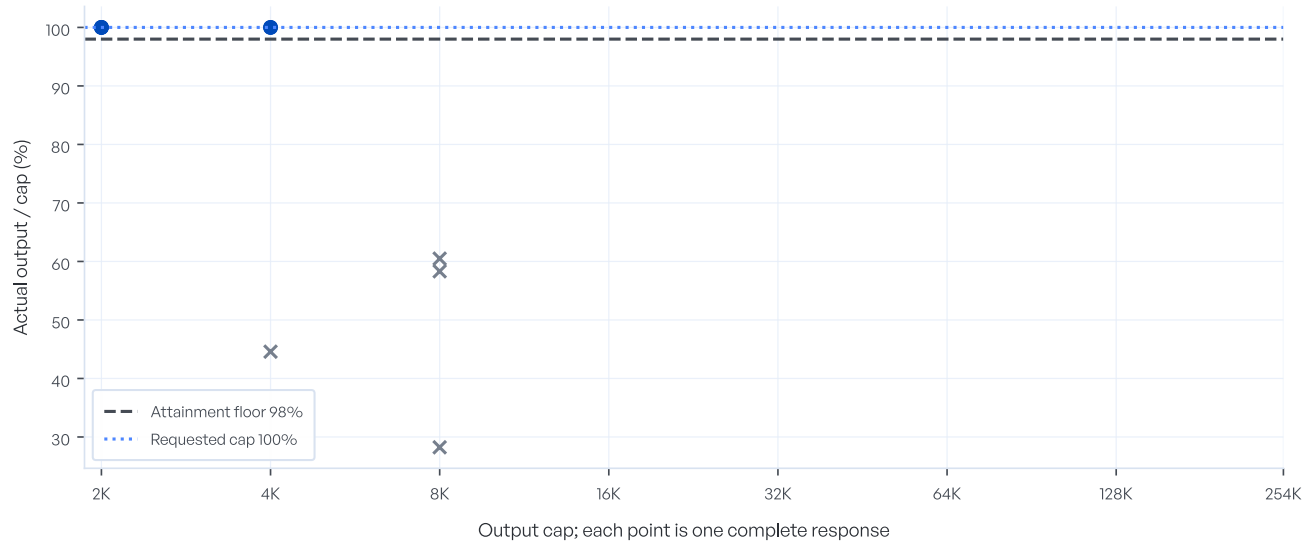


Figure 08–13 Output attainment

Compare original upstream generation with requested caps; missing counts remain blank, not zero.

The following includes only complete requests with input error ⩽1% and realized output at 98–100% of cap. Values are medians; a one-sample cell is that observation, with no tail percentiles.

Table 08–19 Expanded lengths: coverage, realized output and repeatability

Input → output cap	Eligible / complete	Actual output	Visible TTFT s	E2E s	Estimated TPOT ms
1K → 1K	3/3	1,024	1.20	25.14	23.2
2K → 1K	3/3	1,024	1.45	23.38	21.5
4K → 1K	3/3	1,024	2.12	25.90	23.2
8K → 1K	3/3	1,024	3.60	25.49	21.6
16K → 1K	3/3	1,024	5.91	26.80	20.0
32K → 1K	3/3	1,024	12.00	33.21	20.8
64K → 1K	3/3	1,024	23.33	45.91	21.8
128K → 1K	3/3	1,024	46.47	68.05	21.1
1K → 2K	3/3	2,048	1.30	47.45	22.6
1K → 4K	2/3	4,096	1.20	88.58	21.3
2K → 2K	3/3	2,048	1.46	45.56	21.4
4K → 4K	3/3	4,096	2.08	88.44	21.1

Input → output cap	Eligible / complete	Actual output	Visible TTFT s	E2E s	Estimated TPOT ms
254K → 1K	1/1	1,024	99.27	121.45	21.7

6 repeat blocks currently qualify for capacity comparison; boundary probes are excluded from steady-capacity claims.

Table 08-20 **Expanded lengths: coverage, realized output and repeatability**

Input → output / concurrency	Eligible blocks	Aggregate tokens/s	Block-rate range
32K → 1K / 1	3	33.0	31.0–34.8
32K → 1K / 4	3	55.5	54.9–55.9

The 64K/128K prefix diagnostic currently has 17/18 valid probe windows; observed cache fractions and intended overlap remain distinct.

Table 08-21 **Expanded lengths: coverage, realized output and repeatability**

Input / target overlap	Valid / planned pairs	Window cache fraction median / range
64K / 0%	2/3	0.0% / 0.0%–0.0%
64K / 50%	3/3	37.5% / 37.5%–37.5%
64K / 100%	3/3	100.0% / 100.0%–100.0%
128K / 0%	3/3	0.0% / 0.0%–0.0%
128K / 50%	3/3	43.7% / 43.7%–43.7%
128K / 100%	3/3	100.0% / 100.0%–100.0%

Early stopping, parameter rejection, insufficient time and unstarted cells retain their distinct states. New length findings are not extrapolated to untested concurrency or annual return. CSV · JSON

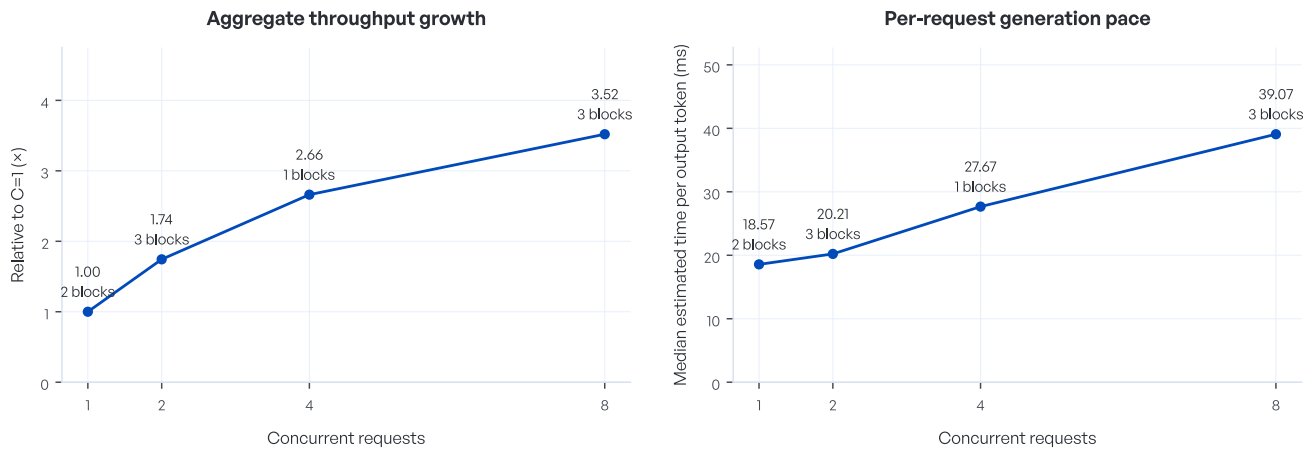


Figure 08-14 **Concurrency benefit and cost**

Left: aggregate throughput relative to C=1 on the same prompt bank. Right: median estimated time per output token. Eligible block counts remain visible; this is descriptive, not a paired causal effect.

Fixed long output: demonstrated native-engine range

This arm uses authenticated native generation with exact token-ID inputs, temperature zero, ignored end tokens and a fixed generation count; the stream interval is 128 tokens. Decode rates use same-process engine counts and times, not public first-content latency. Input plus output stays within 262,144. The 1,024-input / 260,096-output case reserves **1,024 tokens** of headroom, so its longest output target is 254K.

Table 08-22 **Fixed long output: demonstrated native-engine range**

Input → output tokens	Complete / planned	Decode token/s, median [range]	Median engine time s	Draft acceptance fraction
1,024 → 1,024	1/1	62.24 [62.24–62.24]	16.91	0.940
1,024 → 8,192	3/3	55.64 [54.17–59.24]	147.62	0.832
1,024 → 16,384	3/3	61.35 [54.15–61.51]	267.46	0.931
1,024 → 32,768	3/3	60.50 [56.93–61.17]	541.80	0.918
1,024 → 65,536	1/1	60.32 [60.32–60.32]	1,086.82	0.924
1,024 → 131,072	1/1	61.96 [61.96–61.96]	2,115.69	0.959
1,024 → 260,096	1/1	61.61 [61.61–61.61]	4,222.17	0.961
16,384 → 16,384	1/1	62.10 [62.10–62.10]	270.54	0.963
32,768 → 32,768	1/1	60.88 [60.88–60.88]	548.28	0.938
65,536 → 65,536	1/1	60.27 [60.27–60.27]	1,109.63	0.926

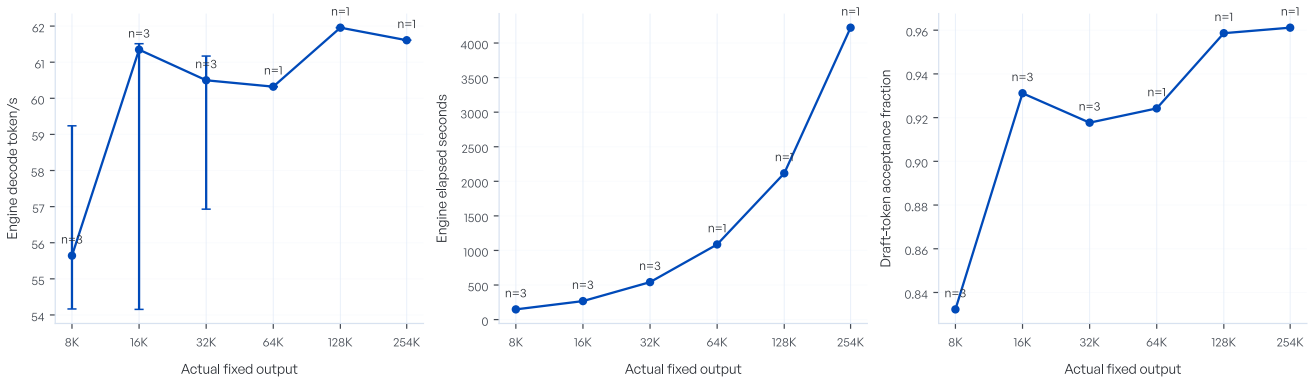


Figure 08-15 **Long-output generation rate and work**

Input is fixed at 1K. Error bars are observed ranges of completed independent requests; n=1 is one observation. Draft acceptance is the native field. Unfinished targets have no performance point.

Forced length may produce repeated text, whitespace or special tokens. Count attainment establishes execution at that length, not useful content. MTP acceptance depends on the generated sequence, so these rates do not directly predict every business workload or annual production. Three-request cells report observed ranges; longer single probes provide no tail percentile or population confidence interval.

Visible-text checks flag 8 complete requests with runs of at least 256 identical characters. Flags accompany the raw outputs and are not answer-quality scores.

Additional requests and timings CSV

GPU profiling across conditions and measurement overhead

Each pair uses identical inputs and 256 output tokens, with alternating arm order. Prefix cache is flushed on an idle service before each arm. Profiling captures prefill and four decode steps; the full-request time ratio measures the impact of this diagnostic on that request, not a correction factor for all historical data.

Table 08-23 GPU profiling across conditions and measurement overhead

Input / concurrency	Pairs	Profiled / unprofiled time ratio	Cached tokens across both arms
1,024 / 1	3	1.162 [1.124–1.220]	0
8,192 / 4	3	1.236 [1.182–1.263]	0
32,768 / 1	3	1.084 [0.931–1.186]	0

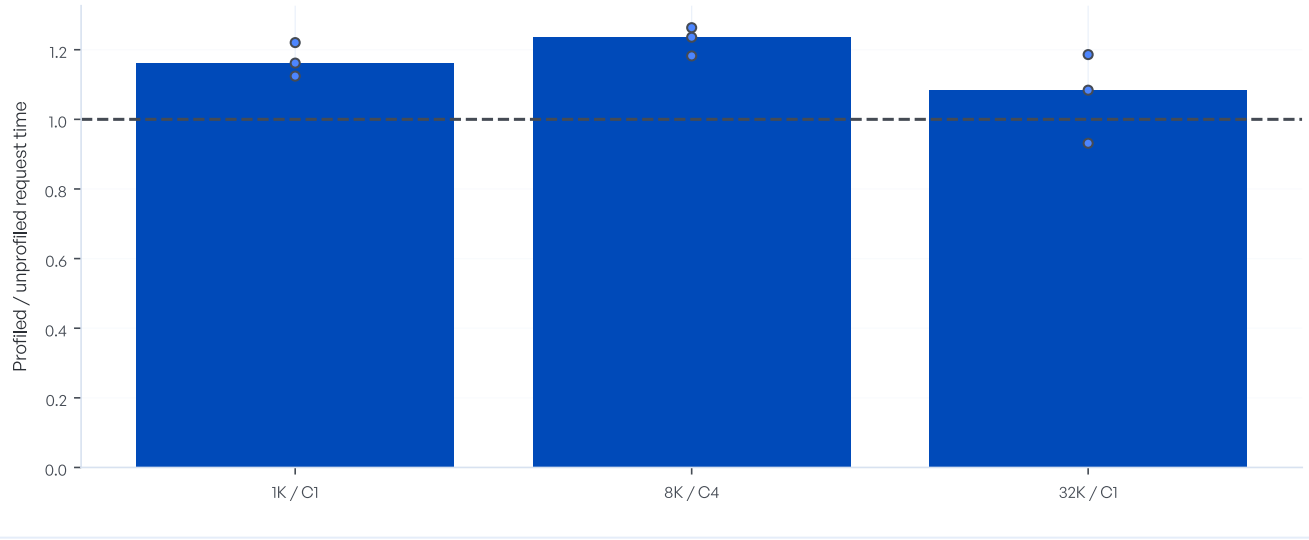


Figure 08-16 Paired profiling overhead

Bars show medians and points show independent pairs; 1 means equal request duration.

Independent check of 64K prefix reuse

Table 08-24 Independent check of 64K prefix reuse

Independent prefix	First / reuse	Actual cached tokens	Cached fraction	Engine first-token seconds
2	First	0	0.000	20.564
2	Reuse	57344	0.875	2.766
3	First	0	0.000	20.746
3	Reuse	57344	0.875	2.682
1	First	0	0.000	20.457
1	Reuse	57344	0.875	2.664

Each prefix is sent after a cache flush and reused immediately. This table tests immediate reuse, not long-term retention or cache capacity.

09 · Sustained performance, thermal state and resources

Independent arrivals and sustained operation

Table 09–01 Independent arrivals and sustained operation

Block	Offered / sent / complete	Local rejections	Window s
Open arrivals and mixed workloads · 0.5 times reference rate	120 / 120 / 120	0	624.9
Open arrivals and mixed workloads · 0.25 times reference rate	57 / 57 / 57	0	607.9
Open arrivals and mixed workloads · 1.25 times reference rate	296 / 243 / 243	53	617.4
Historical observation (5 segments, including gaps)	638 / 638 / 635	0	34,580.3
Open arrivals and mixed workloads · 1 times reference rate	226 / 203 / 203	23	616.4
Open arrivals and mixed workloads · 0.75 times reference rate	160 / 160 / 160	0	612.3

Analytical thresholds: visible TTFT⩽**1/2/5/10 s** and estimated TPOT⩽20/50/100 ms. These are sensitivity conditions, not committed service–level agreements. Global concurrency is capped at eight, queue length at 32 and queue wait at 5 s. The sustained–stability test ends at **1 hour**,**2,701 requests** or the budget/deadline, whichever occurs first. Per–condition qualified–throughput data

The table selects three analytical threshold pairs; cells show qualified requests/s and their fraction of all measured arrivals. The data table contains the remaining pairs. Mixed–load results evaluate matched short requests only.

Table 09–02 Independent arrivals and sustained operation

Condition	1 s / 20 ms	2 s / 50 ms	5 s / 100 ms
Open arrivals and mixed workloads · 0.25 times reference rate	0.012 / 12.2%	0.058 / 61.0%	0.096 / 100.0%
Open arrivals and mixed workloads · 0.5 times reference rate	0.013 / 6.7%	0.150 / 77.9%	0.185 / 96.2%
Open arrivals and mixed workloads · 0.75 times reference rate	0.009 / 3.5%	0.212 / 81.2%	0.253 / 97.2%
Open arrivals and mixed workloads · 1 times reference rate	0.000 / 0.0%	0.208 / 57.1%	0.304 / 83.8%
Open arrivals and mixed workloads · 1.25 times reference rate	0.000 / 0.0%	0.154 / 31.8%	0.326 / 67.5%

Node resources

Compare means aligned to the same concurrency blocks, then inspect changes over the observation period. Telemetry gaps are not interpolated.

Resources aligned with natural–workload blocks

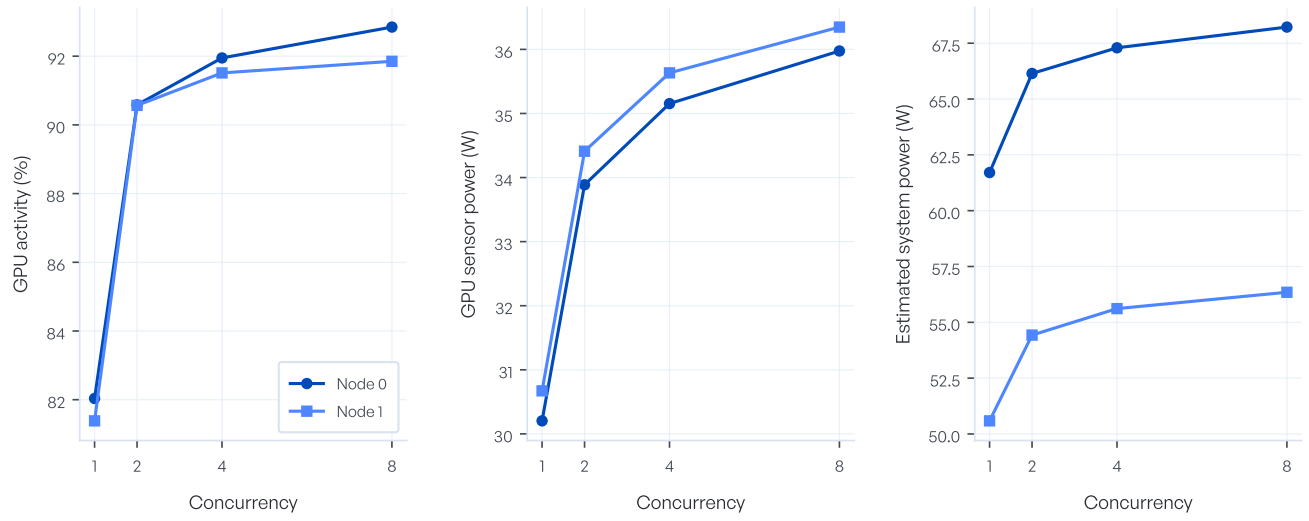


Figure 09–01 Concurrency and node resources

Same natural–workload blocks as the throughput comparison; means are weighted by valid sample duration. Gaps over 20 seconds or missing fields are not interpolated. GPU sensor and estimated system power are separate views and are not added.

Table 09–03 Resources aligned with natural–workload blocks

Concurrency / node	Mean GPU activity	GPU / estimated system W	Peak temperature ˚C	Activity / power coverage
1 / 0	82.0%	30.2 / 61.7	65.0	100.0% / 100.0%
1 / 1	81.4%	30.7 / 50.6	66.0	100.0% / 100.0%
2 / 0	90.6%	33.9 / 66.1	70.0	100.0% / 100.0%
2 / 1	90.6%	34.4 / 54.4	70.0	100.0% / 100.0%
4 / 0	91.9%	35.2 / 67.3	70.0	100.0% / 100.0%
4 / 1	91.5%	35.6 / 55.6	70.0	100.0% / 100.0%
8 / 0	92.8%	36.0 / 68.2	68.0	100.0% / 100.0%
8 / 1	91.9%	36.3 / 56.3	69.0	100.0% / 100.0%

Coverage is the fraction of condition–block time supported by valid adjacent samples; means use the covered intervals. GPU activity near **100%** alone does not establish compute or memory–bandwidth saturation. Between–node power differences also require clocks, work allocation and sensor semantics.

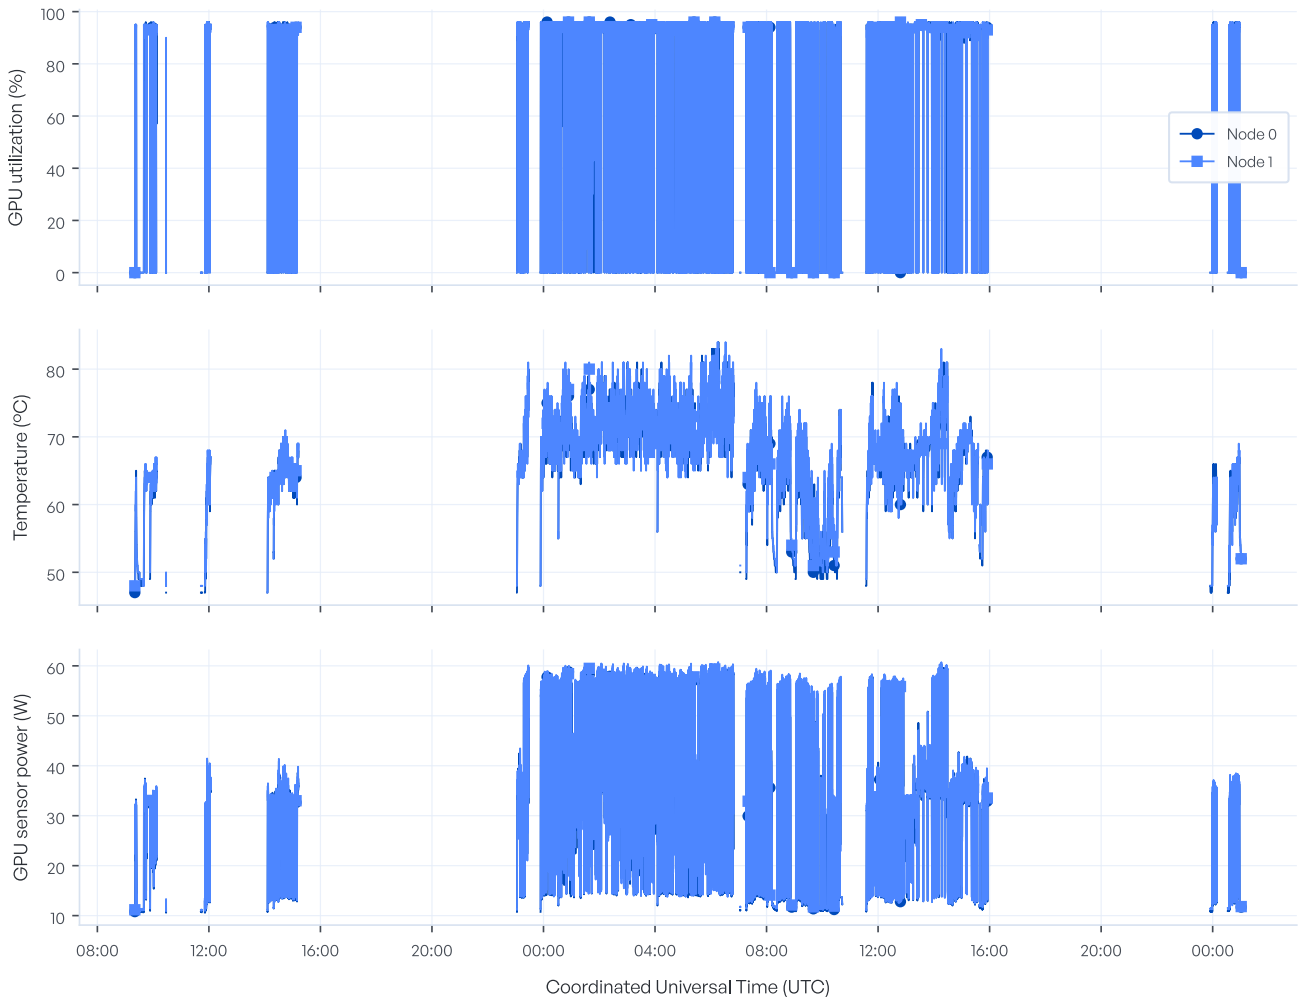


Figure 09–02 **Node timelines**

Includes warmups and idle time; load comparisons use the aligned table above.

Raw engine and memory timelines are in Appendix B. The separate network microbenchmark is about **109 Gb/s** for that transfer test; current evidence does not isolate networking, memory or compute as the sole model bottleneck.

The opening telemetry belongs to serving windows; independent GPU sustained work is reported below by actual coverage. Their work and baselines differ and durations cannot be combined into one stability claim. Keep sensor, estimated system and outlet–metered energy separate; absent outlet measurements, cost remains scenario–based.

Independent hardware sustained load: performance retention and thermal state

The schedule is **60 minutes** of compute, **60 minutes** of bandwidth and **120 minutes** of mixed work. Each phase first records a five–minute matched–work baseline, outside its formal duration. Mixed work alternates compute and bandwidth, retaining separate rates. Retention divides each ten–minute window's median by the same–work baseline median. Windows, phases and serving durations are distinct and cannot be combined into continuous business–service uptime.

Mixed work uses an independent warmup and baseline. GPU submissions pause at **85°C** and resume below75°C after30stable seconds. Timing and cooling belong to this execution; earlier interrupted attempts remain separate. Compute and bandwidth retain their original85°C cutoff, so thermal control is not identical across phases.

Table 09–04 Independent hardware sustained load: performance retention and thermal state

Node	Workload	Observed / planned min	Validated samples	Median / max °C	Median SM MHz	Sensor kWh	Integration–time coverage
0	Compute	60.0 / 60	3359	81.0 / 82.0	2262	0.0888	100.0%
0	Bandwidth	60.0 / 60	3485	65.0 / 66.0	2522	0.0318	100.0%
0	Alternating mixed	120.0 / 120	5832	75.0 / 86.0	2405	0.1082	100.0%
1	Compute	60.0 / 60	3405	83.0 / 84.0	2249	0.0888	100.0%
1	Bandwidth	60.0 / 60	3490	67.0 / 67.0	2515	0.0324	100.0%
1	Alternating mixed	120.0 / 120	6527	74.0 / 86.0	2431	0.1178	100.0%

Mixed–load thermal control and throughput boundaries

Table 09–05 Mixed–load thermal control and throughput boundaries

Control	Rule
Backoff trigger	Temperature ⩾85°C or unavailable: pause GPU submissions
Automatic recovery	Temperature below 75°C for 30 continuous seconds
Maximum individual backoff	600 seconds; end the run and restore service on timeout
Time and stop boundaries	No automatic deadline extension; manual stop, lease and memory guards remain effective

Mixed work uses this policy; the compute and bandwidth phases retain their respective measurement–time thermal conditions. Cooling occurs at CUDA submission boundaries and is excluded from CUDA–event rates. Effective throughput separately counts measured work over the whole phase wall time, including cooling, validation and the other mixed workload. Backoff events are archived separately; unfinished timed calls are not valid performance samples.

Driver limit reasons | Active means the driver reported that reason as active; Not Active means inactive at that sample. Unsupported/missing values are not counted as inactive. Counts do not directly establish throttling duration or performance loss.

Table 09–06 Mixed–load thermal control and throughput boundaries

Node	Phase	Driver reason flag	Active samples	Known–state samples
0	mixed	sw_thermal_slowdown	0	7196
0	mixed	hw_thermal_slowdown	0	7196
0	mixed	hw_slowdown	0	7196
0	mixed	sw_power_cap	0	7196
1	mixed	sw_thermal_slowdown	0	7197
1	mixed	hw_thermal_slowdown	0	7197
1	mixed	hw_slowdown	0	7197
1	mixed	sw_power_cap	0	7197

Mixed work: execution rate and throughput including cooling

Table 09–07 **Mixed work: execution rate and throughput including cooling**

Node	Median execution rate	Phase–wall throughput	Cooling seconds in completed calls	Cooling / phase wall time	Calls containing cooling
0	gemm: 87.57 TFLOP/s; triad: 221.62 GB/s	gemm: 36.39 TFLOP/s; triad: 91.72 GB/s	1178.8	16.4%	38
1	gemm: 86.99 TFLOP/s; triad: 217.28 GB/s	gemm: 40.63 TFLOP/s; triad: 100.93 GB/s	434.4	6.0%	14

Phase–wall throughput divides counted work of each type by the entire mixed–phase duration, including the other workload, validation and cooling. Its difference from execution rate is therefore not solely a cooling effect. Cooling seconds sum completed timed calls; events from interrupted calls remain in the raw records.

Cooling in this run was triggered by the benchmark temperature threshold. Valid driver samples on both nodes recorded no active thermal, hardware–slowdown or power–cap flags. This describes approximately one–second sampling and cannot exclude events between samples. Capacity assessment should consider execution rates together with wall–clock throughput including waits.

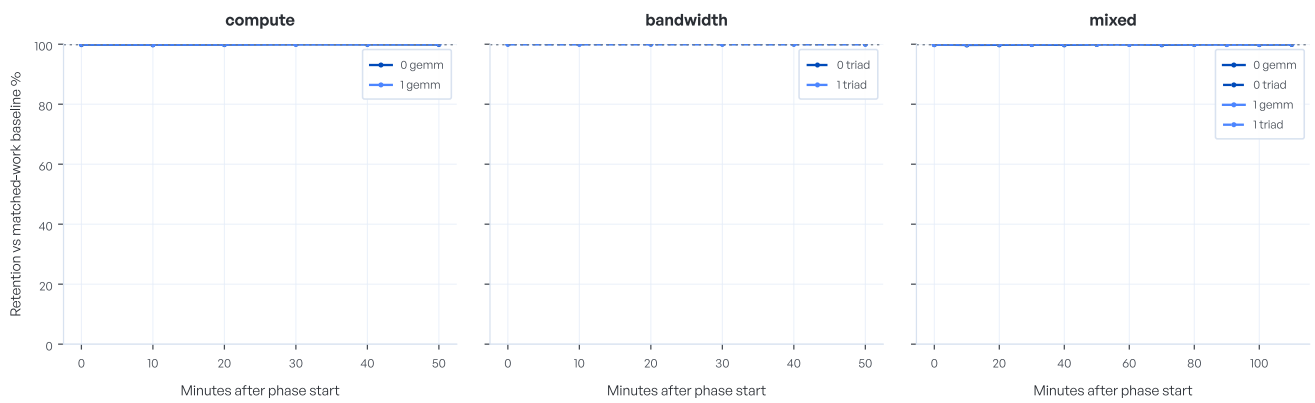


Figure 09–03 **Windowed performance retention**

Only complete ten–minute windows are plotted; partial–tail samples remain in the statistical files. Compare matching node, workload and phase; 100% denotes the five–minute phase baseline, not a theoretical peak.

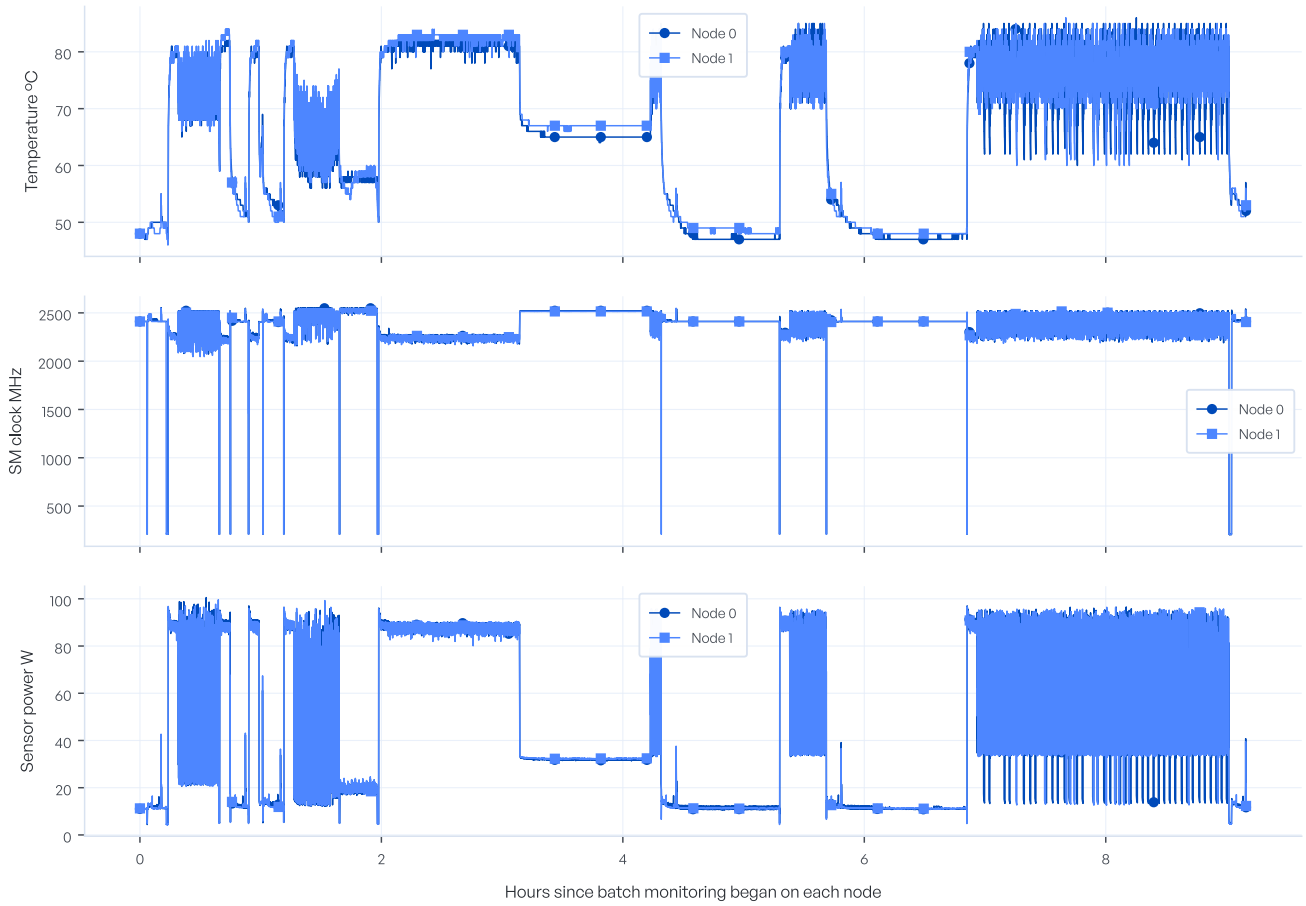


Figure 09–04 **Node resource timeline**

Includes warmup, individual scans and sustained work; phase receipts define each window. Lines break across gaps exceeding five seconds. Sensor power is not whole–system wall power.

Energy uses trapezoidal integration of adjacent valid power samples no more than five seconds apart, reporting time coverage without zero–filling or extrapolating gaps. No wall–meter evidence exists. Depreciation and 1/3/5–year returns retain the explicit whole–system electricity scenarios in section 11; sensor energy does not replace wall–electricity cost.

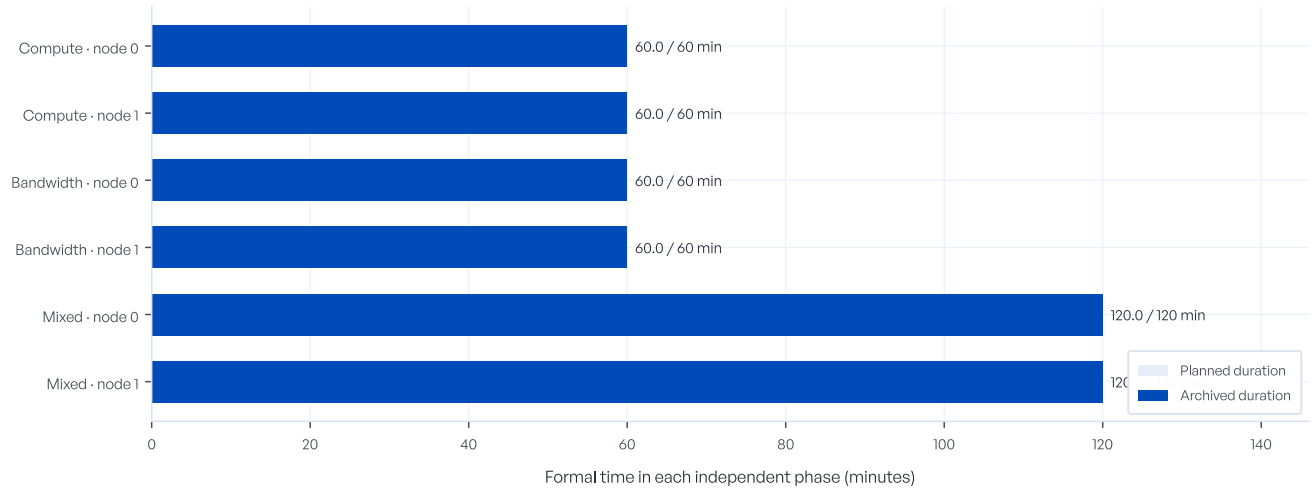


Figure 09–05 Sustained coverage

Solid bars show valid formal duration, pale bars show the planned remainder. Phases are independent and do not establish continuous stability.

The original sustained arm has 638 arrivals in five segments. Its 34,580.3–second span includes interruptions and does not establish continuous stability or continuous goodput. Any completed independent window below is reported separately, without joining it to those segments.

Independent public blocks, interference and continuous window

Table 09–08 Independent public blocks, interference and continuous window

Condition	Sent / complete	Window seconds	Origin output token/s	First–content median / p95 s
Additional business/length block	4/4	62.3	108.44	1.67 / 1.68
Interference pair 4 Isolated	32/32	145.1	42.11	0.85 / 1.20
Interference pair 0 Isolated	32/32	151.2	42.99	0.97 / 1.58
Interference pair 4 Mixed	36/36	154.2	42.07	6.49 / 12.77
Additional business/length block	8/8	398.5	35.26	3.28 / 4.03
Interference pair 2 Isolated	32/32	171.3	35.39	0.90 / 1.15
Continuous run	599/599	3,600.2	64.77	1.01 / 1.49
Short–prompt C4 reference	60/60	196.4	117.65	0.90 / 1.15
Interference pair 2 Mixed	39/39	173.2	39.46	1.92 / 12.57
Short–prompt C4 reference	60/60	194.8	118.45	0.94 / 1.45
Additional business/length block	4/4	255.7	38.53	11.95 / 13.75
Short–prompt C4 reference	60/60	201.1	118.35	0.90 / 1.24
Interference pair 0 Mixed	33/33	159.3	39.01	4.32 / 13.88

Throughput requires original token counts for every complete response; open arrivals retain local rejection counts. Interference effects compare short requests within matched pairs; total mixed throughput is not a substitute for short–request latency differences.

Effect of long requests on short-request waiting

Each pair reuses the same short requests and arrival sequence; the mixed arm adds long requests. Waiting starts at the planned arrival, including client queues and confirmed retries. A ratio requires first-content observations for all 32 short requests in each arm; 1 means equal waiting, and larger values mean longer waits.

Table 09-09 Effect of long requests on short-request waiting

Pair	Complete/offered: isolated → mixed	Short-request median first-content s	Mixed / isolated ratio
1	32/32 → 25/32	0.98 → 1.40	—
3	32/32 → 31/32	0.91 → 2.16	—
5	32/32 → 28/32	0.87 → 5.17	—

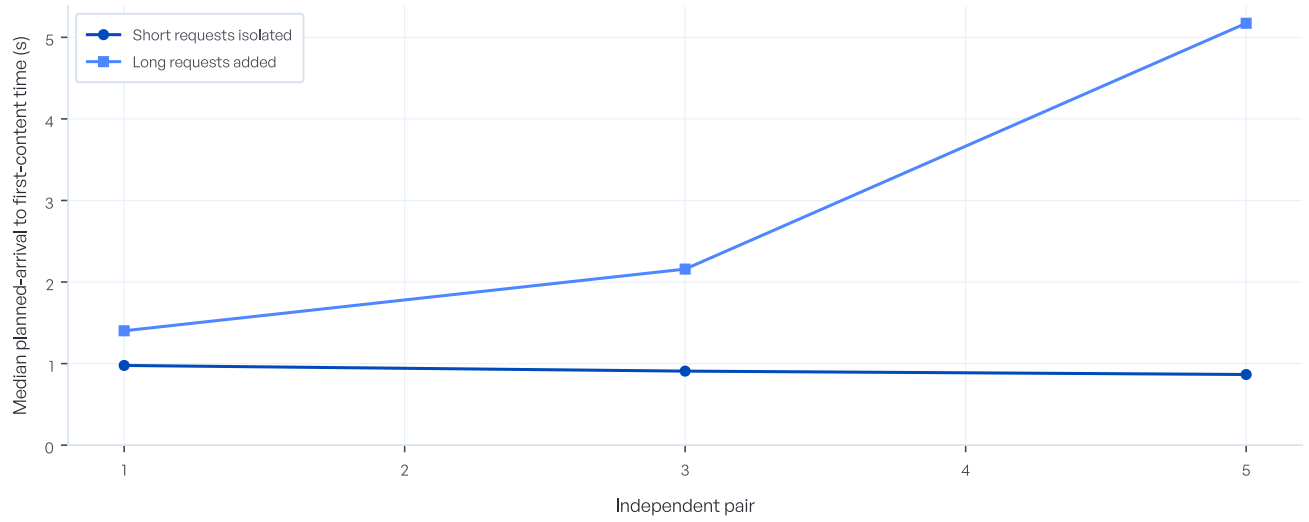


Figure 09-06 Paired short-request waiting

Samples and completeness remain explicit; ratios compare fully observed pairs only.

Five-minute windows of continuous service

Table 09-10 Five-minute windows of continuous service

Independent attempt / minutes	Offered / complete	First-content median / p95 s	Completed-response output token/s
1 / 0–5	55/54	1.03 / 1.67	71.63
1 / 5–10	45/45	0.99 / 1.27	48.54
1 / 10–15	50/50	1.08 / 1.63	67.85
1 / 15–20	59/59	0.99 / 1.56	76.47
1 / 20–25	58/58	1.07 / 1.65	69.37
1 / 25–30	40/40	0.98 / 1.36	55.57
1 / 30–35	55/55	0.97 / 1.38	72.35
1 / 35–40	51/51	1.08 / 1.43	65.93
1 / 40–45	55/55	1.04 / 1.49	67.71

Independent attempt / minutes	Offered / complete	First-content median / p95 s	Completed-response output token/s
1 / 45–50	33/33	1.04 / 1.65	50.93
1 / 50–55	53/53	1.14 / 1.35	70.85
1 / 55–60	46/46	1.12 / 1.58	60.13

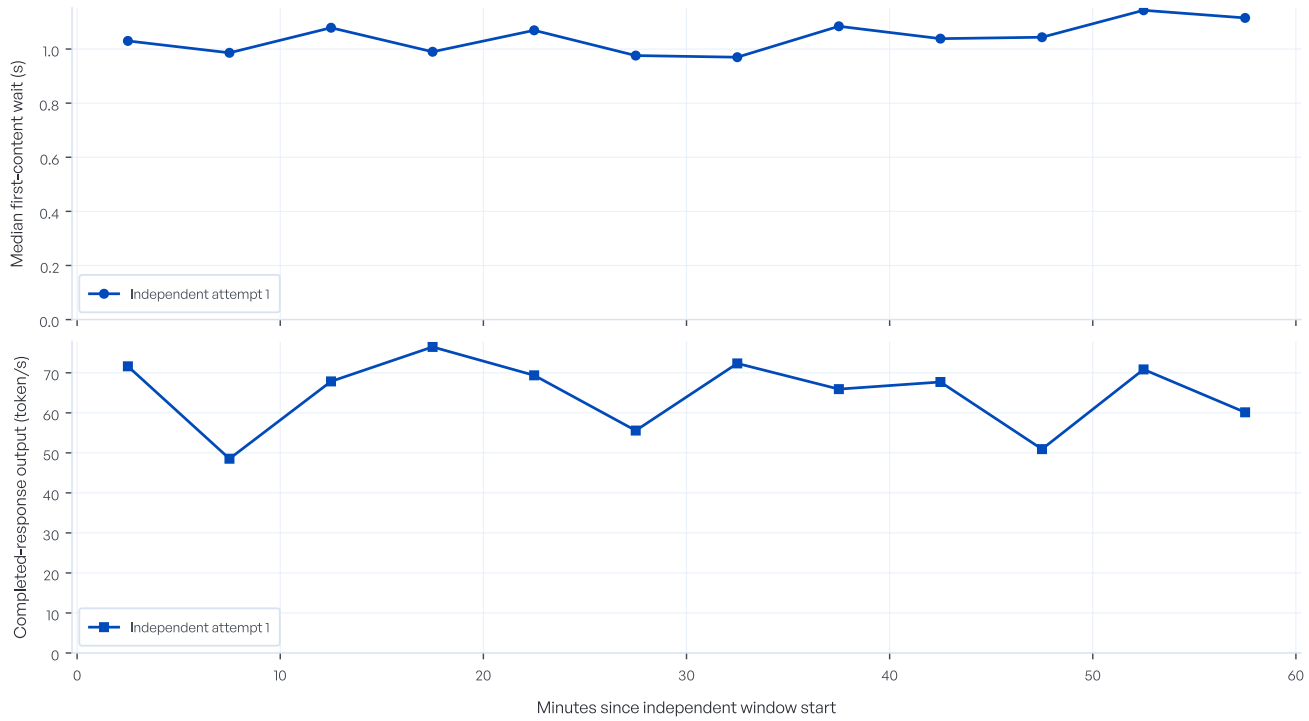


Figure 09–07 **Waiting and completed output in continuous windows**

Latency cohorts use planned arrival time. Tokens are credited when a response completes, not at their individual generation times. Original warmup flags remain; early windows include startup transients.

10 · Cross-layer interpretation and H200 comparison

The experiments establish resource ceilings, execution efficiency and user experience separately. This section connects them: multiple mechanisms can produce the same symptom, and matched-work controls narrow the explanation. The table separates established observations from mechanisms requiring validation.

Table 10-01 **Cross-layer interpretation and H200 comparison**

Dimension	Combined observations	Supported interpretation
Compute and shape	Matrix size/dtype sweeps, standalone operators and engine GPU traces.	Achieved rate depends on shape and implementation; peak TFLOP/s is not model token throughput. Missing pipeline counters prevent unique compute attribution.
Memory and coordination	8 GiB copy: about 223–227 GB/s;1 GiB AllReduce bus rate: about 21.81 GB/s.	Local access and inter-node synchronization have different resource budgets and byte accounting; their ratio is not a model speedup.
Concurrency and waiting	On short prompts,C=1→8 raises aggregate output 45.1→158.9 tokens/s while per-request token time increases.	Concurrency raises system output; user waiting must be checked alongside it. Neither unlimited scaling nor long-term capacity is established.
Cache and context	Public reuse diagnostics, node token counters and native immediate repeats are measured separately.	Repeated input, cached tokens and waiting benefit are distinct. Short reuse tests do not establish retention for arbitrary business traffic.
Temperature and retention	Compute, bandwidth and mixed loads separately complete60/60/120 minutes; driver reasons retain their sampled coverage.	Retention is assessable within observed segments. A temperature threshold, driver throttling and performance loss are distinct observations.

Interpretation distinguishes direct observations, matched comparisons, official specifications and economic scenarios. Only the first two are performance experiments here; ordering, sample-size and resource-state limits accompany comparisons. Hardware counters cover only a single-kernel diagnostic probe; representative hardware/model workloads remain unprofiled, preventing model roofline attribution.

GB10 and H200: an architectural reference aligned by resource function

The basic comparison is one GB10 in one DGX Spark versus one H200 SXM; multi-device topology is stated separately. Metrics align by function—holding data, supplying compute, executing arithmetic, exchanging shards and expanding across servers—rather than technology names. DGX Spark is a complete computer; H200 SXM is a server accelerator. Capacity ownership, power boundaries and fabric scope accompany every value.

Table 10-02 **GB10 and H200: an architectural reference aligned by resource function**

Technical function	DGX Spark / GB10	H200 SXM / stated server	Alignment and interpretation
Hold weights, KV and workspace	128 GB LPDDR5x shared CPU/GPU unified memory	141 GB HBM3e GPU-local memory	Similar nominal capacity; Spark CPU/system use shares the same pool. Two Sparks provide256 GB distributed capacity.
Supply local data to the GPU	273 GB/s LPDDR5x specified	4,800 GB/s HBM3e specified	Both supply local memory data; the 17.58× specification ratio is a bandwidth ratio, not inference speedup.

Technical function	DGX Spark / GB10	H200 SXM / stated server	Alignment and interpretation
Exchange data between CPU and GPU	Shared memory and on-chip NVLink-C2C; no unverified C2C peak is assigned.	PCIe Gen5:128 GB/s bidirectional aggregate; host memory and HBM are separate	Compare cooperation and migration needs; C2C is not the link between two Sparks.
Coordinate shards of one model	Two Sparks use RoCE; the active 200 Gb/s port specifies 25 GB/s in one direction.	H200 SXM uses NVLink:900 GB/s bidirectional aggregate per GPU, equivalent to 450 GB/s aggregate transmit direction.	Same function, different topology: one NIC port versus all GPU links. This is not a single-link or GPU-pair performance ratio.
Expand across more servers	One active 200 Gb/s port per node=25 GB/s; logical interfaces are not additive.	DGX H200 example:400 Gb/s=50 GB/s per NIC; eight NICs in an eight-GPU system total 3.2 Tb/s.	Per-NIC specified rate is 2×; system aggregation depends on topology, concurrency and fabric.
Low-precision matrix acceleration	Blackwell supports FP8 and NVFP4; fixed 8192^3 dense measurements appear below.	Hopper native FP8:about 1,979 dense and 3,958 with2:4 sparsity TFLOP/s (specifications).	FP4 and FP8 serve low-precision acceleration, but formats, kernels and errors differ; their rates are not a speedup ratio.
Power and provisioning	GB10 SoC TDP140 W; Spark power supply rated 240 W.	H200 SXM configurable TDP up to 700 W; host-system power is additional.	Design specifications with different boundaries; neither establishes tokens/J or wall-electricity cost.

DGX Spark · H200 SXM · Spark ConnectX-7 · DGX H200 · Hopper NVLink

All H200 values are official specifications; H200 was not benchmarked in this study. GB10 measurements below use unloaded-model conditions. Functional alignment supports architectural comparison; inference performance ratios require matched model, precision, work and timing on both platforms.

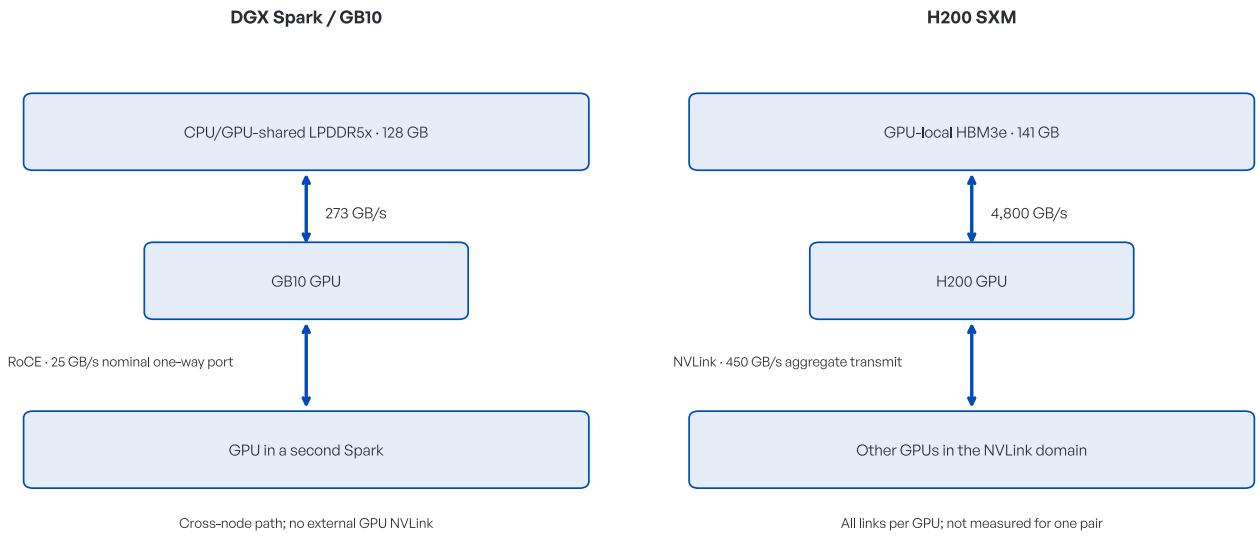


Figure 10-01 Data-path comparison

Both systems supply local data and exchange model shards through different memory and interconnects. Lines denote functional paths and are not scaled by bandwidth.

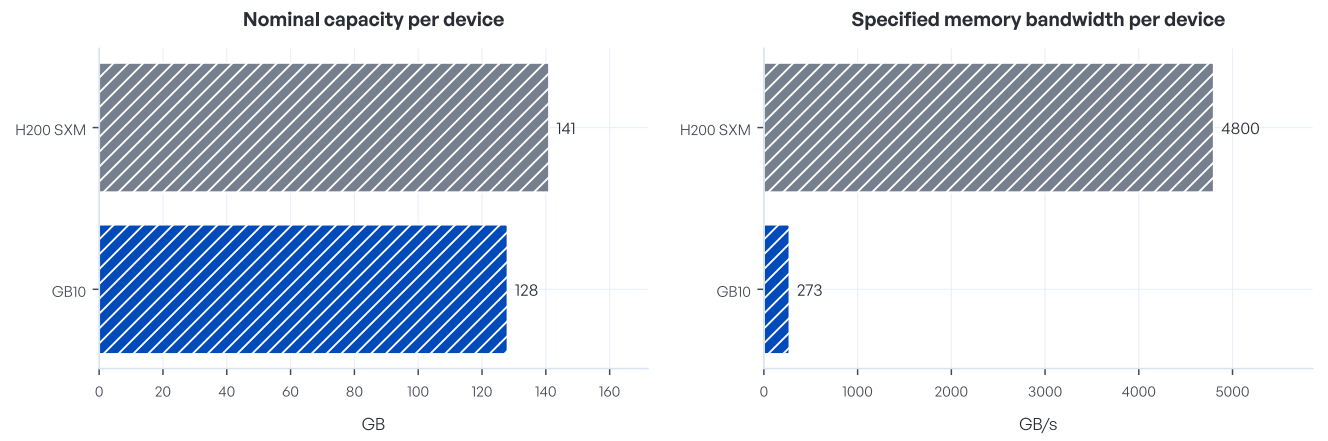


Figure 10–02 Capacity and data supply

Per–device official specifications. Similar capacity does not imply similar bandwidth; both panels begin at zero.

Measured memory rates: fixed 8 GiB working set

Table 10–03 Measured memory rates: fixed 8 GiB working set

Operation	Node 0 GB/s: median [range]; rounds	Node 1 GB/s: median [range]; rounds	H200 HBM3e reference
Contiguous copy: read + write	**223.41** [222.75–223.74]; n=10	**226.95** [226.31–227.88]; n=10	4,800 GB/s specified; this operation unmeasured
Triad: two reads + one write	**214.88** [213.76–215.27]; n=10	**217.66** [217.07–218.13]; n=10	4,800 GB/s specified; this operation unmeasured
Native read	**110.04** [109.88–110.21]; n=5	**110.09** [108.89–110.79]; n=5	4,800 GB/s specified; this operation unmeasured
Native write	**192.63** [191.81–193.07]; n=5	**192.87** [191.51–193.06]; n=5	4,800 GB/s specified; this operation unmeasured

Working sets use binary GiB; rates use decimal GB/s. Values count logical bytes requested by the program, not counter–measured physical DRAM traffic. Native read includes dependent per–thread accumulation and describes that kernel rather than a theoretical read limit. Triad counts actual element–aligned bytes.

Core compute throughput by precision

TFLOP/s counts 10^{12} floating–point operations per second, with FMA counting as two. GEMM fixes M=N=K=8192 and counts 2MNK/time rather than selecting the fastest size. Scalar FP32 FMA is separate. H200 dense Tensor Core figures are approximate halves of the rounded published sparse figures, not measurements. In 2:4 sparsity two of every four weights are nonzero; sparse rates use dense–equivalent work.

Status key: **No native support** means the hardware lacks a native arithmetic path for that format; **Not applicable** means the row does not use that column's dense or 2:4 sparse convention; **Not measured** means this study has no valid measurement yet. None means zero throughput.

Table 10–04 Core compute throughput by precision

Operation / precision	GB10 node 0 measured TFLOP/s	GB10 node 1 measured TFLOP/s	H200 dense specified TFLOP/s	H200 2:4 specified TFLOP/s
Scalar FP32 FMA	28.01 [27.97–28.04]; n=5	27.98 [27.94–28.06]; n=5	67	Not applicable
FP32 GEMM	18.39 [18.16–18.48]; n=10	18.22 [17.94–18.30]; n=10	67	Not applicable
TF32 allowed	33.10 [24.01–38.59]; n=10	37.17 [33.21–38.53]; n=10	≈494.5	989
FP16	93.24 [91.83–93.95]; n=10	92.67 [91.59–92.99]; n=10	≈989.5	1,979
BF16	94.34 [93.50–94.78]; n=10	94.08 [93.63–94.42]; n=10	≈989.5	1,979
FP8 E4M3	192.64 [192.07–195.62]; n=5	191.88 [191.38–193.22]; n=5	≈1,979	3,958
NVFP4	157.38 [148.49–160.44]; n=5	148.02 [143.42–161.21]; n=5	No native support	No native support
FP16 2:4	39.33 [39.25–39.45]; n=5	39.48 [39.38–39.65]; n=5	Not applicable	1,979
BF16 2:4	39.28 [39.03–39.69]; n=5	39.41 [38.82–40.02]; n=5	Not applicable	1,979
FP64	0.4280 [0.4280–0.4281]; n=5	0.4276 [0.4276–0.4277]; n=5	34 scalar; 67 Tensor Core	—

Scalar GB10 FMA uses 256 threads/block, eight independent chains and four blocks per SM. TF32 allowed permits library selection; the flag alone does not establish executed instructions. NVFP4 uses controlled exactly representable operands, not production quantization errors or MoE kernels. GB10's advertised up–to–1 PFLOP/s uses sparse FP4 and is not the denominator for a dense NVFP4 attainment rate.

NVFP4 architecture difference: GB10 is Blackwell and supports native FP4 matrix arithmetic; H200 is Hopper and lacks native NVFP4 Tensor Core arithmetic, so no native dense or sparse peak applies. Compatible software kernels can dequantize FP4 weights to BF16 for computation, but that is a different arithmetic path; H200 FP8 or INT8 specifications cannot fill these cells. This software path was not benchmarked on H200 in this study. NVIDIA NVFP4 · NVIDIA NIM NVFP4 Emulation

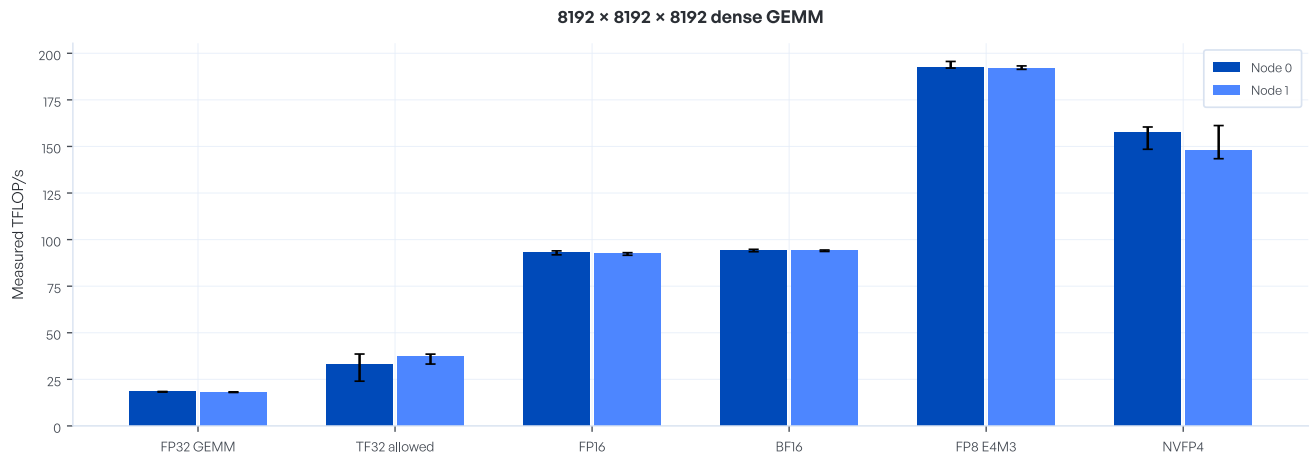


Figure 10–03 Fixed–size compute measurements

Only GB10 measurements are plotted, with observed round ranges. H200 specifications remain in the table and are not presented as a matched benchmark.

Implications for deployment

The measured **8 GiB** copy reaches **223–227 GB/s**, while two-node **1 GiB** AllReduce bus bandwidth is about **21.81 GB**/s. Cross-node synchronization and local memory access operate at different rate scales. H200 specifications offer higher memory and GPU-interconnect ceilings, but benefits depend on sharding, arithmetic intensity, batching and overlap. These data identify resource differences; they do not establish H200 token throughput, latency, energy-efficiency or payback ratios. A matched comparison requires the same model, quantization/precision, input/output, concurrency and cache conditions, plus wall-power measurement.

11 · API cost, ownership and investment return

Observed API charges

Charges are reconciled by the same cohort's request IDs. Blended cost per million generated tokens includes input charges for this input/output mix; it is not a catalog output-token price.

Table 11-01 Observed API charges

Concurrency condition	Bill coverage	Paid USD	USD/1k responses	USD/1M generated
Concurrency 1 · Reasoning disabled	120/120	0.0244	0.204	0.519
Concurrency 2 · Reasoning disabled	180/180	0.0375	0.208	0.517
Concurrency 4 · Reasoning disabled	60/60	0.0124	0.206	0.518
Concurrency 8 · Reasoning disabled	180/180	0.0373	0.207	0.517

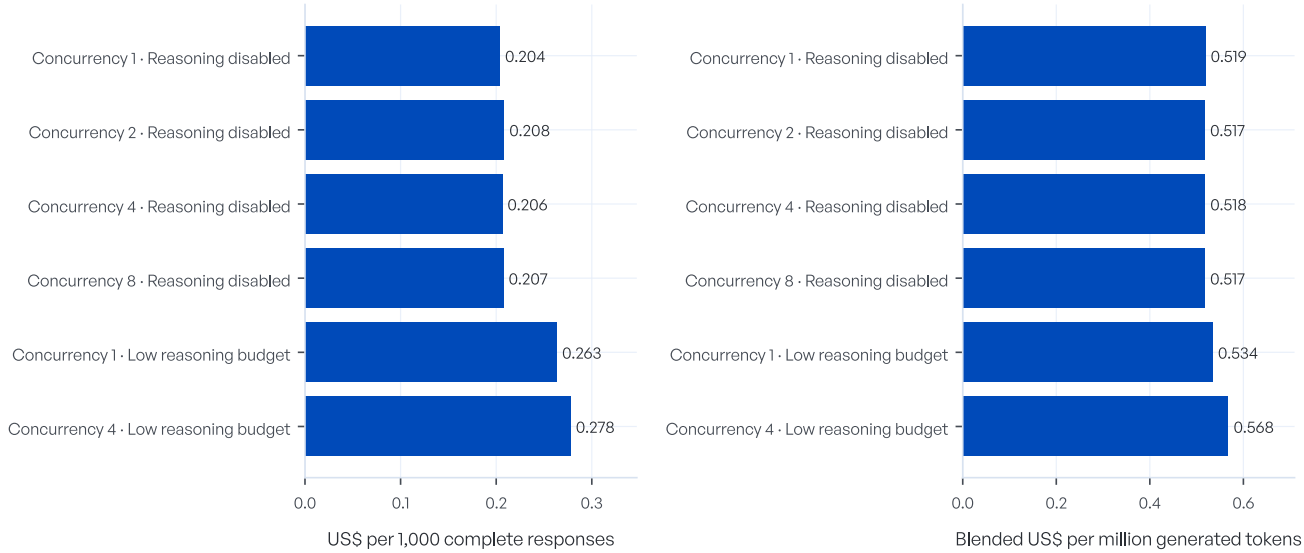


Figure 11-01 Cost normalization

Charges per response and per generated token both include cohort input fees.

Two-node ownership: purchase and electricity

The supplied purchase price is **US$5,000** per device, **US$10,000** for the pair. Electricity uses the public China business reference **US$0.109/kWh (2026 Q3)** as a reproducible proxy, not a site contract or an official consumption-weighted national average. GlobalPetrolPrices

Unit cost uses straight-line depreciation and estimated electricity: annual expense = purchase price / depreciation life + pair kW × 8,760 hours × tariff, divided by useful annual generation. Power uses the matched-condition estimated system sum, assumed constant year-round. Productive fraction is time producing useful work, not GPU activity. Staff, space, cooling, repairs and financing are excluded.

Table 11-02 Two-node ownership: purchase and electricity

Concurrency	Estimated pair W	Energy USD/year	50% · USD/1M	75% · USD/1M	100% · USD/1M
1	112.3	107	4.833	3.222	2.417
2	120.6	115	2.778	1.852	1.389
4	122.9	117	1.821	1.214	0.911
8	124.6	119	1.378	0.919	0.689

The table uses three-year depreciation; the final columns give cost per million generated tokens at each productive fraction.

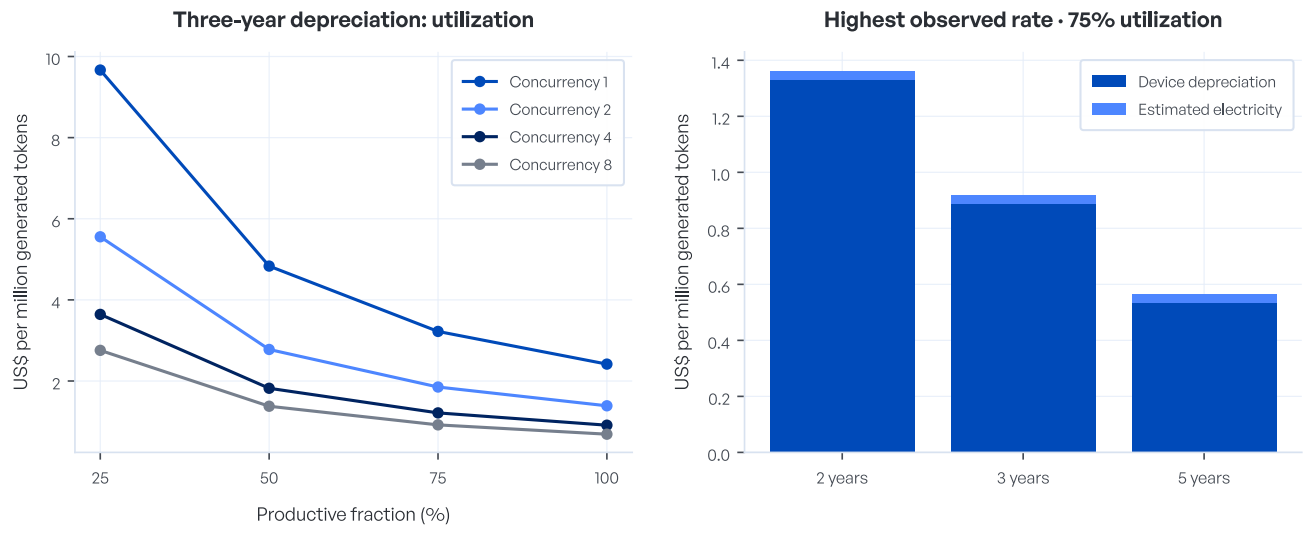


Figure 11-02 Cost sensitivity

Left: utilization with three-year depreciation. Right: depreciation life at 75% utilization.

1 / 3 / 5-year payback and investment returns: self-use and sales

This section asks when cash savings or sales receipts cover the initial two-device investment. Depreciation is used above to allocate unit costs; cash-flow analysis instead deducts the purchase once at acquisition and then only cash operating expenses. Depreciation is not deducted again. All amounts are USD. These are conditional projections from measured throughput, not realized multiyear returns.

Table 11-03 1 / 3 / 5-year payback and investment returns: self-use and sales

Use	Definition of annual benefit	Price and work-volume basis
Self-use: avoided external API spending	API expense that the same work would otherwise incur	Same input/output mix and account-paid blended price; only work actually needed counts
Sales: inference-service receipts	Sold volume × blended sale price, less sales deductions	Illustrative 0.5 / 1 / 2× reference price; sale utilization is time producing collectible work

Fixed projection basis. Natural-bank concurrency 8, observed throughput **158.90 tokens/s**, estimated pair power **124.57 W**, 8,760 provisioned hours/year and **US$10,000** purchase cost. The corresponding account-paid blended API price is **US$0.51739 per million generated tokens**, including input fees for this mix. Sales scenarios use it only as a reference for hypothetical quotes; account spending is not node revenue or a general output-token price.

Table 11-04 1 / 3 / 5-year payback and investment returns: self-use and sales

Metric	Definition and calculation
Annual useful output Q	Q = R × 3,600 × H × u / 1,000,000, in million generated tokens
Annual cash operating expense O	O = pair watts/1,000 × provisioned hours × tariff + other annual expense; sales deductions separate
Annual net cash benefit S	Self-use: S = Q × API reference price − O; sales: S = Q × sale price × (1 − deduction rate) − O
Cumulative net cash at year T	T × S − K, where K is initial capital; negative means capital is not fully recovered
Cumulative simple ROI	(T × S − K) / K; denominator is initial capital, not revenue; this is not annualized
Simple payback	K / S when S > 0; a value over five years means no payback within five years
Minimum blended price for T-year payback	(K/T + O) / [Q × (1 − deduction rate)], covering cash costs included here

Fixed scenarios hold throughput, prices and tariffs constant and ignore the time value of money. They exclude residual proceeds, replacement, financing, taxes, staff, cooling, space and other operations. Baseline sales deductions are set to zero solely to illustrate a ceiling before channel and similar fees; zero is not an observed fee. Each use assumes sustained demand for the stated volume; effective utilization is not GPU activity. Five-year results also require five years of usable device and service life.

Self-use: 1 / 3 / 5-year cash account at 75% productive utilization

Annual useful generation is approximately **3,758.2 million tokens**, avoiding **US$1,944** of API expense, with estimated electricity **US$119** and net annual savings **US$1**,826.

Table 11-05 Self-use: 1 / 3 / 5-year cash account at 75% productive utilization

Holding period	Avoided API expense	Cumulative electricity	Initial purchase	Net cumulative cash	Cumulative ROI
1 year	US$1,944	US$119	US$10,000	US$-8,174	-81.7%
3 years	US$5,833	US$357	US$10,000	US$-4,523	-45.2%
5 years	US$9,722	US$595	US$10,000	US$-872	-8.7%

Utilization sensitivity: with the same purchase and year-round energy assumptions, only useful work creates avoided expenditure.

Table 11-06 Self-use: 1 / 3 / 5-year cash account at 75% productive utilization

Productive fraction	Annual net saving	Simple payback	1 year ROI	3 year ROI	5 year ROI
25%	US$529	18.90 years	-94.7%	-84.1%	-73.5%
50%	US$1,177	8.49 years	-88.2%	-64.7%	-41.1%
75%	US$1,826	5.48 years	-81.7%	-45.2%	-8.7%

Productive fraction	Annual net saving	Simple payback	1 year ROI	3 year ROI	5 year ROI
100%	US$2,474	4.04 years	−75.3%	−25.8%	+23.7%

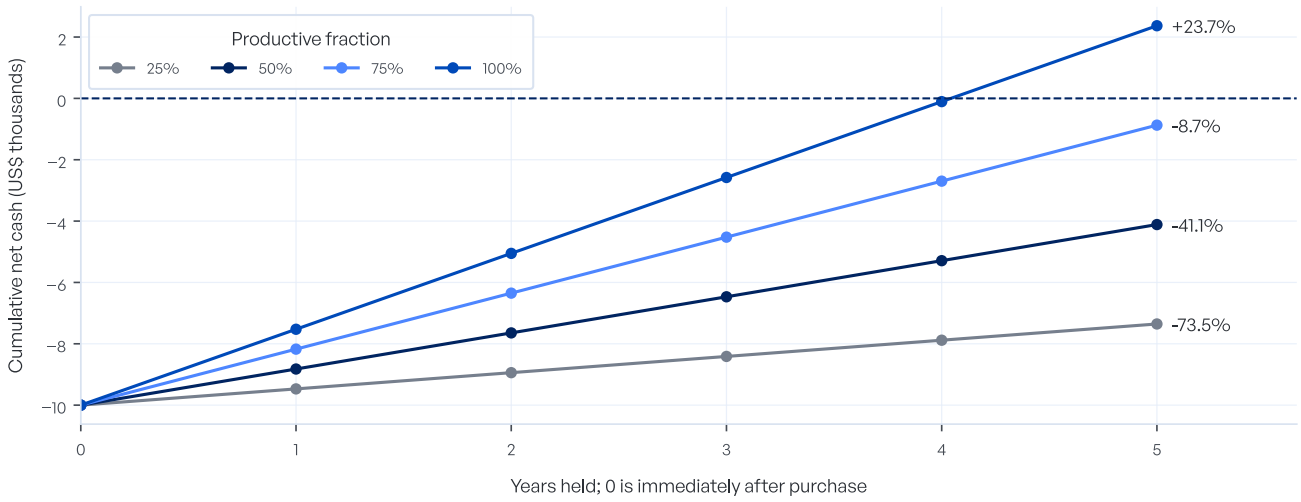


Figure 11–03 **Self–use cumulative cash and payback**

Curves start at the US$10,000 purchase outlay; crossing zero indicates full capital recovery. Endpoint percentages are year–five cumulative ROI. Only productive utilization varies. The curves extrapolate a fixed short–run measured rate; they are not multiyear observations.

Self–use finding. At **75%** productive utilization, simple payback is **5.48 years** and five–year cumulative ROI **−8.7%**. At **100%**, payback is **4.04 years** and five–year ROI **+23.7%**. Payback beyond five years is mathematical extrapolation conditional on unchanged assumptions, not evidence of that device lifetime.

External sales: price and sold utilization jointly determine return

The table fixes sold utilization at **75%** and varies the quote to 0.5 / 1 / 2 times the reference. Annual net cash here is gross sales less estimated electricity; omitted expenses must be added for actual operations. At equal price and volume with no sales deductions, sales and self–use have equal arithmetic, but sales require paying demand while self–use requires internal demand. Their benefits must not be added together.

Table 11–07 **External sales: price and sold utilization jointly determine return**

Sale price USD/1M generated	Annual gross sales	Annual net cash	Simple payback	1 year ROI	3 year ROI	5 year ROI
0.259 (0.5×)	US$972	US$853	11.72 years	−91.5%	−74.4%	−57.3%
0.517 (1×)	US$1,944	US$1,826	5.48 years	−81.7%	−45.2%	−8.7%
1.035 (2×)	US$3,889	US$3,770	2.65 years	−62.3%	+13.1%	+88.5%

Sales scenario finding. At **75%** sold utilization, selling at the reference price still gives **5.48 years** payback. At twice that price (**US$1.035**/million generated tokens), payback is about **2.65 years**, with cumulative net cash at years 1 / 3 / 5 of **US$−6,230** / **US$1,310** / **US$8**,850. Return depends on both price and sustained sales volume; a higher quote alone does not establish demand at that volume.

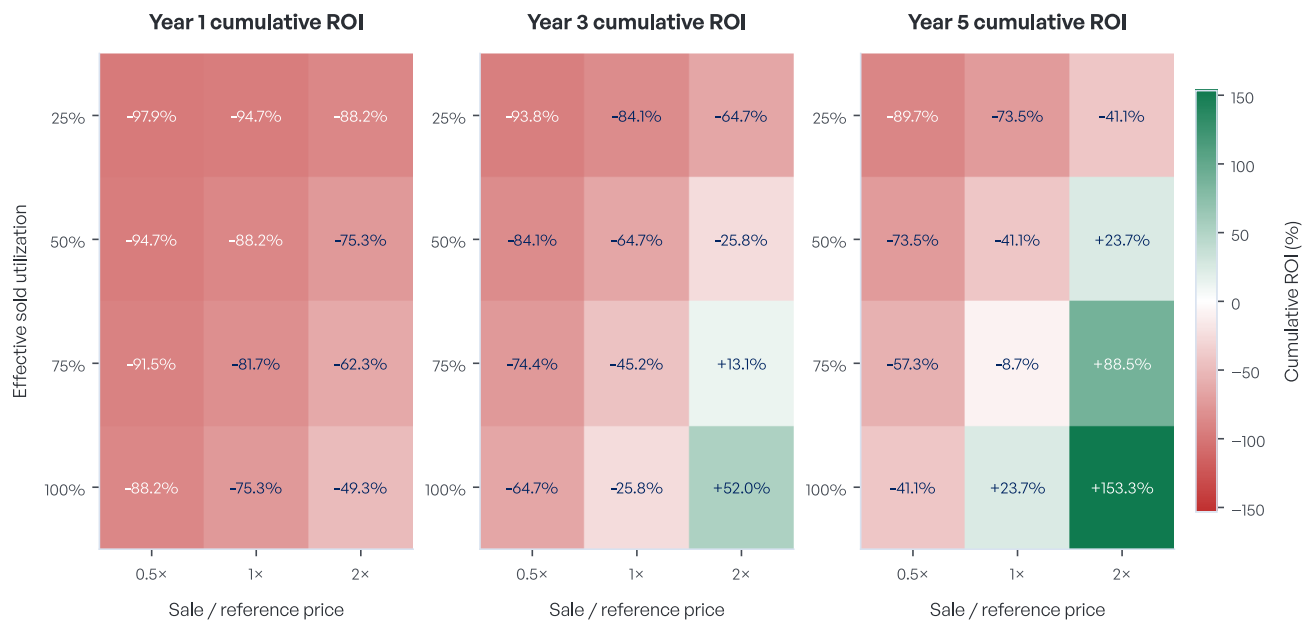


Figure 11–04 Sales ROI sensitivity to price and sold utilization

Panels show cumulative simple ROI at years 1, 3 and 5; negative percentages mean capital is not fully recovered. All panels share one color scale. Prices include the fixed input mix; sales deductions and other operating expenses are excluded.

Blended sale price required to recover capital within a chosen horizon:

Table 11–08 External sales: price and sold utilization jointly determine return

Sold utilization	1–year payback price	3–year payback price	5–year payback price
25%	US$8.078	US$2.756	US$1.691
50%	US$4.039	US$1.378	US$0.846
75%	US$2.693	US$0.919	US$0.564
100%	US$2.019	US$0.689	US$0.423

Prices are blended USD per million actual generated tokens at an unchanged input/output mix. Changes in long–input share, reasoning length or cache conditions require replacing both throughput and blended price with matched–workload values. Minimum prices cover only included costs and are not commercial quote recommendations.

Download all return scenarios and calculation inputs

Reconciled billing by workload · All ownership–cost scenarios

Actual API charges for the additional experiments

Table 11–09 Actual API charges for the additional experiments

Scope	USD
Historical reconciled charges	8.561299
Independent additional batch	0.376474
Cumulative reconciled charges	8.937773

Additional charges are not inserted into earlier workload unit costs or investment returns; those scenarios retain their original matched charges and throughput.

12 · Overall conclusions and next work

This evaluation establishes a layered baseline for two DGX Spark nodes under the tested model, engine and workloads. Deployment decisions require capacity, waiting, throughput and cost to be considered together; each conclusion below maps to reproducible preceding evidence.

Table 12-01 Overall conclusions and next work

Dimension	Finding	Applicability
Data supply and communication	8 GiB contiguous copy223—227 GB/s;1 GiB AllReduce about 21.81 GB/s.	A resource baseline under matching conditions, not a token-throughput prediction.
Engine stages	At C=1 and 256fixed output tokens, median decode rates span 49.8—60.2 tokens/s across tested input lengths.	Native synthetic loads; stage timing includes engine scheduling boundaries.
Public concurrency	C=8 aggregates 158.9 tokens/s with visible-content p95 of 2.39s; throughput is about 3.52×C=1.	The 60-prompt short-input bank and eligible complete blocks; not a service guarantee.
Long input and output	At C=4 and output cap512, input512→32,768 changes median visible wait1.48→37.02s; the largest achieved target output in the extension is 4,096.	256K is the configured boundary; actual256K output is not established.
Prefix reuse	Token-weighted cache fraction is 20.3% in eligible active node windows; native reuse is reported by input length.	Mixed-window statistics are not per-request hit rates or evidence of long retention.
Sustained and thermal behavior	Both nodes completed 60 minutes each of compute and bandwidth; the common mixed-load coverage is 120.0 minutes.	Segmented evidence, not four-hour continuous reliability validation.
H200 technical reference	Per-device capacities are similar; H200 specifies higher local-memory bandwidth, with different cooperation paths and power boundaries.	No matched H200 run supports an inference value-for-money ranking.

Table 12-02 Overall conclusions and next work

Economic use	Scenario result	Key conditions
Self-use economics	At 75% productive utilization, simple payback 5.48 years and five-year ROI −8.7%.	Conditional on steady demand and price; operating expenses reduce return.
Sales economics	At 75% sold utilization and twice the blended API reference price: payback 2.65 years and five-year ROI +88.5%.	A scenario with no channel deductions or other operating costs, not realized revenue.

Limitations and future work

Table 12-03 Limitations and future work

Priority	Open uncertainty	Next experiment and assessment
1	Sustained behavior across periods	Repeat sustained work on independent dates and ambient conditions to assess reproducibility of retention and cooling frequency.
2	Realized long-output range	Verify longer outputs under controlled generation; separate natural stopping from fixed length rather than treating caps as realized work.

Priority	Open uncertainty	Next experiment and assessment
3	Retention and service objectives	Hold a representative traffic distribution fixed, add independent periods and competing requests, and link per-request reuse to waiting and goodput.
4	Hardware mechanism and matched H200 benefit	Compare identical model, precision, lengths, concurrency and cache states; add counters, verified transport and same-trace analysis.
5	Real ownership cost	Replace scenarios with wall-meter energy, actual utilization, purchase quotes and operating expenses before recomputing1/3/5-year returns.

Final coverage and priorities for further measurement

Coverage boundary: execution and archiving have ended. The public-length list retains eight unsent requests and one incomplete response. Three of five candidate interference pairs ran, with zero fully observed pairs. Hardware, native long output, profiling, immediate cache reuse and the continuous hour retain their completed lists; these do not establish validation of every proposed condition.

Prioritize missing public-length combinations and redesign interference observation windows to contain the fixed arrival sequence, retaining overload rejection as a service outcome. Any new experiments require independent blocks and a new production resource arrangement; current non-admitted or incomplete pairs remain unsuccessful evidence.

Additional native-engine result: fixed output has been validated through **260,096 tokens**. Public natural-stop data remain a separate cohort; native results do not replace end-to-end measurements.

An independent full hour extends the sustained-service evidence, with five-minute windows showing waiting and completed-output variation. This extends coverage at the tested fixed rate; it does not establish a full-day service level, peak capacity or annual availability.

13 · Appendix A: complete metric and statistical definitions

This appendix collects metric formulas, data sources and interpretation limits. Seconds are s, milliseconds ms, and token/s denotes tokens per second. Each estimate requires a population, warmup rule, time window and missing-value count. “Collected” means the record type exists, not that every request or the entire experiment has a valid value.

The method follows the distinctions between first response, request latency, output throughput and concurrency/rate in NVIDIA GenAI-Perf and SGLang serving benchmarks. This study uses its own collector and explicitly defines observable fragment boundaries; it does not claim official certification by these tools. References use archived pre-experiment documentation; rolling docs do not prove all behavior of the installed version. NVIDIA GenAI-Perf · SGLang bench_serving · Archived references and hashes

Formula notation: t denotes an event timestamp; subtracting timestamps gives duration. N_in and N_out denote input and output token counts; cap means a ceiling and target a planned value. Σ means summing over the stated samples. Ratios require matching valid samples in numerator and denominator.

Request timing and streaming experience

Table 13-01 Request timing and streaming experience

Metric	Definition, units and limits	Source / availability
Response-header wait	t_headers − t_send; seconds. Includes connection, service waiting and transport.	headers_s; collected per request
First generated-fragment latency	t_generated − t_send; first reasoning, content or tool fragment. Often called TTFT, but not an exact token event.	first_generated_s; collected
First visible-content latency	t_visible − t_send; counts visible content only; missing if no visible content.	first_visible_s; collected
First reasoning-fragment latency	t_reasoning − t_send; missing rather than zero when no reasoning appears.	first_reasoning_s; collected when present
End-to-end latency	t_end − t_send, including stream closure. Open arrivals additionally count pre-send waiting for latency from scheduled arrival.	e2e_s; collected
Estimated time per output token	(E2E − first-generated latency)/(N_out − 1), N_out>1; ms/token. Uses upstream output counts; an approximation.	tpot_compat_s × 1,000; derived
Fragment span and terminal delay	(t_last − t_generated)/(N_out−1) is an alternative decode estimate; t_end−t_last is terminal delay. First-fragment token count is unknown.	decode_tpot_approx_s / terminal_delay_s; recorded/derived
Streaming gaps and stalls	Median, p95 and maximum gaps between nonempty generated events, with counts/durations above 200 ms and 1 s. Not exact inter-token latency.	events.jsonl / chunk_intervals; collected

Tokens, length and caching

Table 13-02 Tokens, length and caching

Metric	Definition, units and limits	Source / availability
Input and output counts	Input includes the chat template; output uses upstream generation counts. API, upstream and billing-normalized counts retain separate provenance.	prompt_tokens / completion_tokens / api_completion_tokens; collected

Metric	Definition, units and limits	Source / availability
Reasoning-token counts	Use explicit fields; do not reconstruct by subtracting incompatible total and visible counts.	reasoning_tokens; reported when available
Planned and realized input/output ratios	Planned=N_in,target / N_out,cap; realized=N_in,observed / N_out,observed. Missing for zero or unavailable output.	Plan manifest and requests.csv; derived
Output-limit fraction	Complete responses with finish_reason=length / complete responses; not an accuracy-failure rate.	truncated / finish_reason; collected
Shared-prefix fraction	Longest common token prefix of complete templates / measured input tokens; separate from the target.	Business manifest matched_prefix_tokens; verified offline
Cached-token fraction	Σ reported cached tokens / Σ corresponding input tokens for valid-field samples, with coverage; per-request distributions are separate.	cached_tokens; requires explicit fields, missing is not zero
Requests with any cache hit	Requests with cached tokens>0 / requests with valid cache fields; distinct from the token-weighted fraction.	Derived from valid cache fields; not a design target

Throughput, arrivals and reliability

Table 13-03 Throughput, arrivals and reliability

Metric	Definition, units and limits	Source / availability
Aggregate block output throughput	Σ output tokens / (block drain end−start), token/s; multiply by 60 for token/min. Not the sum of per-request rates.	Main figure uses complete uninterrupted blocks; full-wall values retained
Request throughput and offered rate	Completed requests / window seconds; scheduled arrivals / window seconds; sent and completed counts remain distinct.	Block receipts and arrivals; latter generated in open-arrival stages
Concurrency, queue and scheduler lag	active counts in-flight requests; client_queue_s measures local waiting; scheduler_delay_s measures planned-to-enqueued delay.	Client monitor; open-queue fields require that stage
Goodput and qualified fraction	Complete, integrity-unflagged requests satisfying both visible-TTFT and estimated-TPOT thresholds / first measured arrival to drain. Qualification fraction divides by all measured offered requests, including local rejections.	Thresholds 1/2/5/10 s × 20/50/100 ms; uses open-stage data
Completion and failure denominators	Complete attempts / terminal attempts; logical first-attempt and eventual completion are separate. In-flight work is not prematurely counted as failure; failed retries are retained.	state / retry_summary / HTTP status; collected
Integrity flags and boundary rejection	Empty content or repetitive fragments are candidate flags. Intended context-limit HTTP 400 rejection is separate from normal generation completion.	integrity_flags / expected_context_rejection; stage-specific
Cost per response and generated output	Same eligible blocks: paid charges / complete responses × 1,000; paid charges / original generated tokens × 1,000,000. Both include input charges.	Exact per-request bill join; missing bills are not zero
Device and electricity scenarios	Purchase cost / useful life + mean kW × provisioned hours/year × tariff, divided by projected generation. Productive fraction is not GPU activity.	User price, public tariff benchmark, estimated power and explicit assumptions; not measured full TCO
Cost and budget reservation	Actual cost is reconciled by request; reservation is a conservative ceiling, not paid cost or hardware total cost of ownership.	billing_latest.json / ledger.json; collected

Inference nodes and mechanisms

Table 13–04 Inference nodes and mechanisms

Metric	Definition, units and limits	Source / availability
GPU activity, clock and temperature	% / MHz / °C by node and load window. GPU activity is not FLOP efficiency or memory–bandwidth utilization.	SparkDash; collected with possible gaps
Power and throttling	GPU sensor watts; systemDraw is estimated system power; thermal/ hwSlowdown identify throttling. No outlet–level tokens/J are derived.	SparkDash; collected
Unified memory and swapping	Available RAM, page cache and used swap; MiB=2^{20} bytes, GiB=2^{30}. Overlapping CPU/GPU memory views must not be added.	SparkDash memory fields; meminfo restricted to sampled windows
KV cache and state slots	Used, available and evictable tokens/slots. Evictable is not immediately free; pools with different units cannot be combined.	Available SparkDash fields; raw SGLang metrics have incomplete coverage
Engine running and queue counts	num_running_reqs / num_queue_reqs; inspect accumulation by window. Missing segments are not filled with zero.	SparkDash snapshots and raw timelines; analyze valid windows
Engine histograms and counters	Boundary deltas for requests/input/generation/cache; histogram quantiles stay within supported bucket bounds. Resets/disappearing series become missing; parallel replicas are not blindly summed.	Incomplete Prometheus archive; raw histograms and counters unavailable through SparkDash
Prediction acceptance	Window acceptance fraction/length follow engine–defined denominators; missing accepted–token counters prevent whole–run acceptance or speedup inference.	Archived raw gauges; SparkDash acceptance fields are null
Memory pressure and faults	PSI measures time stalled on resources; swap–in/out and major faults use valid–window counter deltas per second.	Incomplete vmstat / PSI archive; valid windows only
Errors and fabric	Xid tracks driver errors and OOM counts out–of–memory kills; use changes from baseline. Observe available RoCE traffic/error fields. Isolated RDMA microbenchmarks are not model communication measurements.	Raw SparkDash node snapshots; separate RDMA results

Measurement quality and statistics

Table 13–05 Measurement quality and statistics

Metric	Definition, units and limits	Source / availability
Telemetry coverage and collection cost	Valid sample counts, intervals, last successful sample and collection duration determine which workload windows are alignable.	monitor/status, timestamps and collection_seconds; collected
Client resources and event–loop lag	CPU user/system time, peak resident memory and heartbeat lag help assess client saturation as an alternative explanation.	monitor/client.jsonl; collected
Median, tail quantiles and intervals	p50 is the median; about 95% of observations are at or below p95. A 95% confidence interval for a mean difference uses cluster resampling, not the interval containing 95% of requests.	10,000 seeded bootstrap draws; no robust p99 claim for small samples

Remaining gaps include exact client inter–token intervals, complete engine counters across every stage, wall–meter energy and physical DRAM transactions. Independent collectives and logical memory bandwidth are measured in sections 06—07; they do not directly establish communication cost within a model request or physical bus utilization.

Table 13–06 Measurement quality and statistics

Cache metric	Definition and interpretation
Input–prefix overlap	Longest common token prefix over the full chat template / probe input length. A controlled input property.
Per–request token–cache fraction	Sum of cached input tokens / sum of input tokens for requests with valid fields, accompanied by field coverage / complete responses.
Fraction of requests with a hit	Requests with cached tokens >0 / requests with valid cache fields. A small template hit counts; it does not imply full–input reuse.
Window token–cache fraction	Endpoint cached–counter delta / prompt–counter delta; check freshness, resets, identity and receipt agreement. Count the head endpoint once.
Paired waiting change	Seed visible TTFT minus repeat TTFT; positive means the repeat was faster. Relative changes are separate. Network, order and load may contribute; this is not isolated causal acceleration.
Retention, competition and recovery	Compare cache fractions by actual idle gap; compare post–pressure with baseline for cache loss and extra waiting, then examine recovery. Findings apply only to tested pressure and intervals.
Temporal stability and coverage	Report valid / all windows, observed span, token–weighted mean, range and SD for at least three valid windows. Missing or reset intervals are neither zero–filled nor connected.

Units, inclusion rules and denominators

Primary units are requests and continuous blocks at fixed workload settings. Capacity curves include only blocks with all planned requests completed, no interruption segments or extra client attempts, and usable output counts; latency and throughput use the same blocks. All attempts remain available for completion and cost analysis. Complete–block selection may favor better operating periods, so these rates describe completed conditions rather than service availability. Warmups are excluded from formal business and open–arrival estimates; cache warmups retain pair identity and timing.

Replication, pairing and uncertainty

The 60 natural prompts repeat across conditions; request count is not independent semantic sample count. Reasoning comparisons match prompt, round and concurrency, with low minus none defining the effect. Ten thousand seeded bootstrap resamples draw whole prompt or round clusters to retain dependence. With only three rounds, round–level intervals are unstable; narrow intervals do not establish broad applicability. Comparisons are exploratory, with no multiple–comparison significance claim.

Business cells contain only 12 measured calls and cache cells six pairs. Report sample counts, medians and ranges rather than robust p99 or population–wide claims. Business concurrency conditions use matched material and distinct prefix identifiers. Three blocks provide limited replication, not precise population estimates. When generation stops early, realized output, truncation fraction and throughput must be read together.

Missingness, interim results and causal limits

API, original upstream and billing counts retain distinct provenance. Token rates for low reasoning remain unavailable without original upstream counts; missing cache data is not zero and unmeasured conditions are not interpolated. Telemetry is aligned by source and valid window; association analyses retain temporal coverage limits. Length, content, cache residency, generation and serving paths can jointly affect results. Causal mechanism claims require isolated, matched interventions. Conclusions state the data snapshot, coverage and conditions.

GPU and system metric dictionary

Table 13–07 GPU and system metric dictionary

Metric	Definition, unit and boundary
Compute throughput	Dense GEMM ≈2MNK/t; TFLOP/s = FLOP/s / 10^{12}. State input/accumulator precision and conversion inclusion. Sparse effective–operation counts are separate.
Memory bandwidth	Effective bandwidth = (read + written bytes)/t; copy normally counts 2N bytes. Preserve a tool's one–way payload convention. Physical traffic requires valid counters.
Memory access and cache	Dependent–access latency = chain time/accesses; specify cache, translation–cache and clock state. L1/L2 hit rates retain counter–specific access denominators.
Execution and launch	CUDA events measure device time; a host monotonic clock measures submission and synchronized elapsed time. Separate graph instantiation, first execution and replay.
Concurrency and overlap	Concurrency gain is a throughput ratio at matched total work; overlap benefit = (serial—overlapped time)/serial time. Retain per–task latency.
Communication	algbw = tool–defined message bytes/t; AllReduce busbw = algbw×2(n—1)/n; other operations use their formulas. Distinguish GB/s from Gb/s.
Sustained performance	Ten–minute–window throughput / that stage's post–warmup baseline; separate compute, bandwidth and mixed workloads. Do not mix unlike work.
Energy and efficiency	Energy E=∫P(t)dt in J; efficiency = work/E. Separate chip, estimated system and outlet–metered boundaries; dynamic energy subtracts a matched idle baseline.
Numerics and reliability	Validation errors, valid/invalid rounds, observed duration, new driver faults, allocation failures and throttle reasons. Short error–free tests do not estimate MTBF.
Roofline	Arithmetic intensity I=operations/bytes at the same memory level; reference ceiling min(P,B×I). P and B use matched measured or explicitly rated ceilings; interpret proximity with execution traces.

NVIDIA DGX Spark hardware · Unified–memory reporting limits · CUTLASS architecture and dtype support · nvbandwidth methodology · NCCL communication metrics · Nsight timing, counters and roofline

14 · Appendix B: supplementary figures and statistics

These supplementary figures document distributions, block variation and sampling coverage; primary definitions and findings are in the main text.

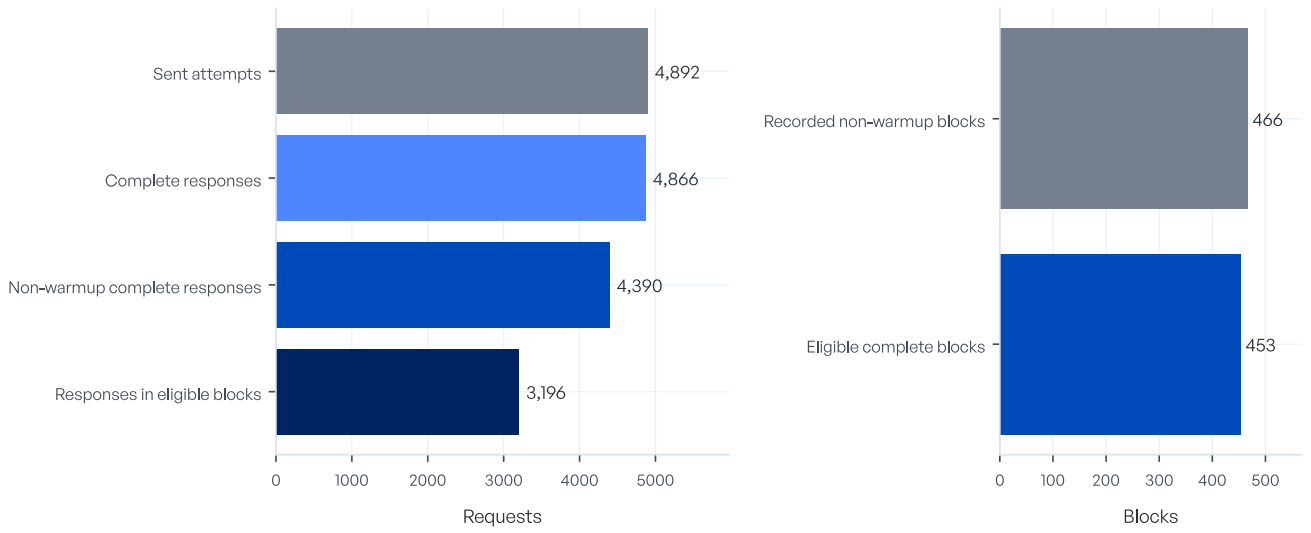


Figure 14–01 **Sample inclusion**

All attempts through eligible complete blocks; requests and blocks are separate units.

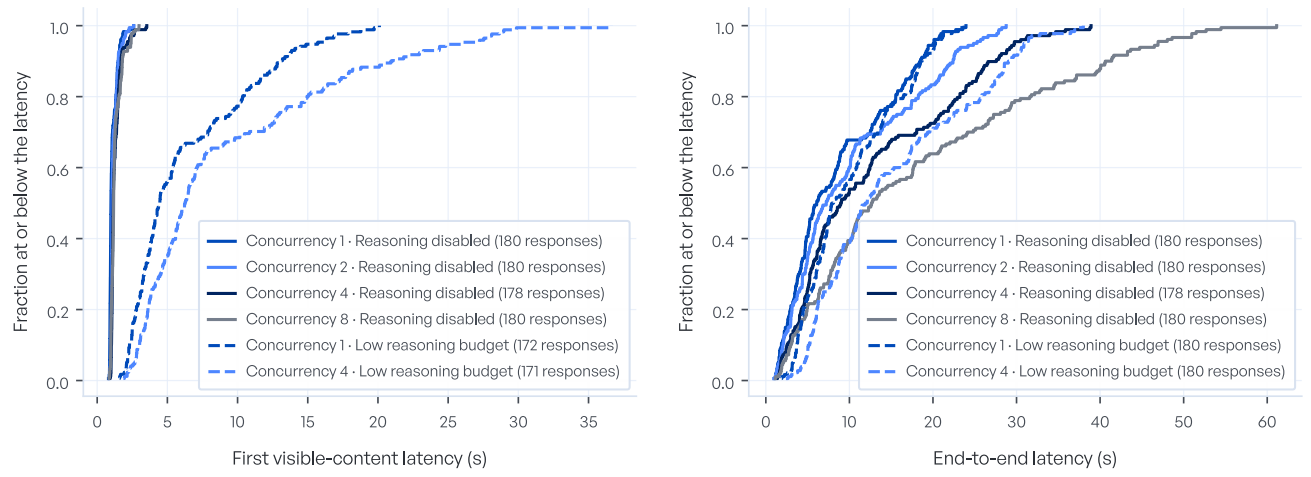


Figure 14–02 **Response distributions**

All non–warmup complete responses, including partial blocks; the population differs from the main capacity curve.

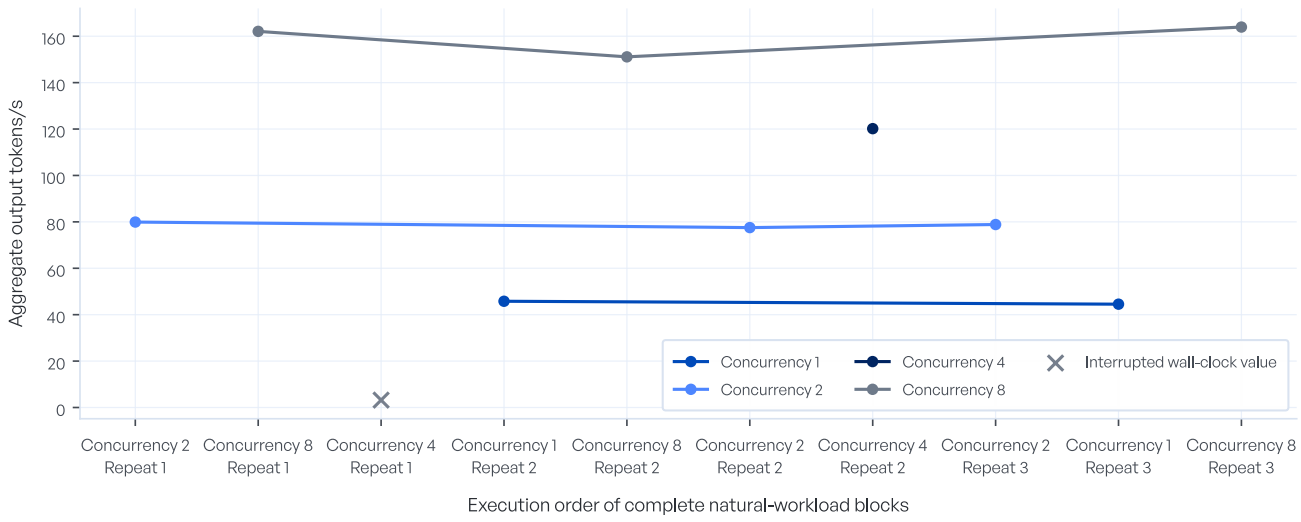


Figure 14–03 **Block variation**

Solid points are uninterrupted blocks; crosses retain interrupted wall–clock values.

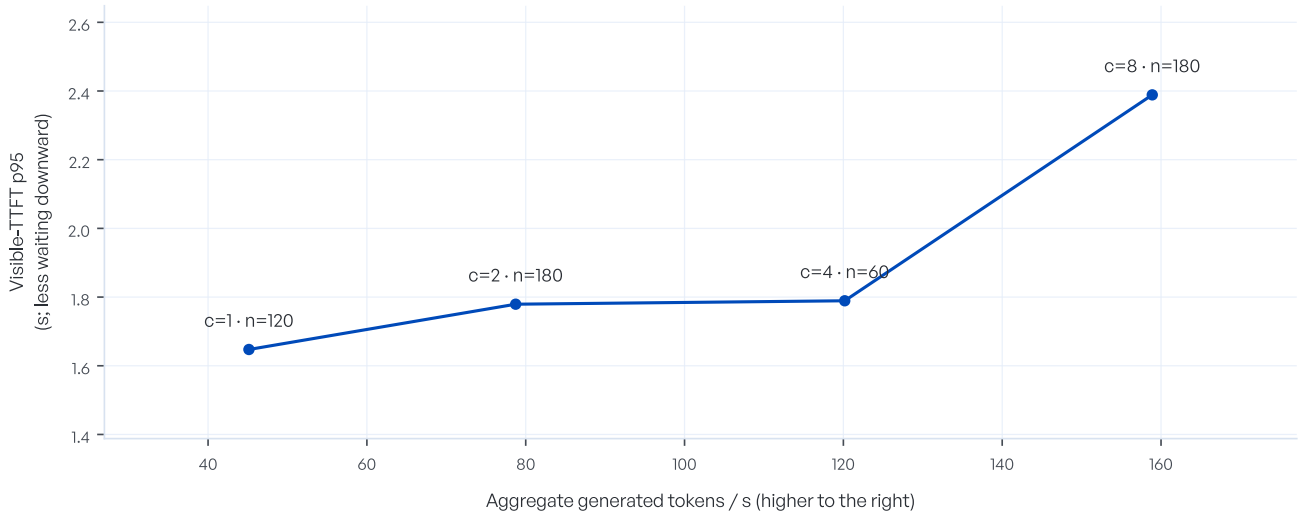


Figure 14–04 **Throughput and waiting**

One point per concurrency condition; lines only guide reading.

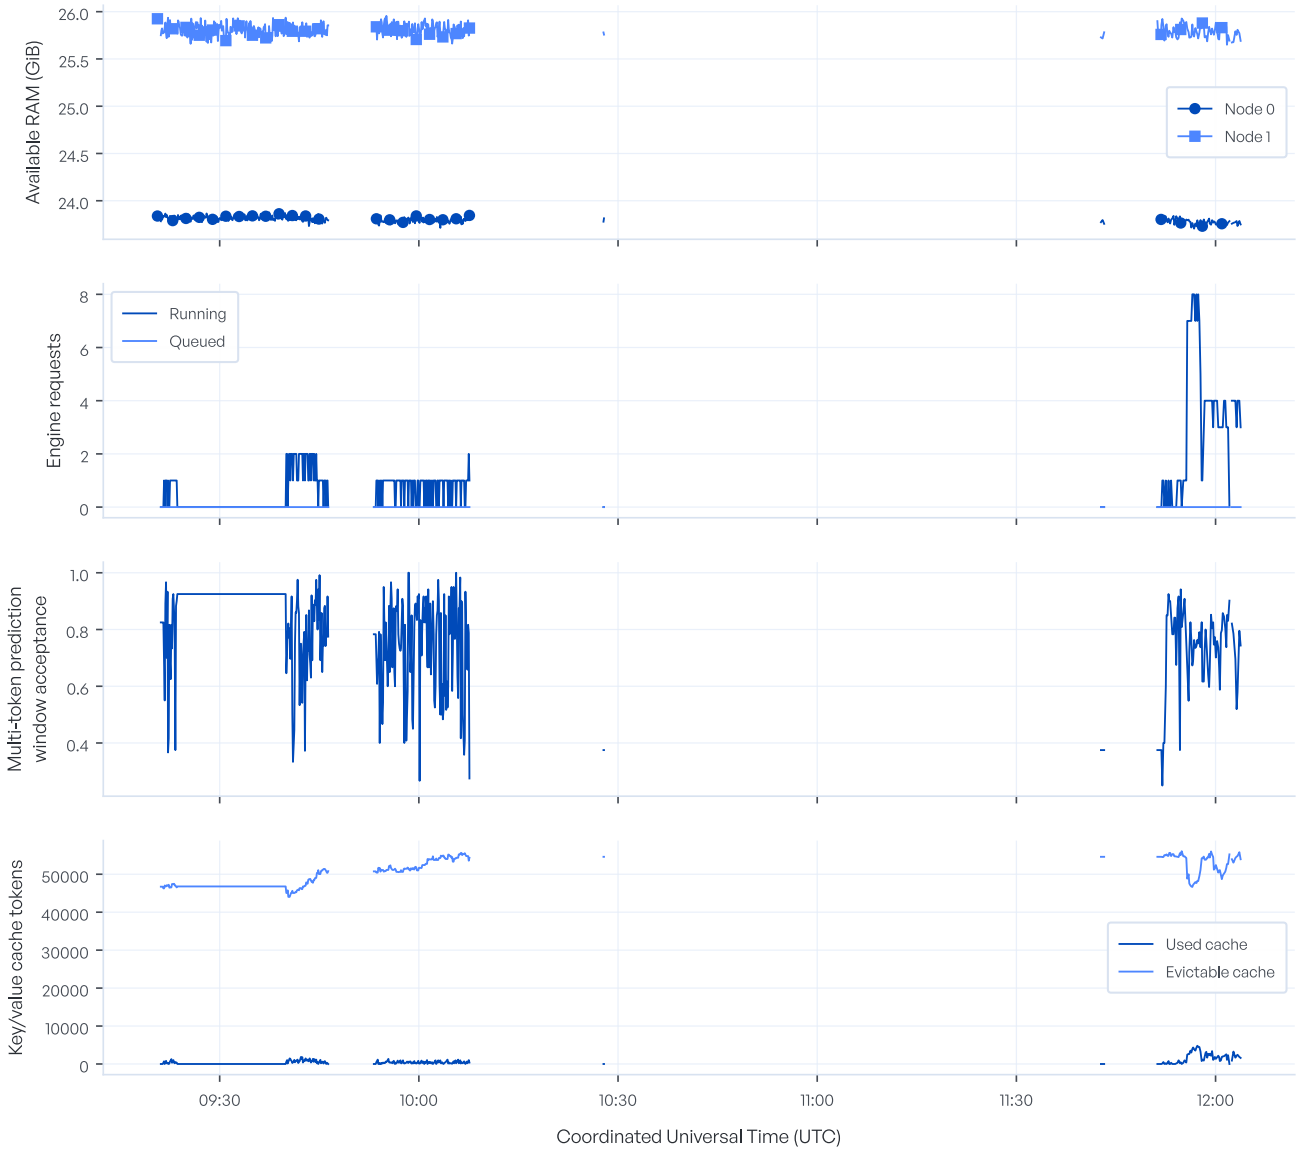


Figure 14–05 **Raw engine and memory samples**

Missing intervals remain blank; acceptance is a sampled–window field.

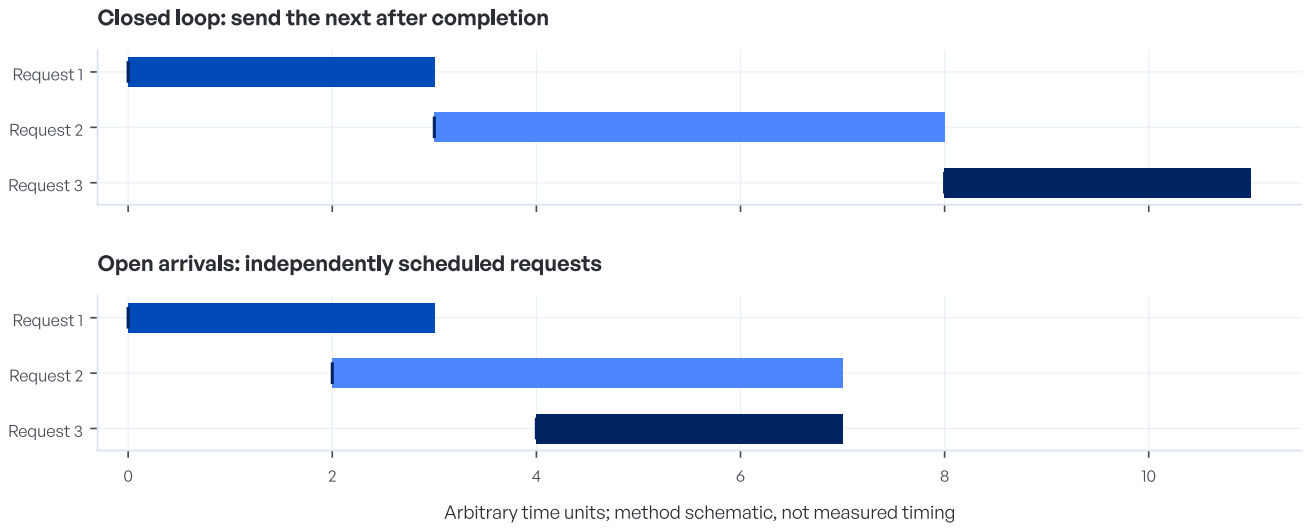


Figure 14–06 Arrival schematic

Independent arrivals can queue; this diagram does not establish test completion.

Full paired reasoning statistics

Table 14–01 Full paired reasoning statistics

Condition	Complete	Median first–content latency s	Median end–to–end latency s	Mean output	Length / integrity flags
Concurrency 1 · Low reasoning budget	180	4.42	8.00	483.6	21 / 8
Concurrency 4 · Low reasoning budget	180	6.26	12.01	489.3	22 / 9

Table 14–02 Full paired reasoning statistics

Concurrency / metric	Pairs	Mean Δ s	Prompt–cluster 95% confidence interval	Repeat–cluster 95% confidence interval
1 / First generated–fragment latency	180	0.042	–0.017—0.098	–0.057—0.125
1 / First visible–content latency	172	5.222	4.180—6.309	4.774—5.462
1 / End–to–end latency	180	1.546	0.610—2.475	1.165—1.779
4 / First generated–fragment latency	178	–0.128	–0.197—–0.058	–0.371—0.173
4 / First visible–content latency	169	7.968	6.342—9.787	7.368—8.776
4 / End–to–end latency	178	2.372	0.773—4.011	1.968—3.010

Input processing, context and reuse details

Table 14–03 Input processing, context and reuse details

Condition	Attempts	Mean input / output	Median first–generated s	Complete stream / other terminal
Long context and boundaries · Context boundary probe	1	— / —	—	0 / 1
Long context and boundaries · Target input 131,072 tokens · Concurrency 1 · Generation after long input	3	131,073.0 / 1,024.0	48.34	3 / 0
Long context and boundaries · Target input 131,072 tokens · Concurrency 2 · Generation after long input	3	131,073.0 / 1,024.0	90.65	3 / 0
Long context and boundaries · Target input 16,384 tokens · Concurrency 1 · Generation after long input	3	16,385.0 / 1,024.0	6.77	3 / 0
Long context and boundaries · Target input 16,384 tokens · Concurrency 2 · Generation after long input	3	16,384.7 / 1,024.0	11.28	3 / 0
Long context and boundaries · Target input 200,000 tokens · Concurrency 1 · Generation after long input	3	200,000.3 / 1,024.0	74.05	3 / 0
Long context and boundaries · Target input 250,000 tokens · Concurrency 1 · Generation after long input	3	250,000.7 / 1,024.0	94.08	3 / 0
Long context and boundaries · Target input 32,768 tokens · Concurrency 1 · Generation after long input	3	32,769.0 / 1,024.0	12.55	3 / 0
Long context and boundaries · Target input 32,768 tokens · Concurrency 2 · Generation after long input	3	32,769.0 / 1,024.0	22.01	3 / 0
Long context and boundaries · Target input 4,096 tokens · Concurrency 1 · Generation after long input	3	4,096.7 / 1,024.0	2.54	3 / 0
Long context and boundaries · Target input 4,096 tokens · Concurrency 2 · Generation after long input	3	4,097.0 / 1,024.0	2.41	3 / 0
Long context and boundaries · Target input 65,536 tokens · Concurrency 1 · Generation after long input	3	65,536.3 / 1,024.0	24.49	3 / 0
Long context and boundaries · Target input 65,536 tokens · Concurrency 2 · Generation after long input	5	65,537.0 / 1,024.0	42.67	3 / 2
Prefix caching · History text 0 tokens · First prefix	10	148.0 / 210.8	1.07	10 / 0
Prefix caching · History text 0 tokens · Repeated prefix	10	148.0 / 213.8	0.96	10 / 0
Prefix caching · History text 0 tokens · Appended text	10	404.7 / 206.9	1.06	10 / 0
Prefix caching · History text 32,768 tokens · First prefix	10	32,917.2 / 229.3	11.69	10 / 0
Prefix caching · History text 32,768 tokens · Repeated prefix	10	32,917.2 / 262.2	1.55	10 / 0
Prefix caching · History text 32,768 tokens · Appended text	10	33,174.1 / 193.6	1.61	10 / 0
Prefix caching · History text 8,192 tokens · First prefix	10	8,340.8 / 222.1	3.61	10 / 0
Prefix caching · History text 8,192 tokens · Repeated prefix	10	8,340.8 / 234.5	1.08	10 / 0
Prefix caching · History text 8,192 tokens · Appended text	10	8,597.6 / 222.3	1.17	10 / 0

Lengths and latency use complete streams. Other terminal states include expected boundary rejections and are not all ordinary request failures. Cache fields remain in per-request data; faster responses alone do not prove cache hits or retention.

Complete statistics for key matrix shapes

Table 14-04 Complete statistics for key matrix shapes

Node	Precision	M×N×K	Rounds	Median TFLOP/s	Observed range
0	fp32	128×4096×4096	10	13.127	12.758–13.267
0	fp32	4096×4096×4096	10	16.762	16.628–16.861
0	fp16	1×4096×4096	10	0.276	0.227–0.304
0	fp16	4096×4096×4096	10	87.556	86.791–88.325
0	tf32_allowed	128×4096×4096	10	15.154	14.365–15.389
0	fp32	1×4096×4096	10	0.123	0.119–0.126
0	bf16	1×4096×4096	10	0.284	0.265–0.292
0	bf16	128×4096×4096	10	33.110	31.757–34.217
0	bf16	4096×4096×4096	10	88.641	86.682–89.490
0	fp16	128×4096×4096	10	32.312	30.924–34.355
0	tf32_allowed	1×4096×4096	10	0.125	0.119–0.127
0	tf32_allowed	4096×4096×4096	10	39.380	38.995–39.802
1	fp32	128×4096×4096	10	12.992	12.576–13.150
1	fp32	4096×4096×4096	10	16.579	16.470–16.726
1	fp16	1×4096×4096	10	0.281	0.255–0.297
1	fp16	4096×4096×4096	10	87.287	85.604–87.888
1	tf32_allowed	128×4096×4096	10	14.387	13.757–15.252
1	fp32	1×4096×4096	10	0.121	0.118–0.124
1	bf16	1×4096×4096	10	0.283	0.260–0.303
1	bf16	128×4096×4096	10	32.418	31.179–33.816
1	bf16	4096×4096×4096	10	87.727	87.410–88.985
1	fp16	128×4096×4096	10	32.275	30.953–34.407
1	tf32_allowed	1×4096×4096	10	0.122	0.119–0.128
1	tf32_allowed	4096×4096×4096	10	39.194	38.901–39.687
0	bf16_sparse_2of4	4096×4096×4096	10	34.292	33.662–35.655
0	nvfp4	1×4096×4096	10	0.767	0.766–0.772
0	nvfp4	128×4096×4096	10	98.417	97.847–98.677
0	bf16_sparse_2of4	1×4096×4096	10	0.010	0.010–0.011

Node	Precision	M×N×K	Rounds	Median TFLOP/s	Observed range
0	nvfp4	4096×4096×4096	10	318.537	311.189–322.891
0	fp16_sparse_2of4	1×4096×4096	10	0.010	0.010–0.010
0	bf16_sparse_2of4	128×4096×4096	10	1.307	1.271–1.340
0	fp8_e4m3	4096×4096×4096	10	176.414	175.105–177.438
0	fp16_sparse_2of4	128×4096×4096	10	1.320	1.297–1.330
0	fp16_sparse_2of4	4096×4096×4096	10	33.755	33.070–34.757
0	fp8_e4m3	1×4096×4096	10	0.467	0.465–0.468
0	fp8_e4m3	128×4096×4096	10	118.961	116.431–120.597
1	bf16_sparse_2of4	4096×4096×4096	10	34.179	32.971–35.480
1	nvfp4	1×4096×4096	10	0.771	0.766–0.776
1	nvfp4	128×4096×4096	10	98.541	97.539–98.784
1	bf16_sparse_2of4	1×4096×4096	10	0.010	0.010–0.011
1	nvfp4	4096×4096×4096	10	317.672	308.915–320.311
1	fp16_sparse_2of4	1×4096×4096	10	0.010	0.010–0.010
1	bf16_sparse_2of4	128×4096×4096	10	1.285	1.245–1.312
1	fp8_e4m3	4096×4096×4096	10	175.027	170.906–176.635
1	fp16_sparse_2of4	128×4096×4096	10	1.305	1.271–1.314
1	fp16_sparse_2of4	4096×4096×4096	10	33.868	32.344–34.901
1	fp8_e4m3	1×4096×4096	10	0.464	0.463–0.465
1	fp8_e4m3	128×4096×4096	10	116.633	115.258–117.238

15 · Appendix C: test code and data archive index

This appendix retains the original study archive index. The GitHub project distributes principal algorithms, measurement tables, protocols and full reports. Raw HTTP bodies, large traces, production controllers and the complete local inventory are not included. Archive paths are provenance references, not attachments to this HTML.

The inventory contains **50,430 files**; scan completed (UTC): 2026-10-17T08:58:51.005978+00:00

Table 15-01 Appendix C: test code and data archive index

Contents	Archive path
Current experiment root	export/qwen38_perf_p3_20261014_0905/
Model, dataset and protocol source root	export/qwen38_abliterated_eval_20261014_mtp/
Current execution test plan	export/qwen38_perf_p3_20261014_0905/test_plan/test-plan.en.html
Searchable complete file inventory	export/qwen38_perf_p3_20261014_0905/reports/artifact-index.en.html
Complete path inventory (CSV)	export/qwen38_perf_p3_20261014_0905/reports/artifact-index.csv
Complete paths and metadata (JSON)	export/qwen38_perf_p3_20261014_0905/reports/artifact-index.json

Test code

Workload construction and execution

Table 15-02 Workload construction and execution

Contents	Archive path
length_plan.py	export/qwen38_perf_p3_20261014_0905/length_plan.py
length_runner.py	export/qwen38_perf_p3_20261014_0905/length_runner.py
length_rejections.py	export/qwen38_perf_p3_20261014_0905/length_rejections.py
prepare_performance.py	export/qwen38_perf_p3_20261014_0905/prepare_performance.py
business_plan.py	export/qwen38_perf_p3_20261014_0905/business_plan.py
refine_business_plan.py	export/qwen38_perf_p3_20261014_0905/refine_business_plan.py
runner.py	export/qwen38_perf_p3_20261014_0905/runner.py
driver.py	export/qwen38_perf_p3_20261014_0905/driver.py
business_runner.py	export/qwen38_perf_p3_20261014_0905/business_runner.py
extended.py	export/qwen38_perf_p3_20261014_0905/extended.py
arrival.py	export/qwen38_perf_p3_20261014_0905/arrival.py

Monitoring and run management

Table 15-03 Monitoring and run management

Contents	Archive path
cache_runner.py	export/qwen38_perf_p3_20261014_0905/cache_runner.py
cache_counters.py	export/qwen38_perf_p3_20261014_0905/cache_counters.py
collector.py	export/qwen38_perf_p3_20261014_0905/collector.py
public_monitor.py	export/qwen38_perf_p3_20261014_0905/public_monitor.py
hardware_diag.py	export/qwen38_perf_p3_20261014_0905/hardware_diag.py
observe_tunnel.py	export/qwen38_perf_p3_20261014_0905/observe_tunnel.py
supervise.py	export/qwen38_perf_p3_20261014_0905/supervise.py
recover.py	export/qwen38_perf_p3_20261014_0905/recover.py
continuation.py	export/qwen38_perf_p3_20261014_0905/continuation.py
cf_retry.py	export/qwen38_perf_p3_20261014_0905/cf_retry.py
terminal_review.py	export/qwen38_perf_p3_20261014_0905/terminal_review.py
incident_report.py	export/qwen38_perf_p3_20261014_0905/incident_report.py
common.py	export/qwen38_perf_p3_20261014_0905/common.py

Statistics, token counts and billing

Table 15-04 Statistics, token counts and billing

Contents	Archive path
cache_analysis.py	export/qwen38_perf_p3_20261014_0905/cache_analysis.py
analyze.py	export/qwen38_perf_p3_20261014_0905/analyze.py
metrics.py	export/qwen38_perf_p3_20261014_0905/metrics.py
origin.py	export/qwen38_perf_p3_20261014_0905/origin.py
reconcile.py	export/qwen38_perf_p3_20261014_0905/reconcile.py
business_results.py	export/qwen38_perf_p3_20261014_0905/business_results.py
report_model.py	export/qwen38_perf_p3_20261014_0905/report_model.py
report_diagnostics.py	export/qwen38_perf_p3_20261014_0905/report_diagnostics.py
report_cost_model.py	export/qwen38_perf_p3_20261014_0905/report_cost_model.py

Report and visualization

Table 15-05 Report and visualization

Contents	Archive path
report_benchmarks.py	export/qwen38_perf_p3_20261014_0905/report_benchmarks.py

Contents	Archive path
report_gpu_system.py	export/qwen38_perf_p3_20261014_0905/report_gpu_system.py
system_figures.py	export/qwen38_perf_p3_20261014_0905/system_figures.py
integrate_system_plan.py	export/qwen38_perf_p3_20261014_0905/integrate_system_plan.py
length_analysis.py	export/qwen38_perf_p3_20261014_0905/length_analysis.py
report_length.py	export/qwen38_perf_p3_20261014_0905/report_length.py
length_figures.py	export/qwen38_perf_p3_20261014_0905/length_figures.py
cache_plan.py	export/qwen38_perf_p3_20261014_0905/cache_plan.py
report_cache.py	export/qwen38_perf_p3_20261014_0905/report_cache.py
cache_figures.py	export/qwen38_perf_p3_20261014_0905/cache_figures.py
report_endpoint_details.py	export/qwen38_perf_p3_20261014_0905/report_endpoint_details.py
report_emphasis.py	export/qwen38_perf_p3_20261014_0905/report_emphasis.py
build_report.py	export/qwen38_perf_p3_20261014_0905/build_report.py
plots.py	export/qwen38_perf_p3_20261014_0905/plots.py
research_figures.py	export/qwen38_perf_p3_20261014_0905/research_figures.py
diagnostic_figures.py	export/qwen38_perf_p3_20261014_0905/diagnostic_figures.py
report_foundations.py	export/qwen38_perf_p3_20261014_0905/report_foundations.py
report_measurements.py	export/qwen38_perf_p3_20261014_0905/report_measurements.py
report_narrative.py	export/qwen38_perf_p3_20261014_0905/report_narrative.py
report_results_expansion.py	export/qwen38_perf_p3_20261014_0905/report_results_expansion.py
report_returns.py	export/qwen38_perf_p3_20261014_0905/report_returns.py
report_editorial.py	export/qwen38_perf_p3_20261014_0905/report_editorial.py
build_reviewed_plan.py	export/qwen38_perf_p3_20261014_0905/build_reviewed_plan.py
report_labels.py	export/qwen38_perf_p3_20261014_0905/report_labels.py
report_chart_catalog.py	export/qwen38_perf_p3_20261014_0905/report_chart_catalog.py
report_cost_widget.py	export/qwen38_perf_p3_20261014_0905/report_cost_widget.py
report_appendix.py	export/qwen38_perf_p3_20261014_0905/report_appendix.py

Offline validation and other scripts

Table 15-06 Offline validation and other scripts

Contents	Archive path
boundary_review.py	export/qwen38_perf_p3_20261014_0905/boundary_review.py
cf_edge_review.py	export/qwen38_perf_p3_20261014_0905/cf_edge_review.py
engine_split_report.py	export/qwen38_perf_p3_20261014_0905/engine_split_report.py

Contents	Archive path
finalize_hardware.py	export/qwen38_perf_p3_20261014_0905/finalize_hardware.py
gpu_system_bench.py	export/qwen38_perf_p3_20261014_0905/gpu_system_bench.py
gpu_system_collectives.py	export/qwen38_perf_p3_20261014_0905/gpu_system_collectives.py
gpu_system_control.py	export/qwen38_perf_p3_20261014_0905/gpu_system_control.py
gpu_system_guard.py	export/qwen38_perf_p3_20261014_0905/gpu_system_guard.py
gpu_system_observer.py	export/qwen38_perf_p3_20261014_0905/gpu_system_observer.py
gpu_system_preflight.py	export/qwen38_perf_p3_20261014_0905/gpu_system_preflight.py
gpu_system_remote.py	export/qwen38_perf_p3_20261014_0905/gpu_system_remote.py
hardware_collective_analysis.py	export/qwen38_perf_p3_20261014_0905/hardware_collective_analysis.py
hardware_completion_report.py	export/qwen38_perf_p3_20261014_0905/hardware_completion_report.py
hardware_extension_figures.py	export/qwen38_perf_p3_20261014_0905/hardware_extension_figures.py
hardware_h200_report.py	export/qwen38_perf_p3_20261014_0905/hardware_h200_report.py
hardware_partial.py	export/qwen38_perf_p3_20261014_0905/hardware_partial.py
quarantine.py	export/qwen38_perf_p3_20261014_0905/quarantine.py
redraw_report_figures.py	export/qwen38_perf_p3_20261014_0905/redraw_report_figures.py
report_coverage.py	export/qwen38_perf_p3_20261014_0905/report_coverage.py
report_gap_completion.py	export/qwen38_perf_p3_20261014_0905/report_gap_completion.py
report_overview.py	export/qwen38_perf_p3_20261014_0905/report_overview.py
report_review.py	export/qwen38_perf_p3_20261014_0905/report_review.py
report_visual_layout.py	export/qwen38_perf_p3_20261014_0905/report_visual_layout.py
report_visual_theme.py	export/qwen38_perf_p3_20261014_0905/report_visual_theme.py
review_figures.py	export/qwen38_perf_p3_20261014_0905/review_figures.py
test_arrival.py	export/qwen38_perf_p3_20261014_0905/test_arrival.py
test_business.py	export/qwen38_perf_p3_20261014_0905/test_business.py
test_cache_diagnostics.py	export/qwen38_perf_p3_20261014_0905/test_cache_diagnostics.py
test_cf_retry.py	export/qwen38_perf_p3_20261014_0905/test_cf_retry.py
test_common_atomic.py	export/qwen38_perf_p3_20261014_0905/test_common_atomic.py
test_incident.py	export/qwen38_perf_p3_20261014_0905/test_incident.py
test_metrics.py	export/qwen38_perf_p3_20261014_0905/test_metrics.py
test_public_monitor.py	export/qwen38_perf_p3_20261014_0905/test_public_monitor.py
test_quarantine.py	export/qwen38_perf_p3_20261014_0905/test_quarantine.py
test_recover.py	export/qwen38_perf_p3_20261014_0905/test_recover.py
test_report_diagnostics.py	export/qwen38_perf_p3_20261014_0905/test_report_diagnostics.py

Contents	Archive path
test_runner.py	export/qwen38_perf_p3_20261014_0905/test_runner.py
test_terminal_review.py	export/qwen38_perf_p3_20261014_0905/test_terminal_review.py
validate_h200_comparison.py	export/qwen38_perf_p3_20261014_0905/validate_h200_comparison.py
verify_scientific_report.py	export/qwen38_perf_p3_20261014_0905/verify_scientific_report.py
verify_visual_system.py	export/qwen38_perf_p3_20261014_0905/verify_visual_system.py

Shared code, execution snapshots and runtime environment

Table 15-07 Shared code, execution snapshots and runtime environment

Contents	Archive path
Project settings entry point	src/insightloop/config/loader.py
ClickHouse read-query client entry point	src/insightloop/runtime/clickhouse.py
ClickHouse client implementation	src/insightloop/datasources/clickhouse.py
Core execution source snapshot	export/qwen38_perf_p3_20261014_0905/evidence/executor_core/
Frozen business execution source	export/qwen38_perf_p3_20261014_0905/evidence/business_amendment/
Execution source and configuration evidence	export/qwen38_perf_p3_20261014_0905/evidence/
Test Python interpreter	/tmp/gp-qwen38-eval/bin/python
Report Python interpreter	/tmp/growthpulse-morphic-report-311/bin/python
Project dependencies	pyproject.toml
Project dependency lock	uv.lock
Recorded source environment versions	export/qwen38_abliterated_eval_20261014_mtp/sources/environment_versions.json

Input data, sampling and frozen request lists

Table 15-08 Input data, sampling and frozen request lists

Contents	Archive path
GPU/system module design manifest	export/qwen38_perf_p3_20261014_0905/plans/gpu_system/manifest.json
GPU-method official-source snapshots	export/qwen38_perf_p3_20261014_0905/evidence/gpu_system_revision/sources/
Length expansion requests, parameters and hashes	export/qwen38_perf_p3_20261014_0905/plans/length_expansion/
Expanded length observations	export/qwen38_perf_p3_20261014_0905/reports/length-observations.csv
Framework methodology sources	export/qwen38_perf_p3_20261014_0905/evidence/length_expansion/sources/
Fixed 60 prompts from six suites	export/qwen38_perf_p3_20261014_0905/protocol/workloads.jsonl

Contents	Archive path
Prompt sampling and length manifest	export/qwen38_perf_p3_20261014_0905/protocol/workload_manifest.json
Benchmark size and selected-sample audit	export/qwen38_perf_p3_20261014_0905/evidence/benchmark_documentation/inventory.json
Original dataset directory	export/qwen38_abliterated_eval_20261014_mtp/datasets/
MMLU-Pro source data	export/qwen38_abliterated_eval_20261014_mtp/datasets/TIGER-Lab__MMLU-Pro/
MATH-500 source data	export/qwen38_abliterated_eval_20261014_mtp/datasets/HuggingFaceH4__MATH-500/
GSM8K source data	export/qwen38_abliterated_eval_20261014_mtp/datasets/openai__gsm8k/
IFEval source inputs	export/qwen38_abliterated_eval_20261014_mtp/sources/instruction_following_eval/data/input_data.jsonl
HumanEval+ source data	export/qwen38_abliterated_eval_20261014_mtp/datasets/humaneval_plus.jsonl
MBPP+ source data	export/qwen38_abliterated_eval_20261014_mtp/datasets/mbpp_plus.jsonl
Length-test text corpus	export/qwen38_perf_p3_20261014_0905/plans/corpus.txt
Tokenized length-test corpus	export/qwen38_perf_p3_20261014_0905/plans/corpus_tokens.json
Frozen cache diagnostic requests and design	export/qwen38_perf_p3_20261014_0905/plans/cache_diagnostics/
Complete request-plan directory	export/qwen38_perf_p3_20261014_0905/plans/
Core schedule and seeds	export/qwen38_perf_p3_20261014_0905/plans/schedule.json
Frozen request-file hashes	export/qwen38_perf_p3_20261014_0905/plans/frozen.json
Business and prefix-ratio requests	export/qwen38_perf_p3_20261014_0905/plans/business/
Business shapes and sample counts	export/qwen38_perf_p3_20261014_0905/plans/business/manifest.json
Model configuration, tokenizer and chat template	export/qwen38_abliterated_eval_20261014_mtp/sources/model/
Model and dataset source hashes	export/qwen38_abliterated_eval_20261014_mtp/sources/manifest.json
Full protocol and references	export/qwen38_abliterated_eval_20261014_mtp/test_plan/

Raw responses, telemetry and run records

Table 15-09 Raw responses, telemetry and run records

Contents	Archive path
Per-request raw data directory	export/qwen38_perf_p3_20261014_0905/runs/requests/
Block execution records	export/qwen38_perf_p3_20261014_0905/runs/blocks/
Before/after block resource boundaries	export/qwen38_perf_p3_20261014_0905/runs/boundaries/
Request ledger and reservations	export/qwen38_perf_p3_20261014_0905/runs/ledger.json
Logical-case outcomes	export/qwen38_perf_p3_20261014_0905/runs/logical_outcomes/

Contents	Archive path
Raw node-dashboard samples	export/qwen38_perf_p3_20261014_0905/monitor/sparkdash.jsonl
Raw engine and operating-system samples	export/qwen38_perf_p3_20261014_0905/monitor/engine-os.jsonl
Client activity samples	export/qwen38_perf_p3_20261014_0905/monitor/client.jsonl
Cache diagnostic windows and observed timing	export/qwen38_perf_p3_20261014_0905/runs/cache_windows/
Raw cache counters and freshness	export/qwen38_perf_p3_20261014_0905/monitor/cache_counters.jsonl
Cache provenance and execution evidence	export/qwen38_perf_p3_20261014_0905/evidence/cache_evaluation/
All run logs and receipts	export/qwen38_perf_p3_20261014_0905/runs/
All monitoring files	export/qwen38_perf_p3_20261014_0905/monitor/
Exact-request billing records	export/qwen38_perf_p3_20261014_0905/evidence/billing_latest.json
Configuration, counts and provenance evidence	export/qwen38_perf_p3_20261014_0905/evidence/

Files for one request

The paths below show one saved request. Other requests use the same directory structure and are individually searchable in the complete inventory.

Table 15-10 Files for one request

Contents	Archive path
Sent input and parameters	export/qwen38_perf_p3_20261014_0905/runs/requests/P0_0_HumanEval_51/request.json
Send-start record	export/qwen38_perf_p3_20261014_0905/runs/requests/P0_0_HumanEval_51/started.json
Raw server-sent-event response stream	export/qwen38_perf_p3_20261014_0905/runs/requests/P0_0_HumanEval_51/response.sse
Parsed stream events	export/qwen38_perf_p3_20261014_0905/runs/requests/P0_0_HumanEval_51/events.jsonl
Chunk receive timing	export/qwen38_perf_p3_20261014_0905/runs/requests/P0_0_HumanEval_51/chunks.jsonl
Final state, usage and latency	export/qwen38_perf_p3_20261014_0905/runs/requests/P0_0_HumanEval_51/result.json

Statistics, figures and delivered reports

Table 15-11 Statistics, figures and delivered reports

Contents	Archive path
Complete cache analysis	export/qwen38_perf_p3_20261014_0905/reports/cache_analysis.json
Cache condition results	export/qwen38_perf_p3_20261014_0905/reports/cache_conditions.csv
Cache per-request results	export/qwen38_perf_p3_20261014_0905/reports/cache_requests.csv
Cache time windows	export/qwen38_perf_p3_20261014_0905/reports/cache_time_windows.csv

Contents	Archive path
Cache field coverage	export/qwen38_perf_p3_20261014_0905/reports/cache_field_coverage.csv
First/repeat pairs	export/qwen38_perf_p3_20261014_0905/reports/cache_reuse_pairs.csv
Per-request analysis	export/qwen38_perf_p3_20261014_0905/reports/requests.csv
Complete analysis	export/qwen38_perf_p3_20261014_0905/reports/analysis.json
Eligible complete-block population	export/qwen38_perf_p3_20261014_0905/reports/presentation.json
Joined cost, length and resource data	export/qwen38_perf_p3_20261014_0905/reports/diagnostics.json
Reconciled cost by workload	export/qwen38_perf_p3_20261014_0905/reports/cohort-costs.csv
Block-aligned node resources	export/qwen38_perf_p3_20261014_0905/reports/cohort-resources.csv
Length matrix	export/qwen38_perf_p3_20261014_0905/reports/length-metrics.csv
Device and electricity scenarios	export/qwen38_perf_p3_20261014_0905/reports/cost-scenarios.csv
Investment-return and payback scenarios	export/qwen38_perf_p3_20261014_0905/reports/investment-scenarios.csv
Business and cache measurements	export/qwen38_perf_p3_20261014_0905/reports/business_results.json
Hardware timeline	export/qwen38_perf_p3_20261014_0905/reports/hardware.csv
Engine and operating-system timeline	export/qwen38_perf_p3_20261014_0905/reports/engine-os.csv
Qualified throughput	export/qwen38_perf_p3_20261014_0905/reports/goodput.csv
Chinese report (HTML)	export/qwen38_perf_p3_20261014_0905/reports/report.zh.html
Chinese report (MD)	export/qwen38_perf_p3_20261014_0905/reports/report.zh.md
Chinese report (PDF)	export/qwen38_perf_p3_20261014_0905/reports/report.zh.pdf
English report (HTML)	export/qwen38_perf_p3_20261014_0905/reports/report.en.html
English report (MD)	export/qwen38_perf_p3_20261014_0905/reports/report.en.md
English report (PDF)	export/qwen38_perf_p3_20261014_0905/reports/report.en.pdf
Standalone SVG/PNG figures and report directory	export/qwen38_perf_p3_20261014_0905/reports/
Report validation record	export/qwen38_perf_p3_20261014_0905/reports/scientific_qa.json

The inventory reflects its scan time; live files may continue growing. Credentials, lock/temp files, bytecode and third-party installation directories are excluded. Historical material in the source root is not automatically included in this experiment's performance statistics.

Hardware completion: code, data and reproduction paths

Table 15–12 Hardware completion: code, data and reproduction paths

Purpose	Archive path
Instruction, scope and fixed deadline	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/authorization.json
Base conditions and guards	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/manifest.json
Extension conditions, repeats and source hashes	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/extensions_manifest.json
Bounded server/LAN coordination	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/remote_coordinator.py
External status collection and phase archival	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/collect.py
Base compute/bandwidth/operators and sustained work	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/gpu_system_bench.py
Extension execution and numerical checks	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/gpu_extensions.py
Scalar FMA, pointer and shared-memory CUDA source	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/kernels.cu
Four-arm compute/communication pairs	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/gpu_overlap.py
Per-phase two-node raw telemetry, timings and logs	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/raw
Per-record CSV, summaries, coverage and energy	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/analysis
Offline report update and verification	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/finalize.py
Service capture, guarded restoration and log redaction	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/service_custodian.py
Final two-node restoration verification	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/service_restoration_verification.json
Source hashes, launch and validation receipts	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/evidence
Automatic thermal backoff and events	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/thermal_backoff.py
Effective thermal thresholds and recovery rules	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/thermal_backoff_policy.json
Finite mixed queue, retries and terminal validation	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/resume_thermal_retry.py
Failure classification and retry admission	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/retry_policy.py
Live retry-supervisor state	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/retry_supervisor_status.json

Purpose	Archive path
Continuation authorization, launch, failure decisions and restoration evidence	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/evidence/thermal_retry_continuation
Overnight collection/report retries and coverage verification	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/overnight_finish.py
Bounded counter diagnostic recheck	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/counter_retry_remote.py
Completion coverage, diagnostic counters and terminal verification evidence	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/evidence/overnight_completion
Initialization-error diagnosis for the fixed counter probe	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/evidence/overnight_completion/diagnose_counter_verbose.py
CUDA-status diagnostic entry point preserving the original kernel	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/evidence/overnight_completion/counter_probe_diagnostic.cu
Post-diagnostic serving verification and study-observer closure	export/qwen38_perf_p3_20261014_0905/hardware_unloaded_confirmed_20261016/evidence/overnight_completion/close_observers.py

Remote directory on both nodes: /home/benchmark/infron-benchmarks/hardware_unloaded_confirmed_20261016

H200 comparison code, specifications and fixed-condition evidence: export/qwen38_perf_p3_20261014_0905/hardware_h200_report.py · export/qwen38_perf_p3_20261014_0905/evidence/h200_comparison

Resident-condition evidence: export/qwen38_perf_p3_20261014_0905/hardware_completion_20261016

Maintenance preflight and control-network evidence: export/qwen38_perf_p3_20261014_0905/hardware_unloaded_20261016

Independent engine code and raw evidence

Local root: export/qwen38_perf_p3_20261014_0905/engine_split_20261016

Executor: engine_probe.py and run_engine.py; collector: collect_engine.py; offline analysis: analyze_engine.py; raw requests/responses and GPU traces: raw/node0 and raw/node1; timing tables: analysis/*.csv; installed-source snapshot: evidence/source. Remote directory: /home/benchmark/infron-benchmarks/engine-split-20261016

Full SGLang launch and resolved settings

Configuration recorded (UTC): 2026-10-14T09:08:19.004649+00:00。 Launch arguments and resolved engine values are distinguished below; values not explicitly passed come from the engine configuration snapshot.

Table 15-13 Full SGLang launch and resolved settings

Item	Recorded value
SGLang version	`0.5.21`
Runtime entrypoint	`/home/benchmark/Projects/qwen38/.sglang/bin/python3 -m sglang.launch_server`
Ranks and served name	`tp_size=2; nnodes=2; node_rank=0 / 1; served_model_name=qwen38-flash`
Listener and distributed initialization	`host=0.0.0.0; port=8000; dist_init_addr=192.0.2.10:5000`

Item	Recorded value
Prefill and running requests	`chunked_prefill_size=8192; max_prefill_tokens=16384; max_running_requests=48`
KV pages and eviction	`page_size=64; kv_cache_dtype=auto; disable_radix_cache=false; radix_eviction_policy=lru`
Scheduling and CUDA Graph	`schedule_policy=fcfs; schedule_conservativeness=1.0; disable_overlap_schedule=false; enable_mixed_chunk=false; disable_cuda_graph=false`
Compute backends	`attention_backend=flashinfer; sampling_backend=flashinfer; moe_runner_backend=flashinfer_cutlass; moe_a2a_backend=none; --fp4-gemm-backend flashinfer_cutlass`
MTP launch / resolved value	`--speculative-algorithm NEXTN → speculative_algorithm=EAGLE; num_steps=3; eagle_topk=1; num_draft_tokens=4`
Output parsers	`reasoning_parser=qwen3; tool_call_parser=qwen3_coder`
Metrics and cache reporting	`enable_metrics=true; enable_cache_report=false`
Other resolved settings	`trust_remote_code=true; random_seed=523471992`

Node 0 launch command (credential redacted); node 1 uses the same arguments with `--node-rank 1`.

```
MODEL_PATH=/home/benchmark/.cache/huggingface/hub/
MODEL_PATH+=models--dealignai--Qwen3.8-Flash-Next-ABLITERATED-NVFP4/
MODEL_PATH+=snapshots/be794b990578ef3031eccf9f28e675a289a09ee9/
/home/benchmark/Projects/qwen38/.sglang/bin/python3 -m sglang.launch_server \
  --model-path "$MODEL_PATH" \
  --served-model-name qwen38-flash \
  --api-key '[REDACTED]' \
  --quantization modelopt_fp4 \
  --fp4-gemm-backend flashinfer_cutlass \
  --tp 2 \
  --nnodes 2 \
  --node-rank 0 \
  --dist-init-addr 192.0.2.10:5000 \
  --page-size 64 \
  --context-length 262144 \
  --mem-fraction-static 0.75 \
  --reasoning-parser qwen3 \
  --speculative-algorithm NEXTN \
  --speculative-num-steps 3 \
  --speculative-eagle-topk 1 \
  --speculative-num-draft-tokens 4 \
  --tool-call-parser qwen3_coder \
  --trust-remote-code \
  --host 0.0.0.0 \
  --port 8000 \
  --enable-metrics
```

Prefix reuse and per–request reporting are separate settings: `disable_radix_cache=false` permits reuse, while `enable_cache_report=false` means per–request cache usage reporting was not enabled in this configuration.

Table 15–14 Full SGLang launch and resolved settings

Control network for the hardware window	Effective values and scope
Inter–node control address	SGLANG_HOST_IP=192.0.2.10 / 192.0.2.11

Control network for the hardware window	Effective values and scope
Inference configuration verification	Restored model path, version, TP, MTP, context and core cache settings match the pre-stop configuration. Public-endpoint and independent-engine performance data were collected before this network binding.

Node 0 configuration and command · Node 1 command

Independent completion arm: code, raw evidence and analysis

Table 15-15 Independent completion arm: code, raw evidence and analysis

Content	Archive path
Gap audit	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/gap_audit.zh.md
Protocol	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/manifest.json
Budget amendment	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/budget_amendment.json
Supervisor	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/supervisor.py
Native fixed output	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/native_completion.py
Paired profiles and cache	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/profile_completion.py
Public endpoint	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/public_completion.py
Hardware coordinator	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/hardware_coordinator.py
FP64	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/fp64_completion.py
Counter kernels	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/counter_workload.cu
Counter worker	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/hardware_worker.py
Offline analysis	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/analyze_completion.py
Statistical summary	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/analysis/summary.json
Native request CSV	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/analysis/native_requests.csv
Node 0 raw evidence	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/raw/node0
Node 1 raw evidence	export/qwen38_perf_p3_20261014_0905/gap_completion_20261017/raw/node1